

# EL APRENDIZAJE AUTOMATICO

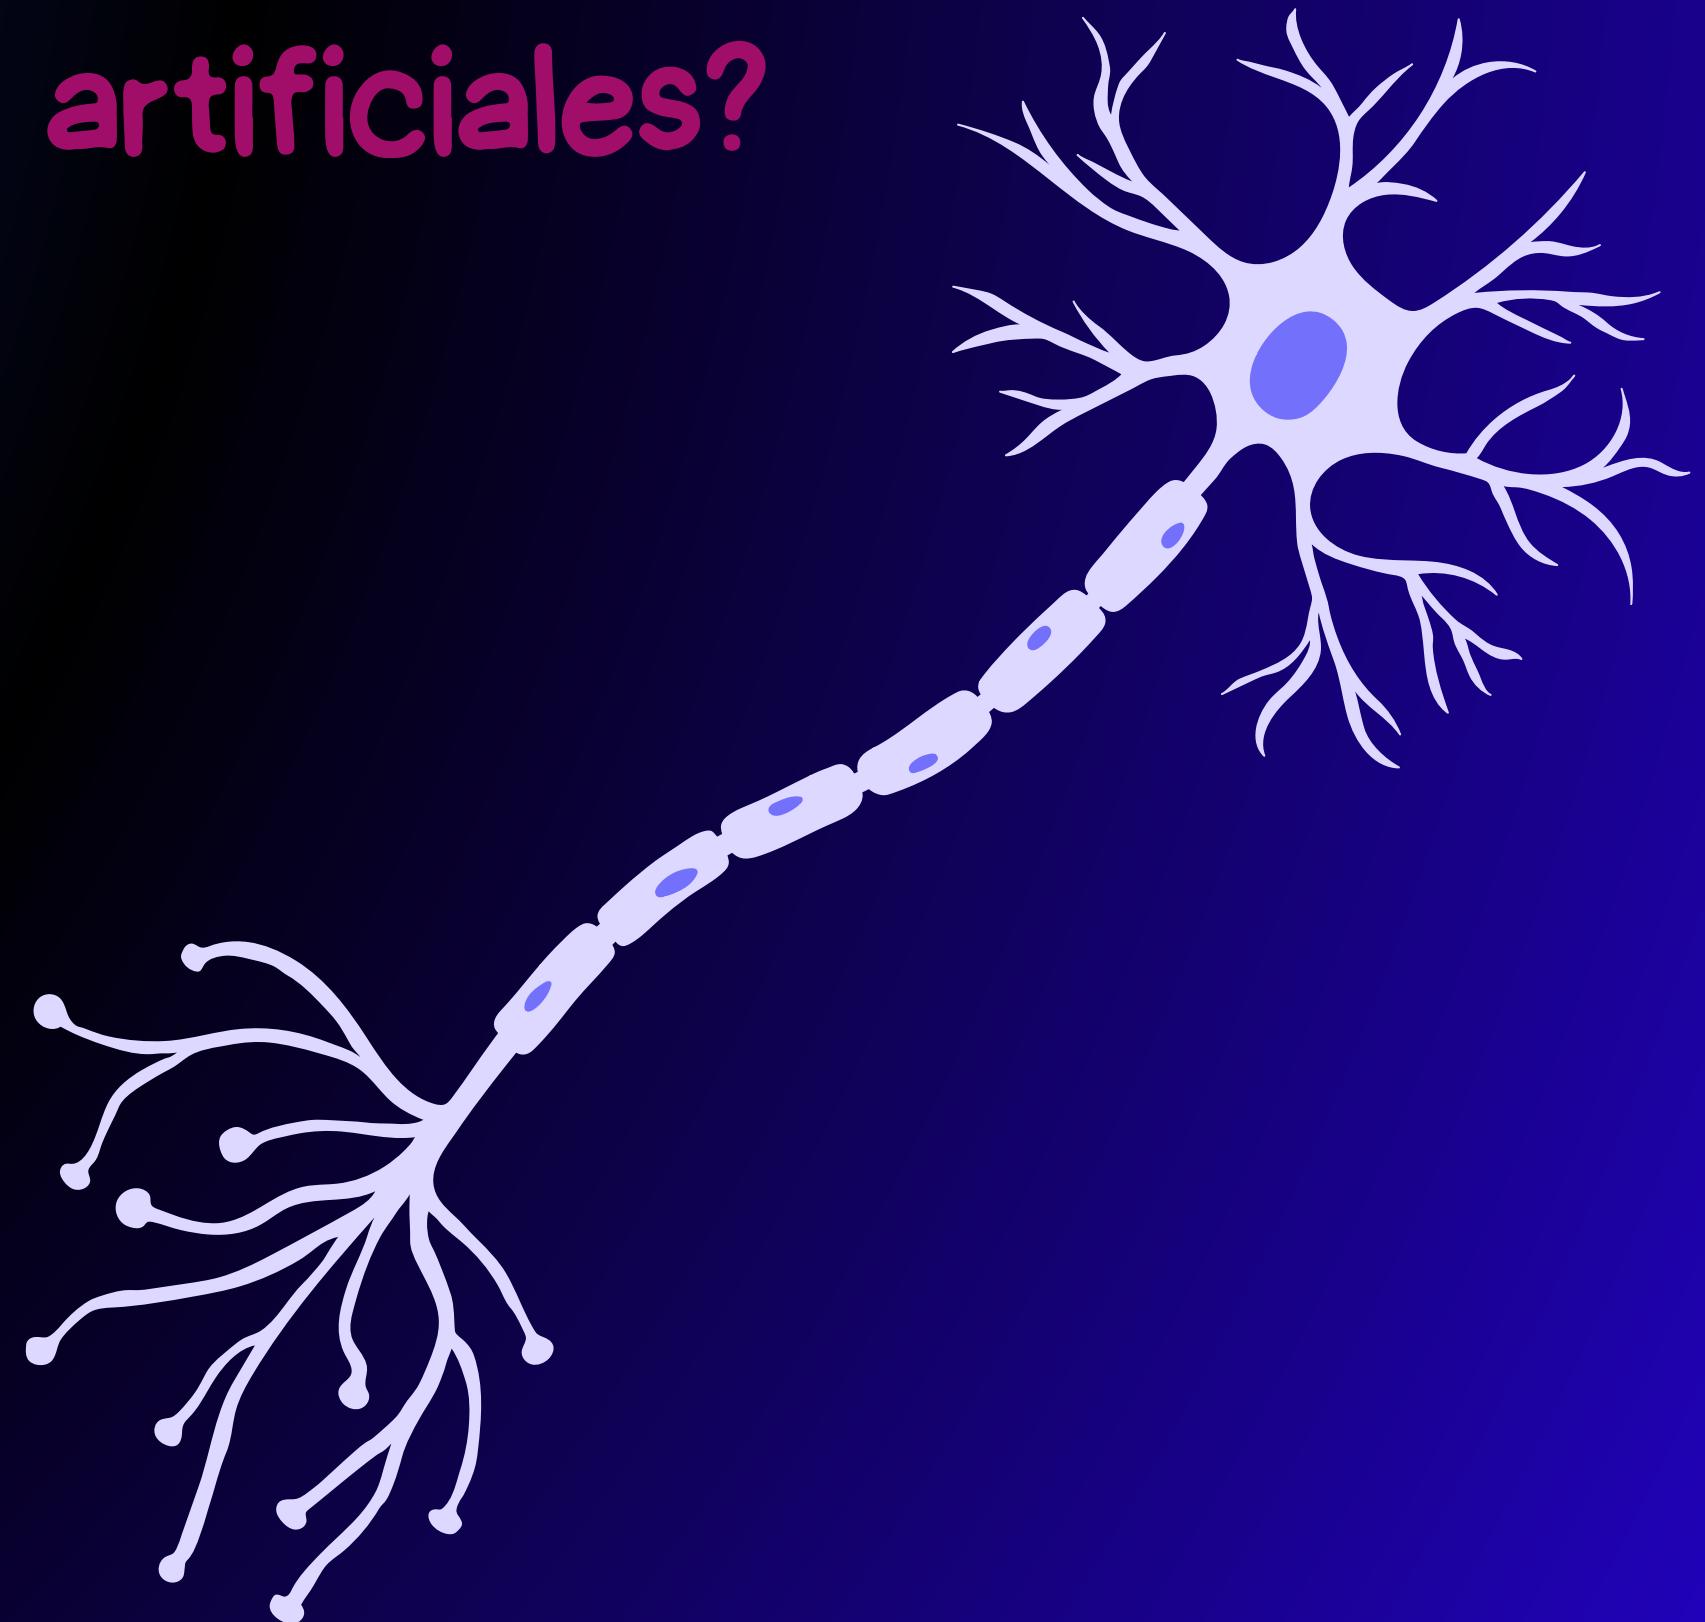
## ESTUDIO DE LA DETECCIÓN

### ATAQUES MALWARE

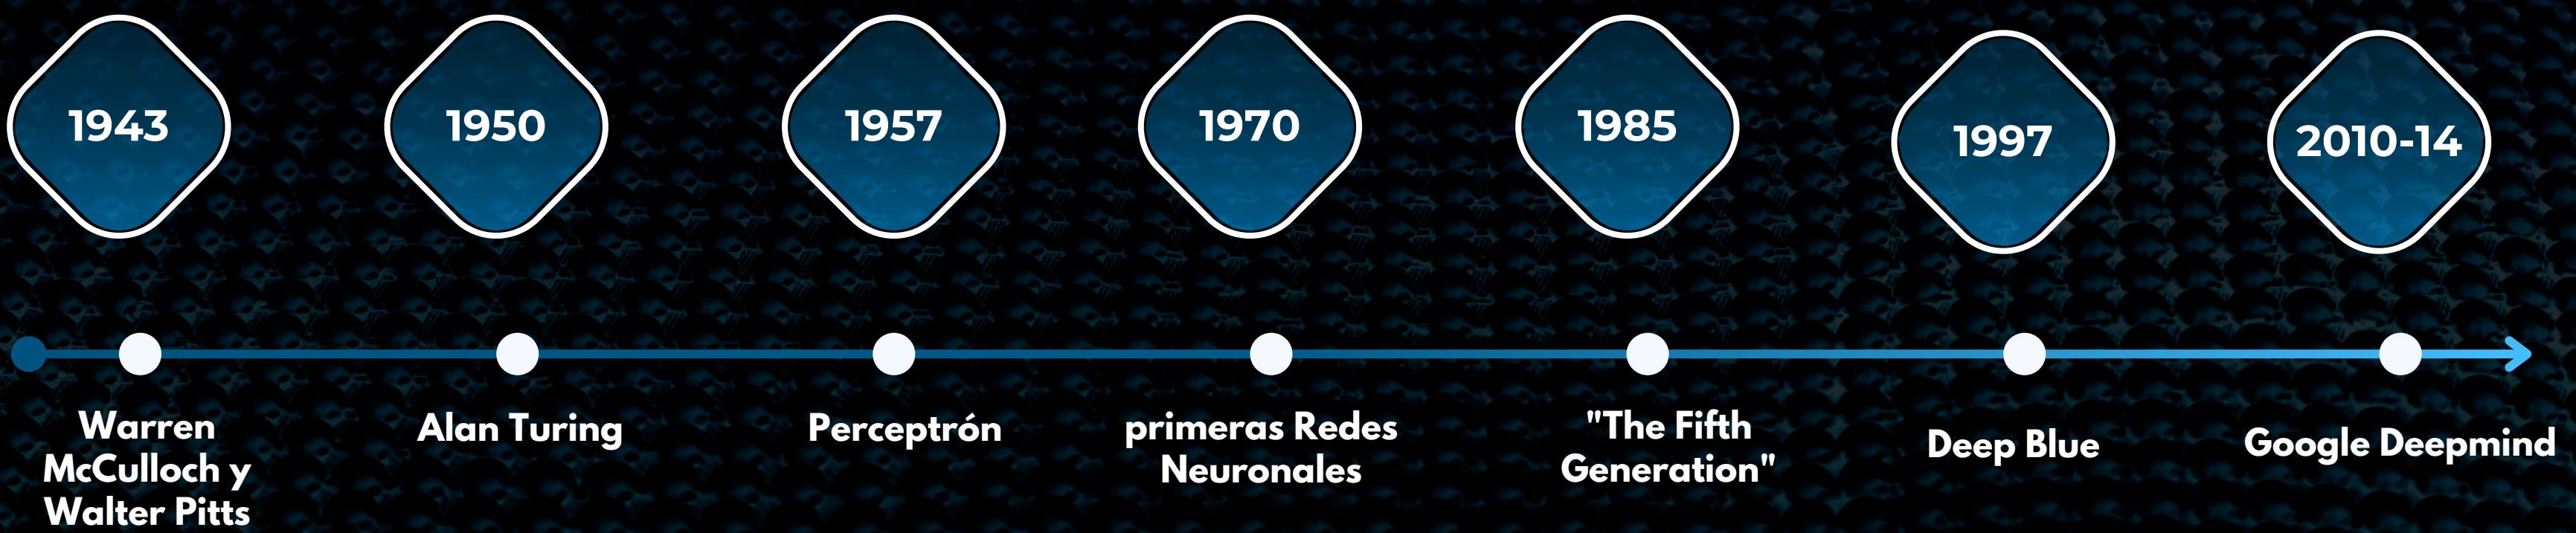


¿Qué son las neuronas?

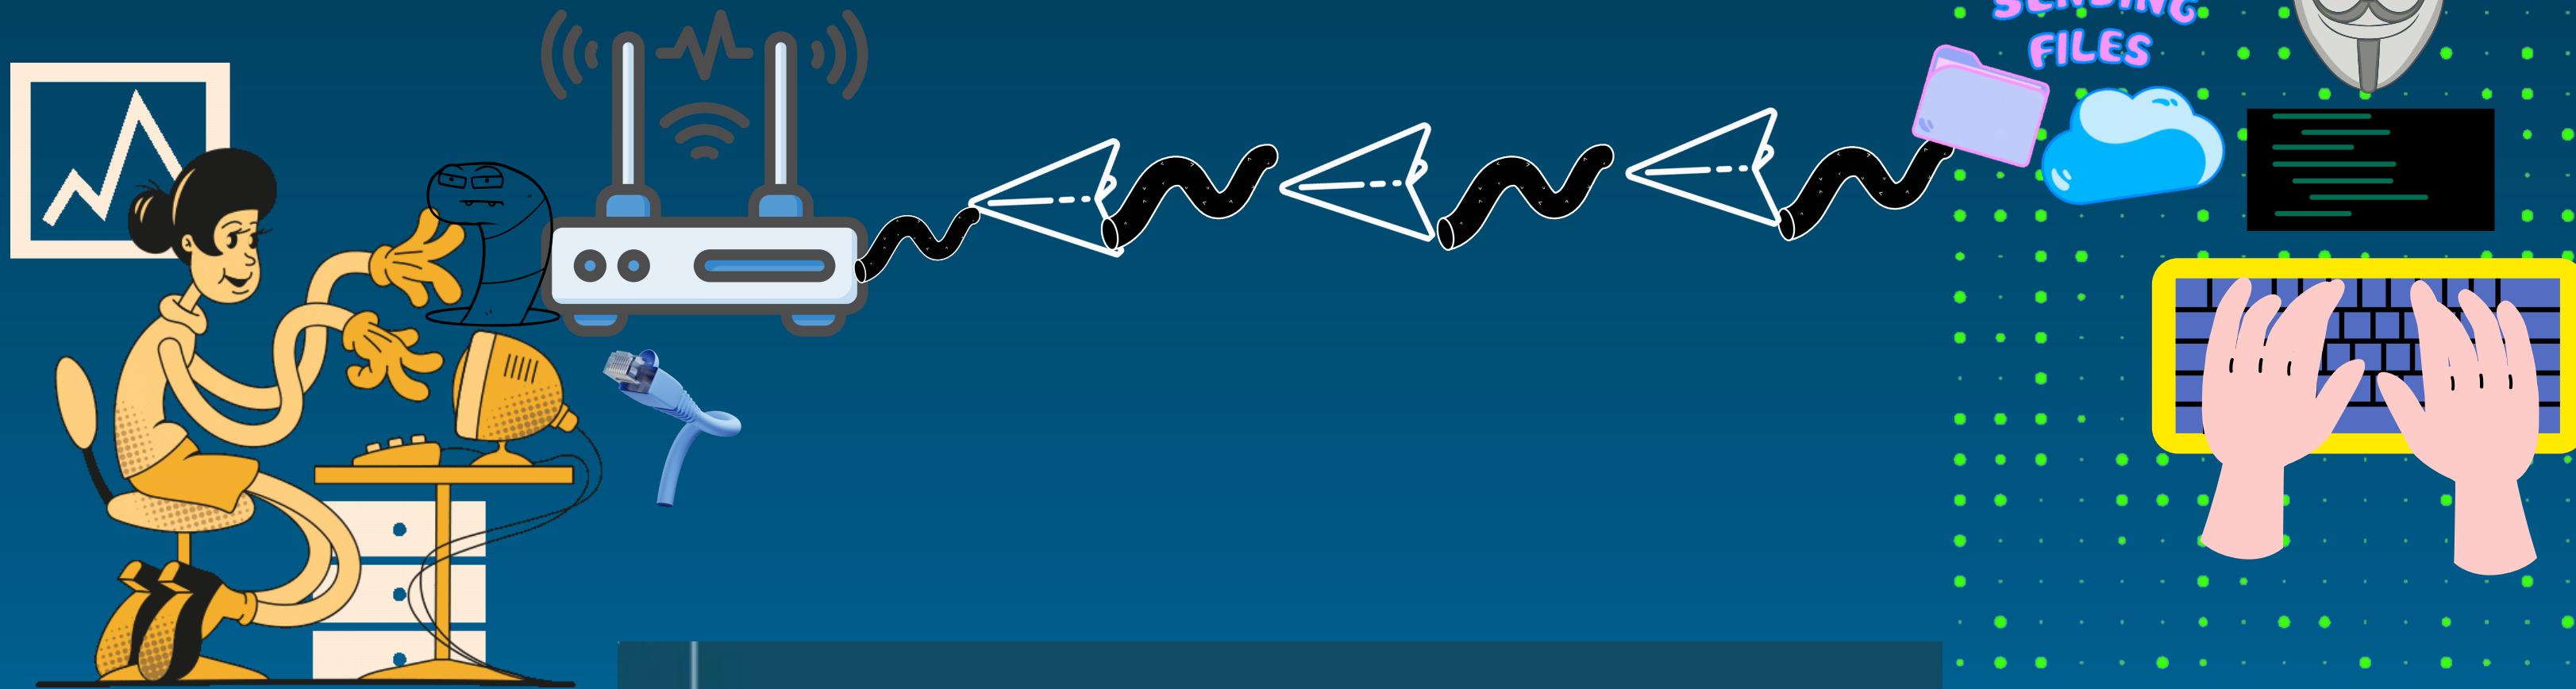
¿Qué son las neuronas artificiales?



# Breve historia de la IA



# ¿Qué es el malware?



# PROCESO NEGOCIO, RECOPIACIÓN, PREPARACIÓN ANÁLISIS Y MODELADO DE DATOS E INFORME RESULTADOS



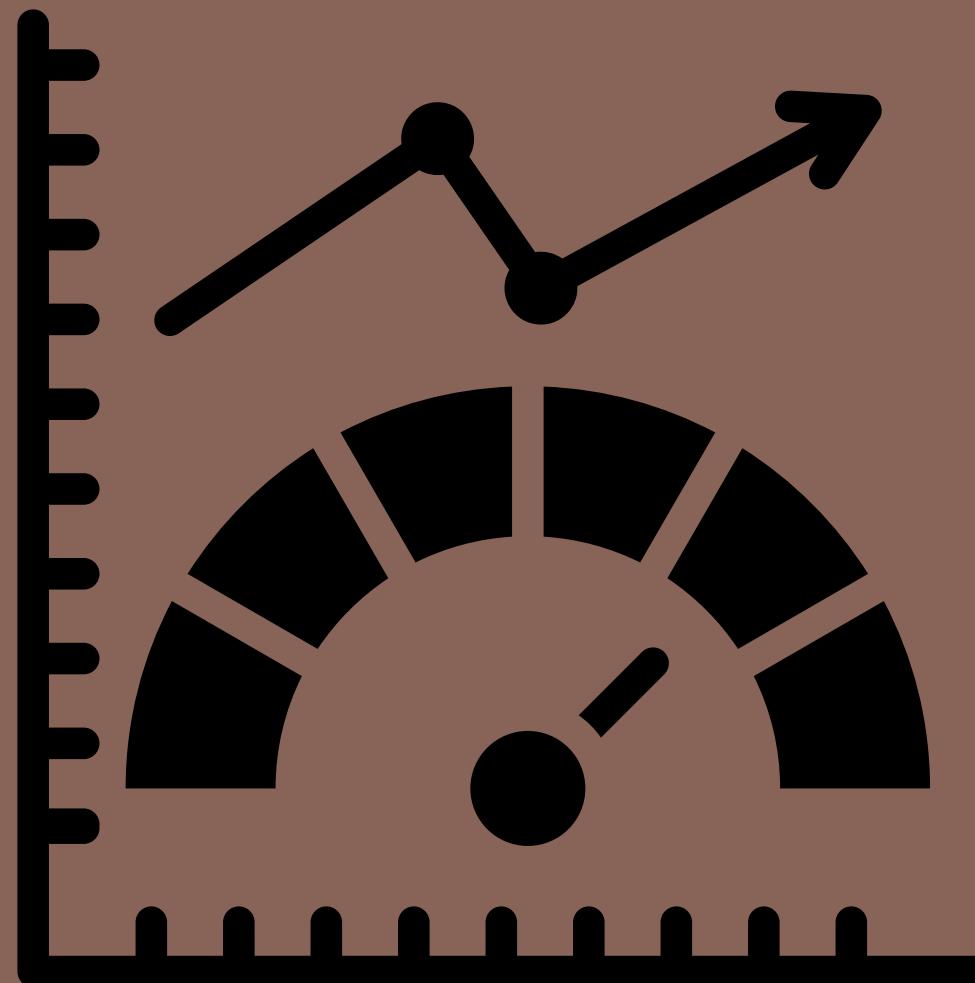


# PROBLEMA DE NEGOCIO

ESTE PROYECTO BUSCA LA EFICIENCIA EN  
LA DETECCION DE MALWARE TENIENDO EN  
CUENTA LOS DATOS OBRANTES EN EL DATASET  
OBJETO DE ESTUDIO, BUSCANDO QUE DETECTE  
TODOS LOS ATAQUES DE ESTE TIPO



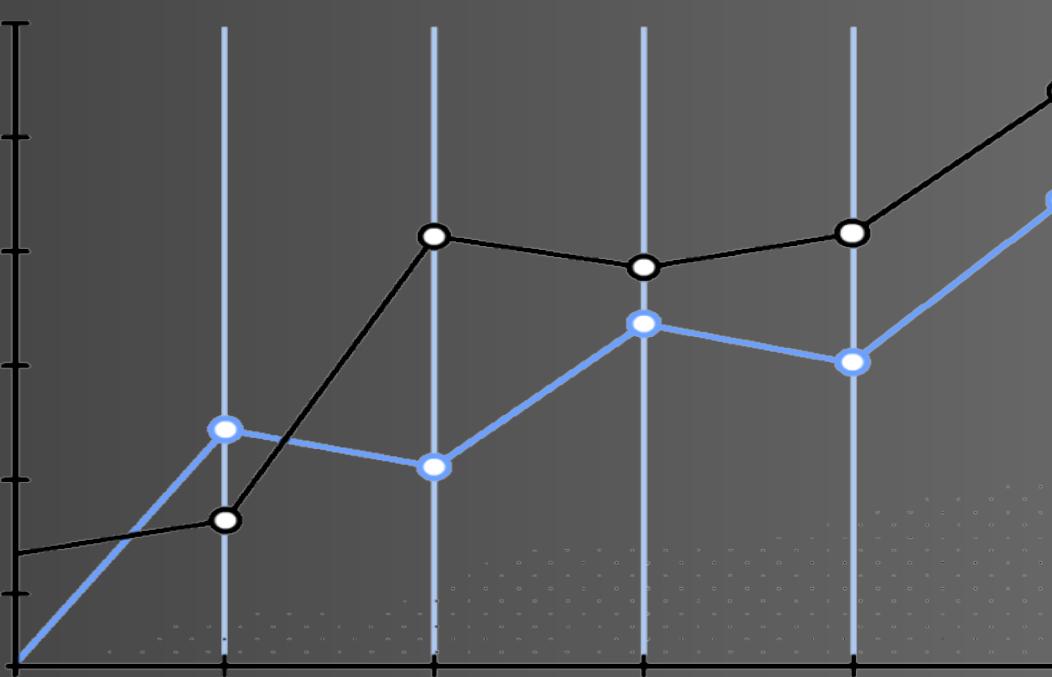
METRICA= RECALL



# EL DATASET

Columna:	Descripción:	Tipo:	Orden de entrada:	Información posterior a la etiqueta "label"
id.resp_h	Identificador único del host de destino	Categórico	1	No
proto	Protocolo utilizado (TCP, UDP, ICMP, etc.)	Categórico	2	No
conn_state	Estado de la conexión (SYN, ACK, FIN, etc.)	Categórico	3	No
orig_pkts	Número de paquetes enviados por el host de origen	Numérico	4	No
orig_ip_bytes	Número de bytes enviados por el host de origen	Numérico	5	No
resp_pkts	Número de paquetes recibidos por el host de destino	Numérico	6	No
resp_ip_bytes	Número de bytes recibidos por el host de destino	Numérico	7	No
anyo	Año de la conexión	Numérico	8	No
mes	Mes de la conexión	Numérico	9	No
dia	Día de la conexión	Numérico	10	No
hora	Hora de la conexión	Numérico	11	No
id.resp_p	Identificador único del puerto de destino	Categórico	12	No
missed_bytes	Número de bytes perdidos en la conexión	Numérico	13	No
id.orig_p	Identificador único del puerto de origen	Categórico	14	No
hostory	conexiones y patrones detectados de conexiones anteriores	Categórico	15	No
label	Etiqueta que indica si la conexión es normal o anómala	Categórica	Target	

# MINI- EDA



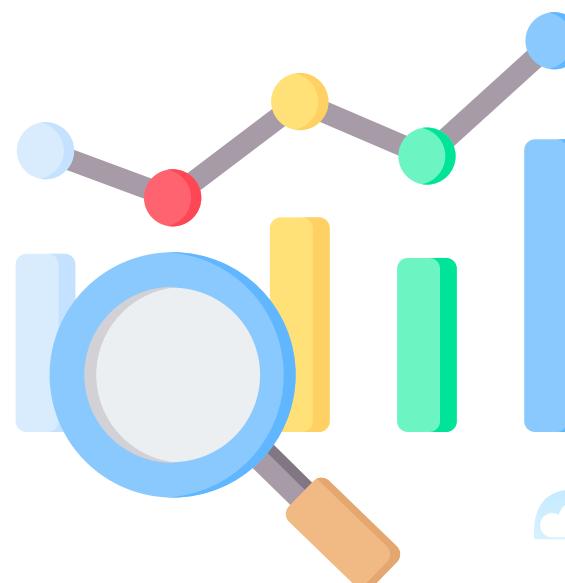


# TARTAMIENTO DE NULOS

## NULOS

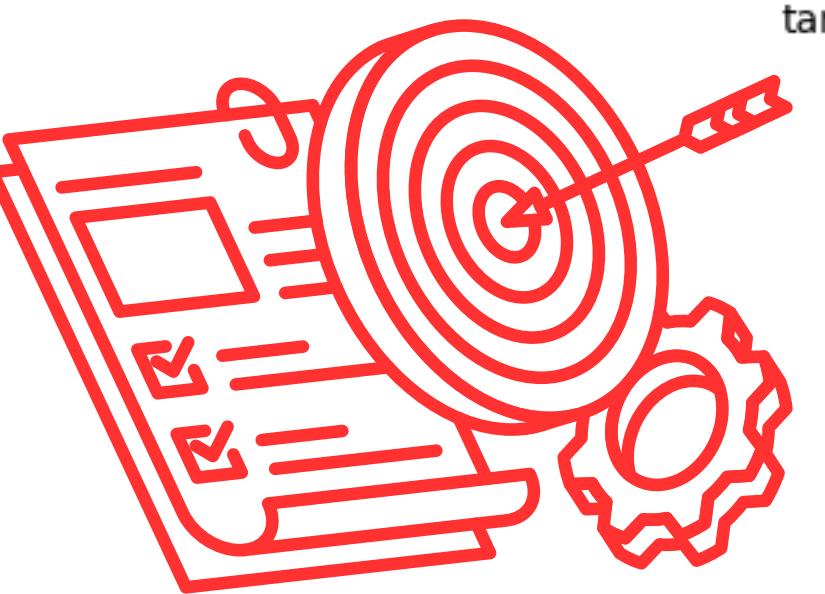
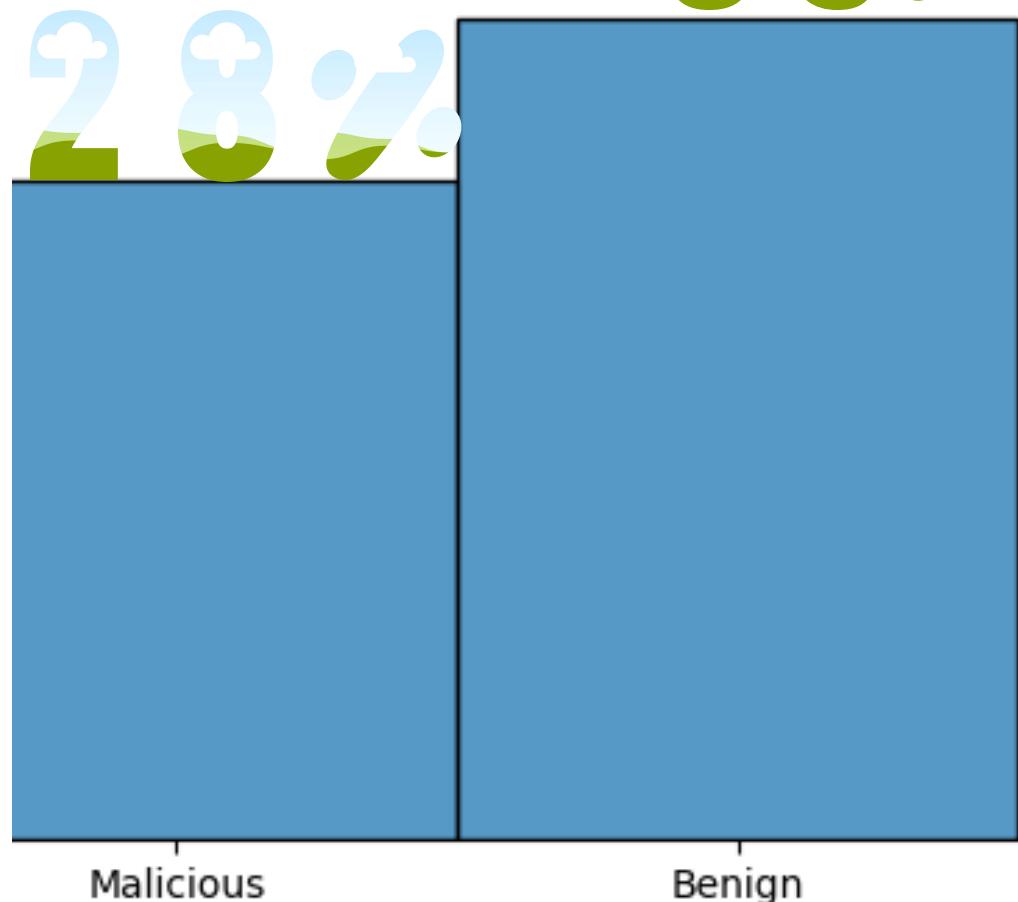
ts	0
uid	0
id.orig_h	0
id.orig_p	0
id.resp_h	0
id.resp_p	0
proto	0
service	24993006
duration	15272073
orig_bytes	15272073
resp_bytes	15272073
conn_state	0
local_orig	25011603
local_resp	25011603
missed_bytes	0
history	25116
orig_pkts	0
orig_ip_bytes	0
resp_pkts	0
resp_ip_bytes	0
tunnel_parents	25011603
label	0
detailed-label	17954112

# TARGET

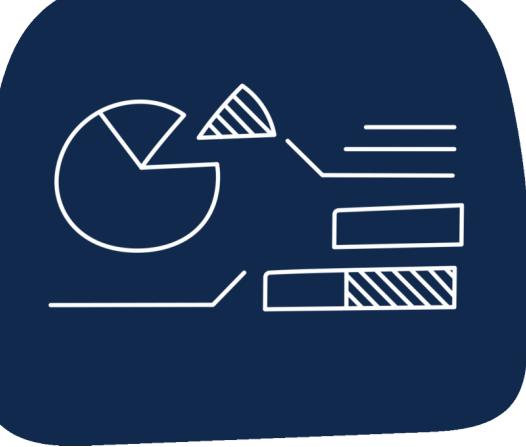
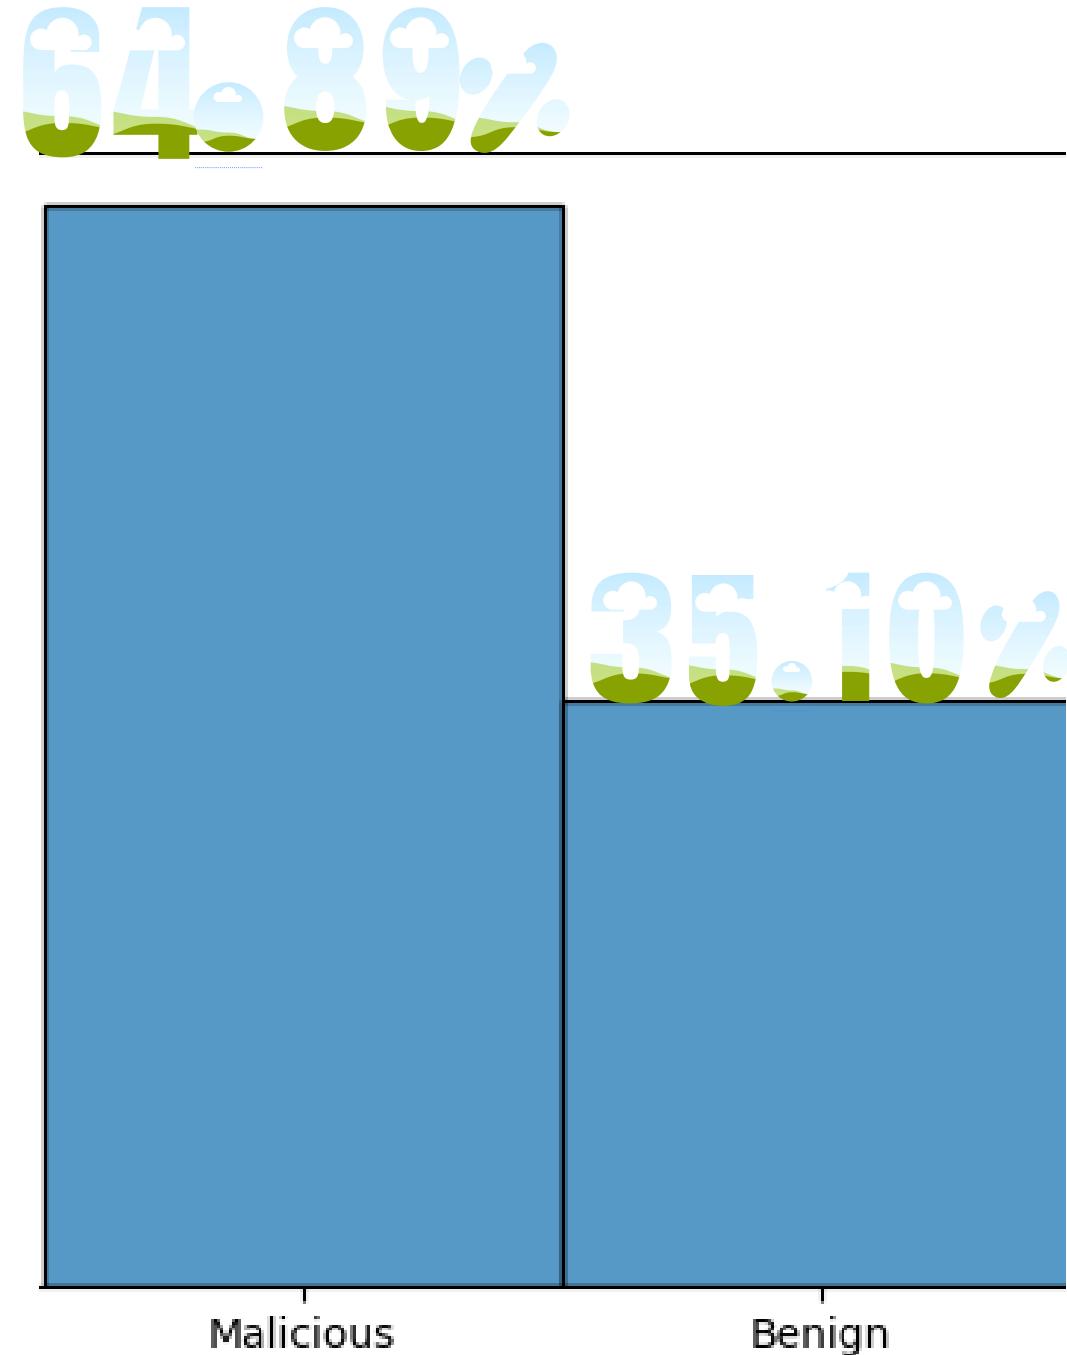


35%

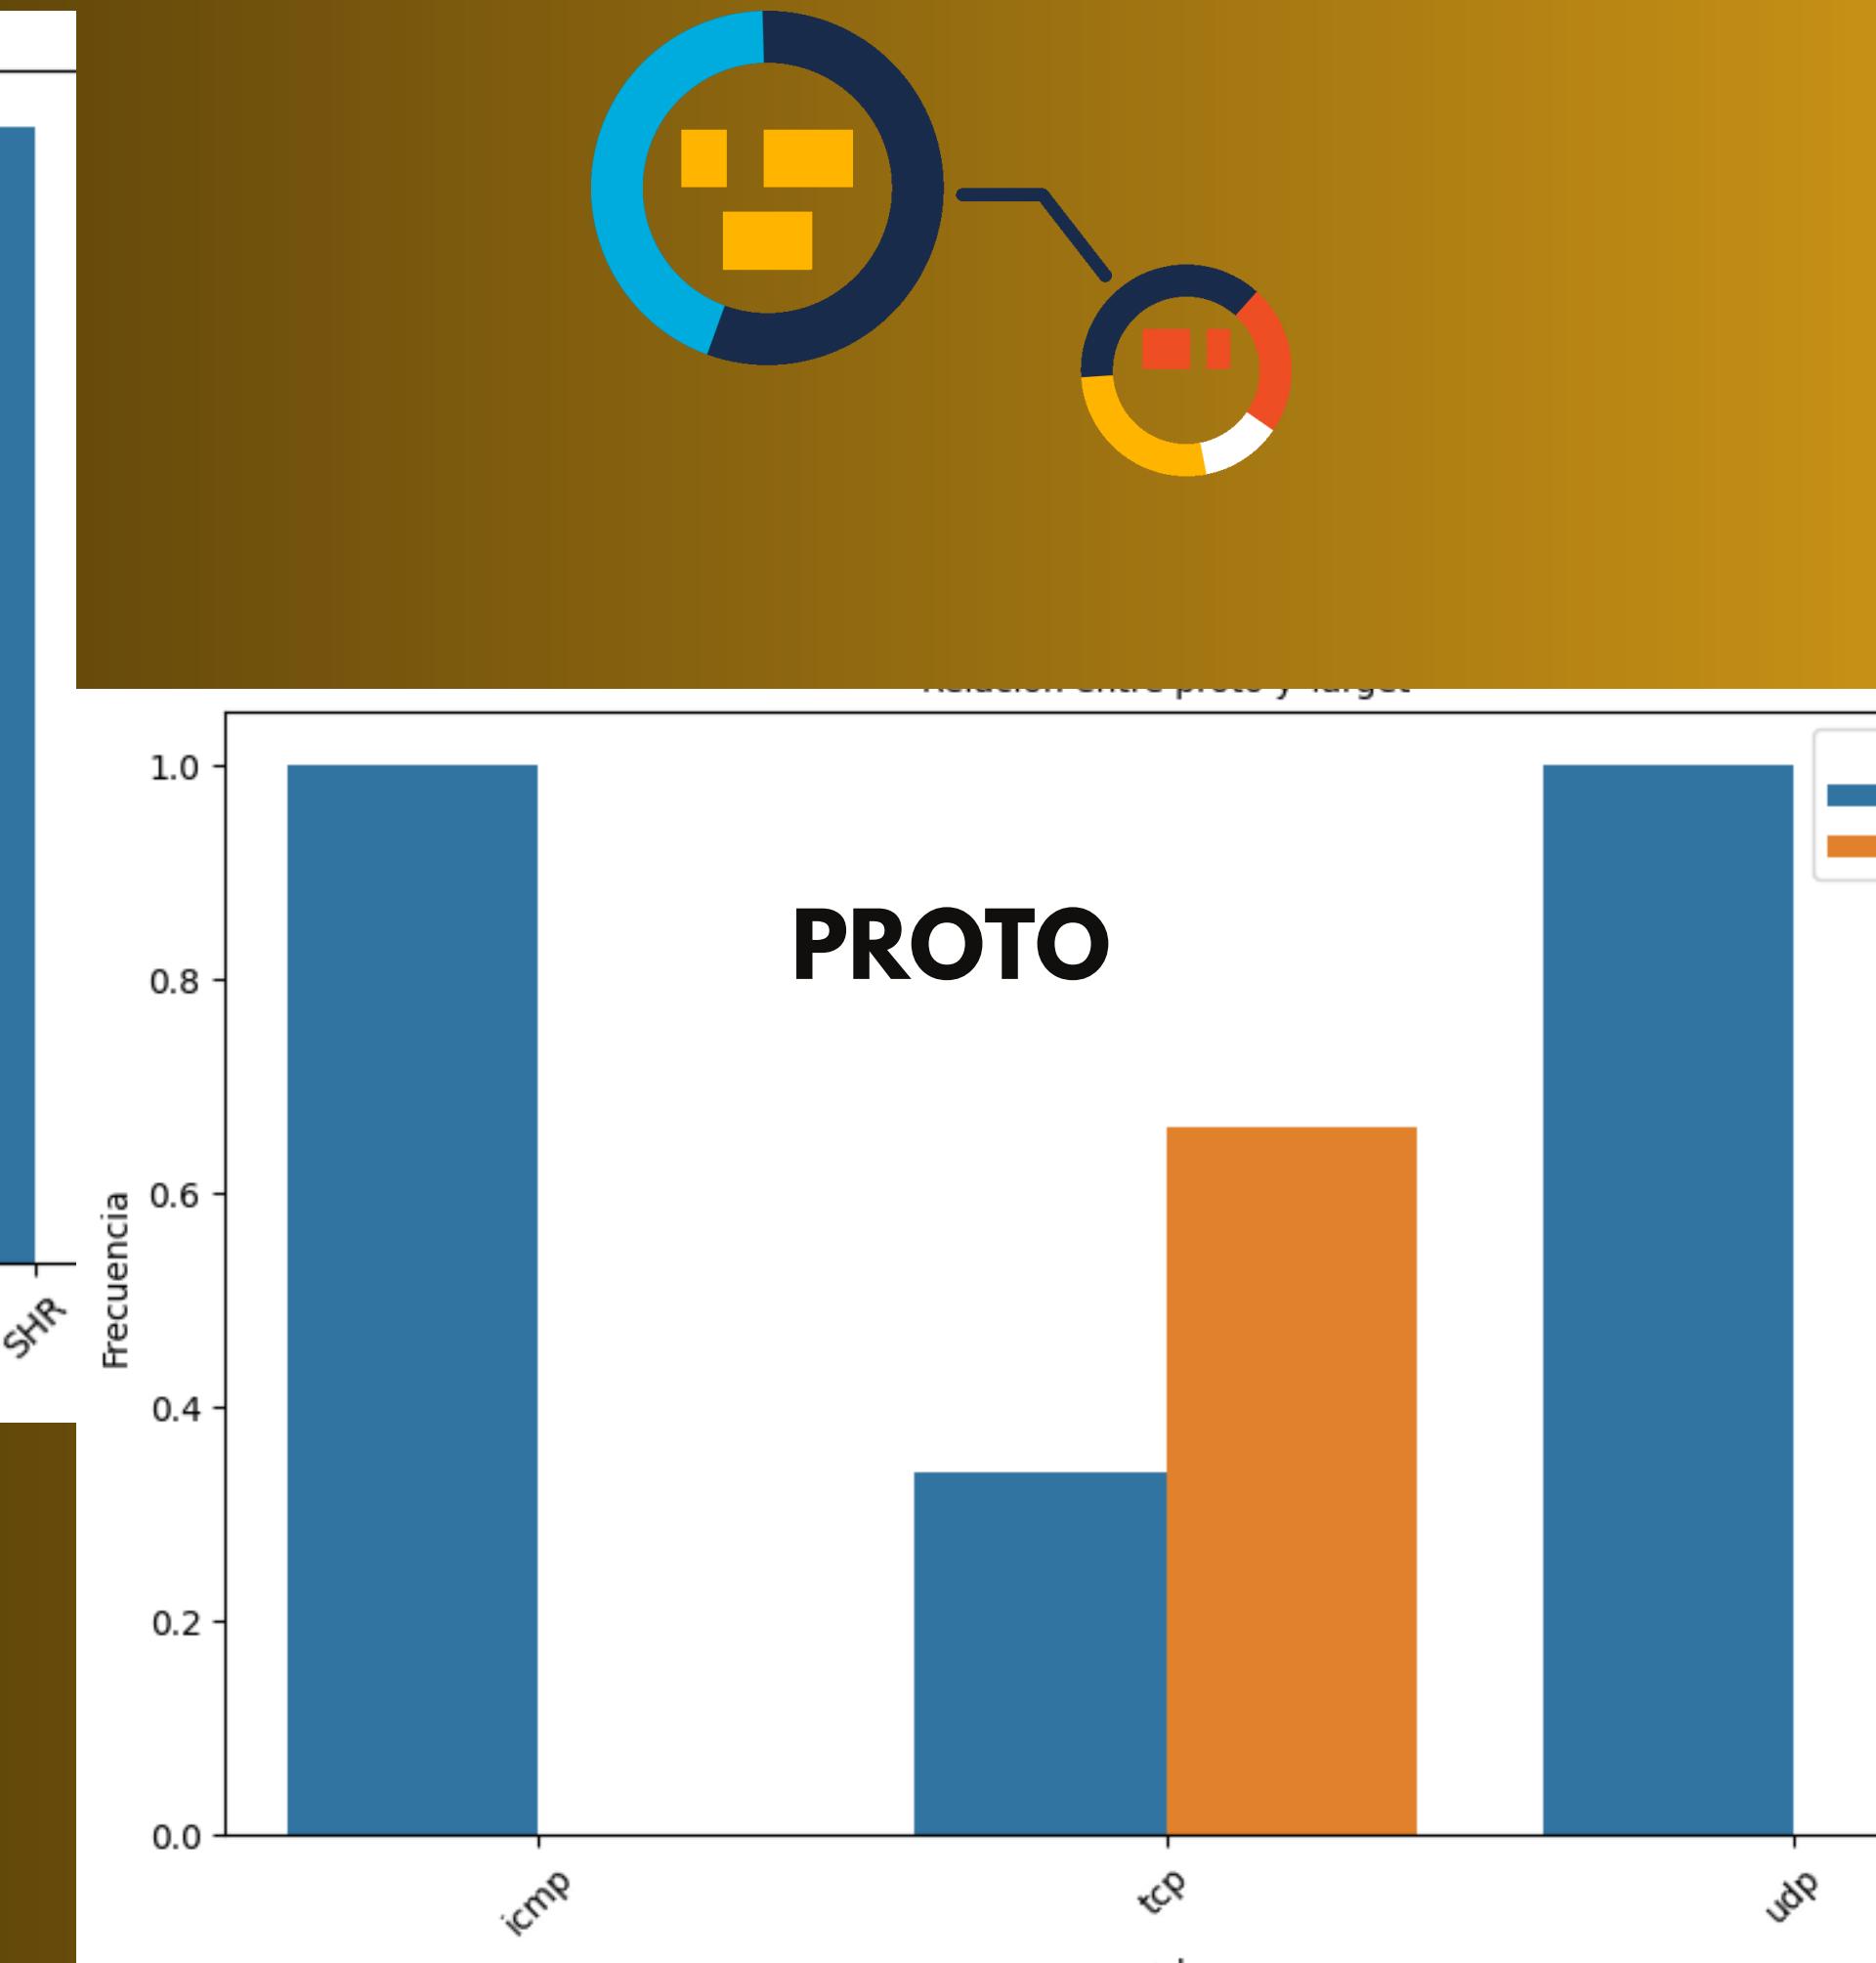
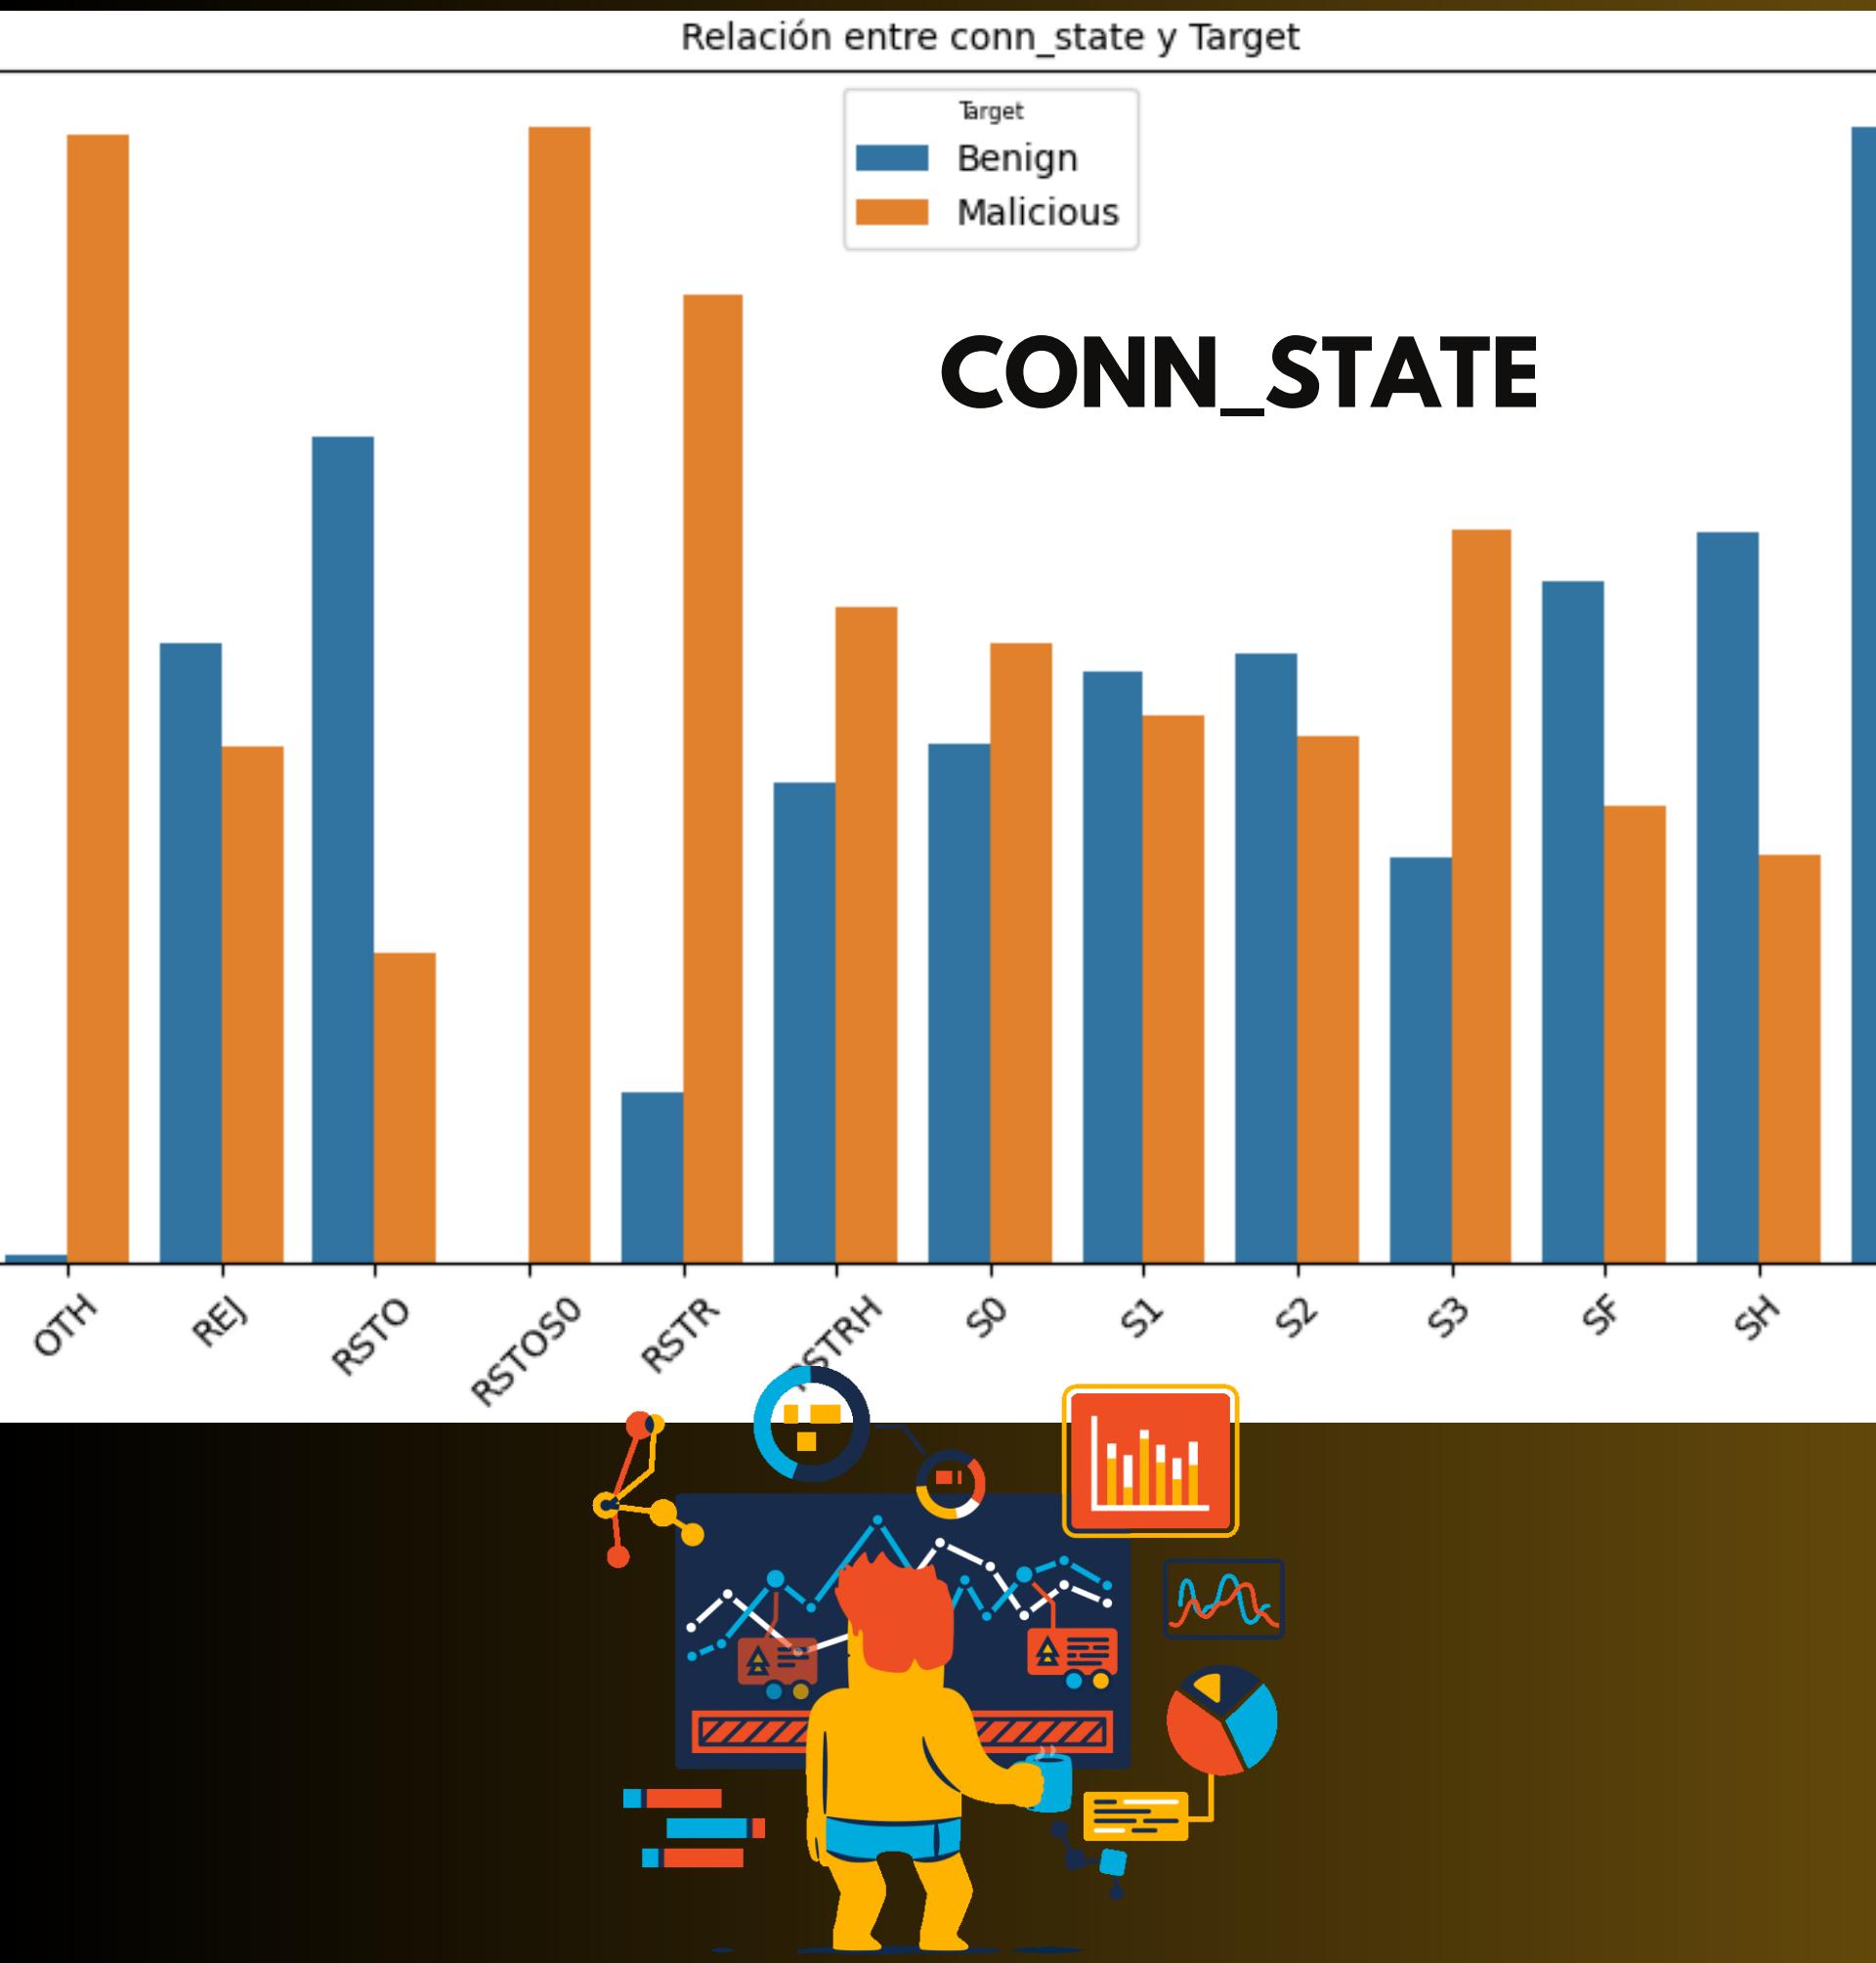
SIN TRATAR



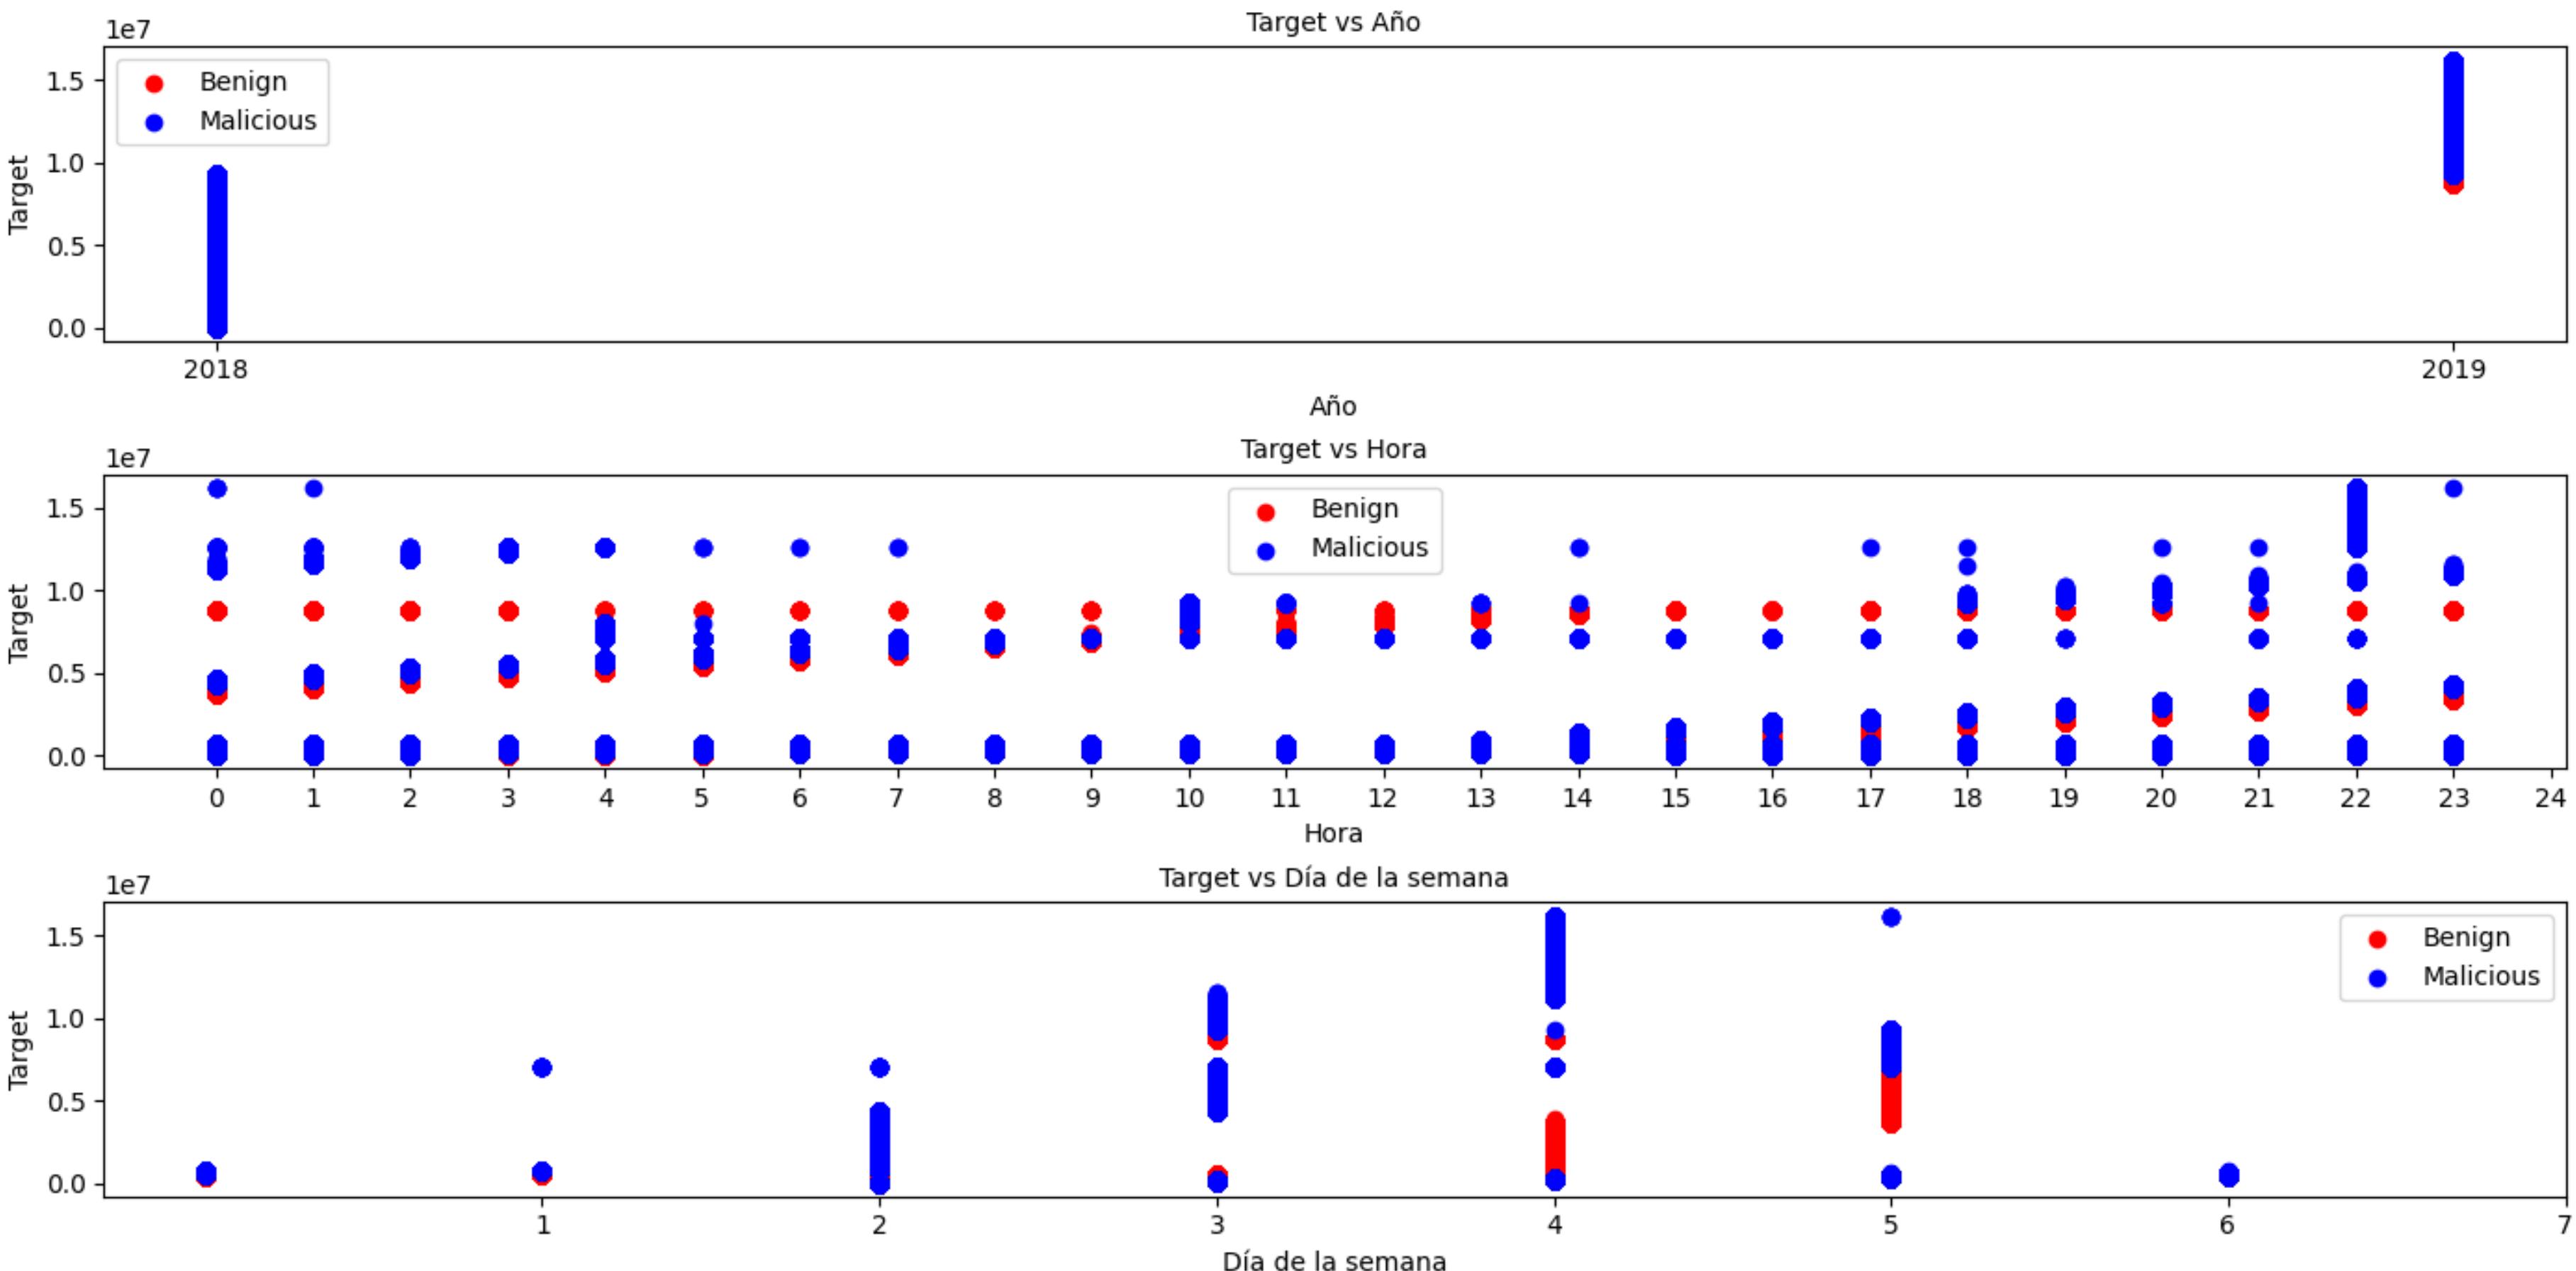
Malicious C&C Malicious PartOfAHorizontalPortScan Malicious DDoS  
target



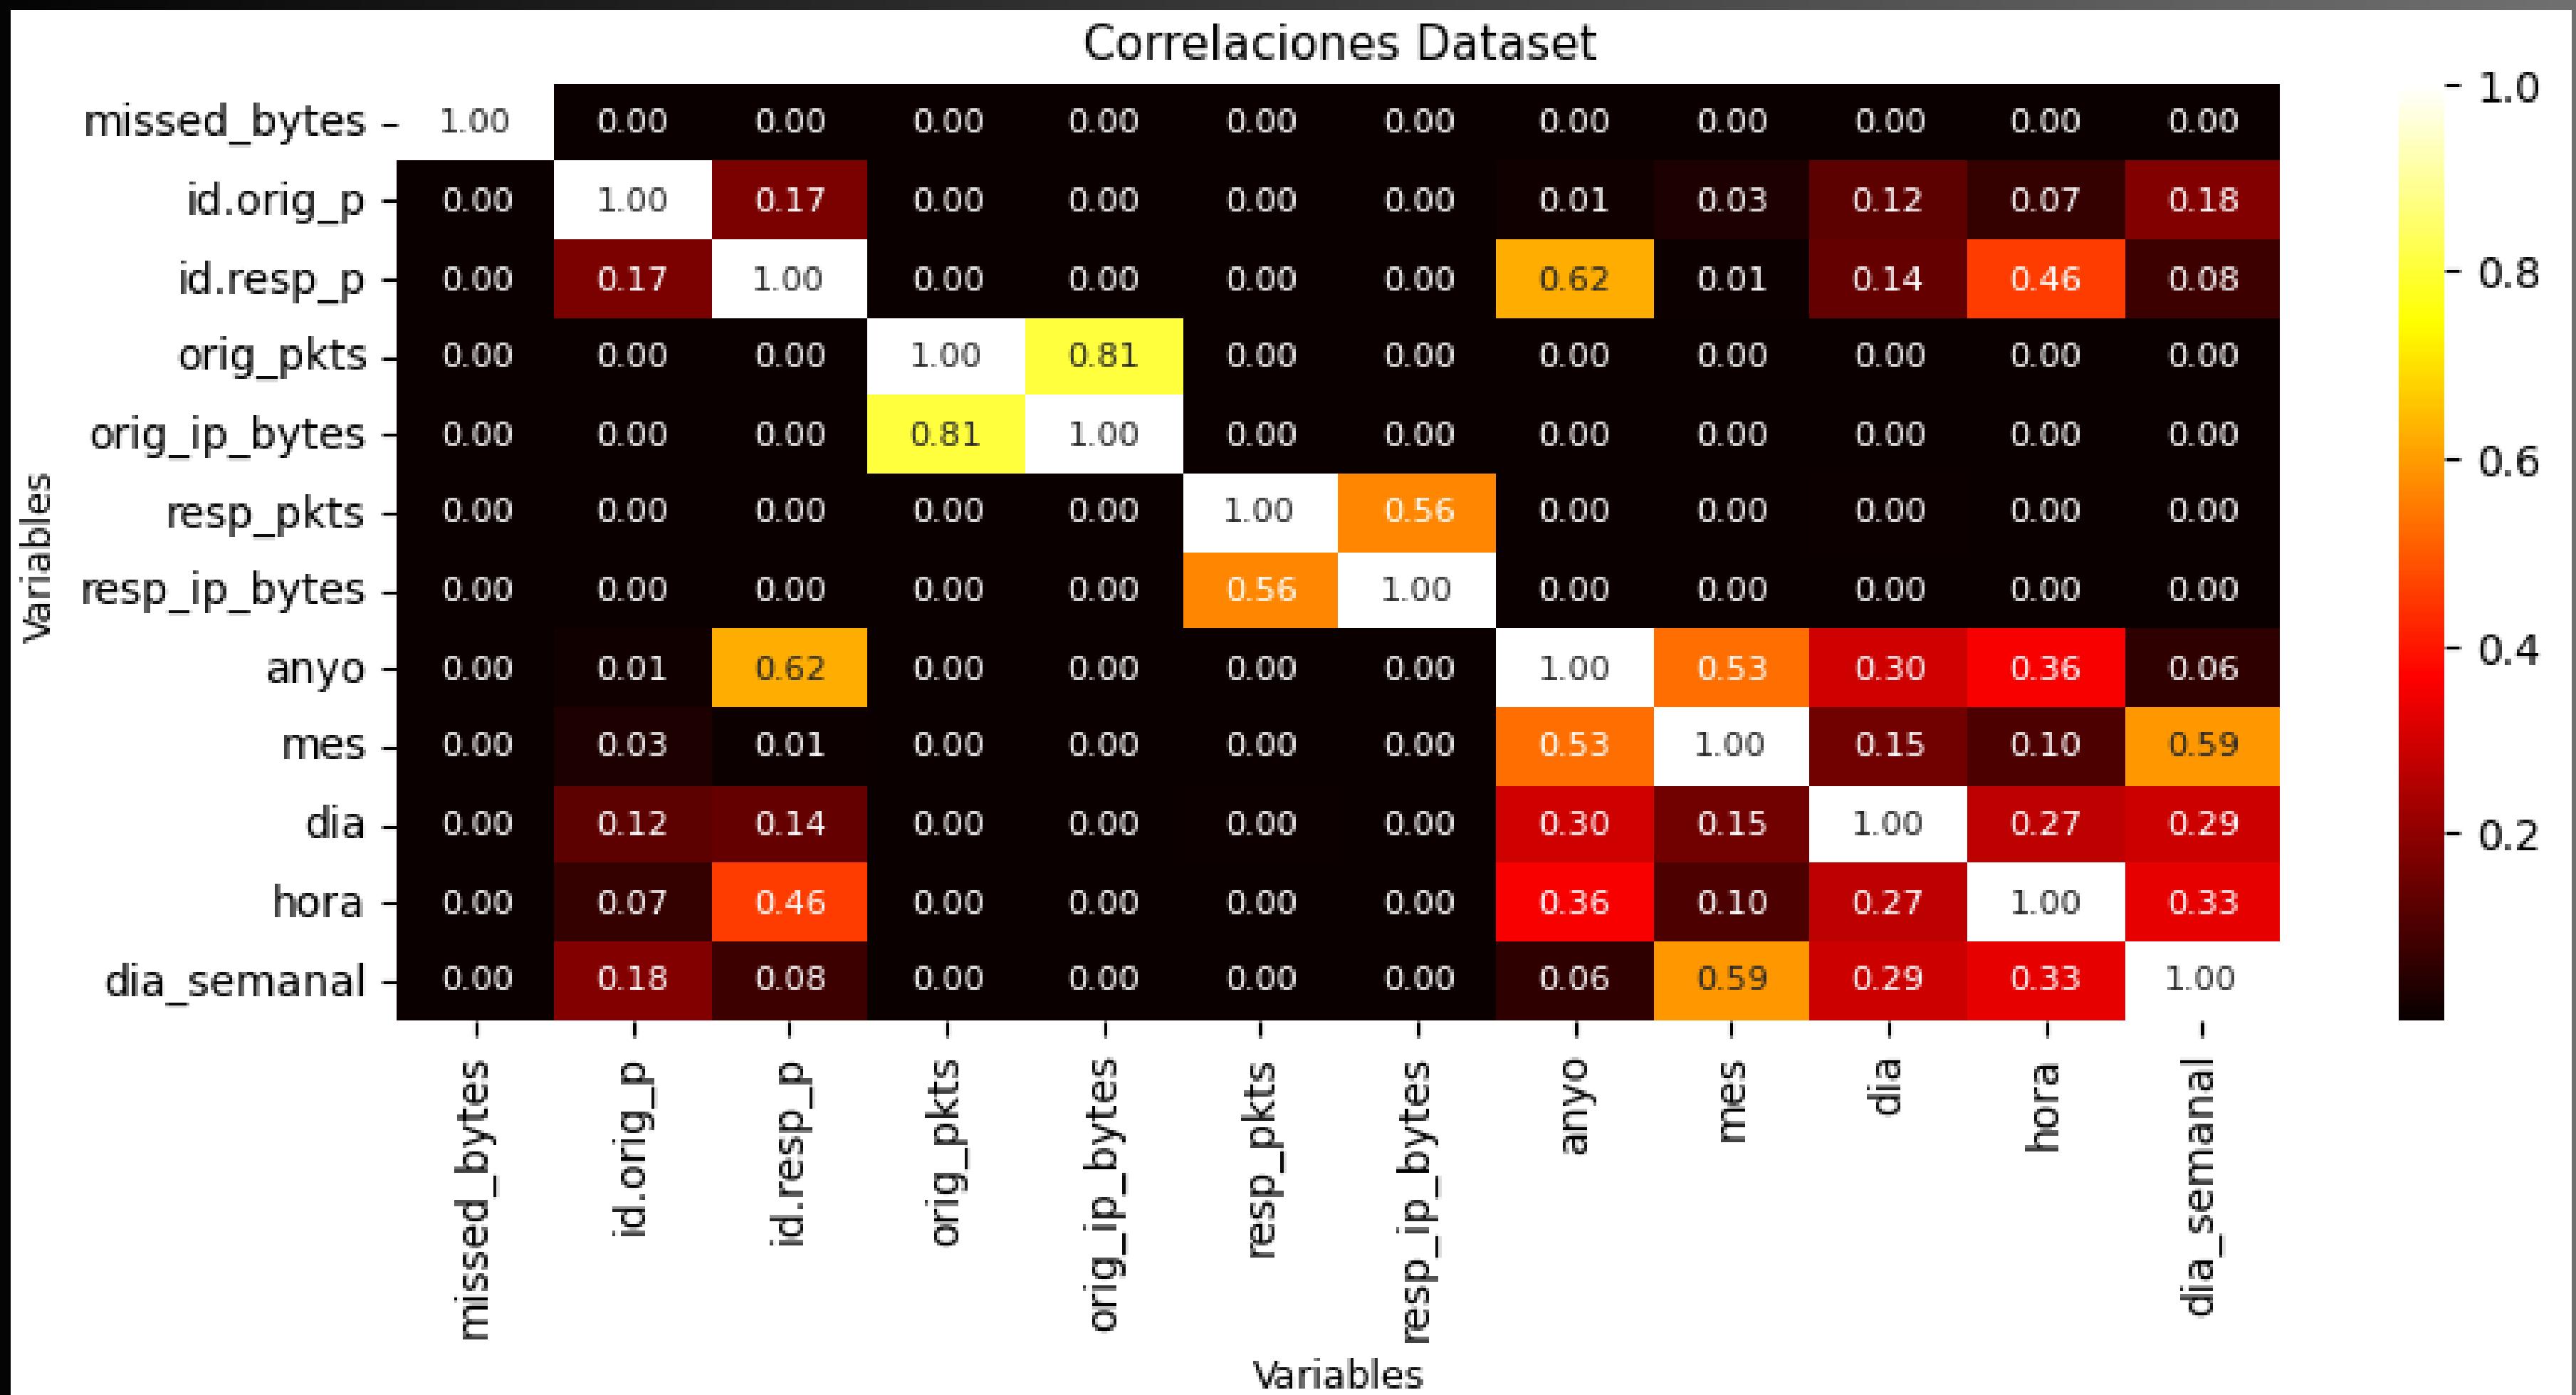
# RELACIONES TARGET CON COLUMNAS `CONN_STATE` Y `PROTO`

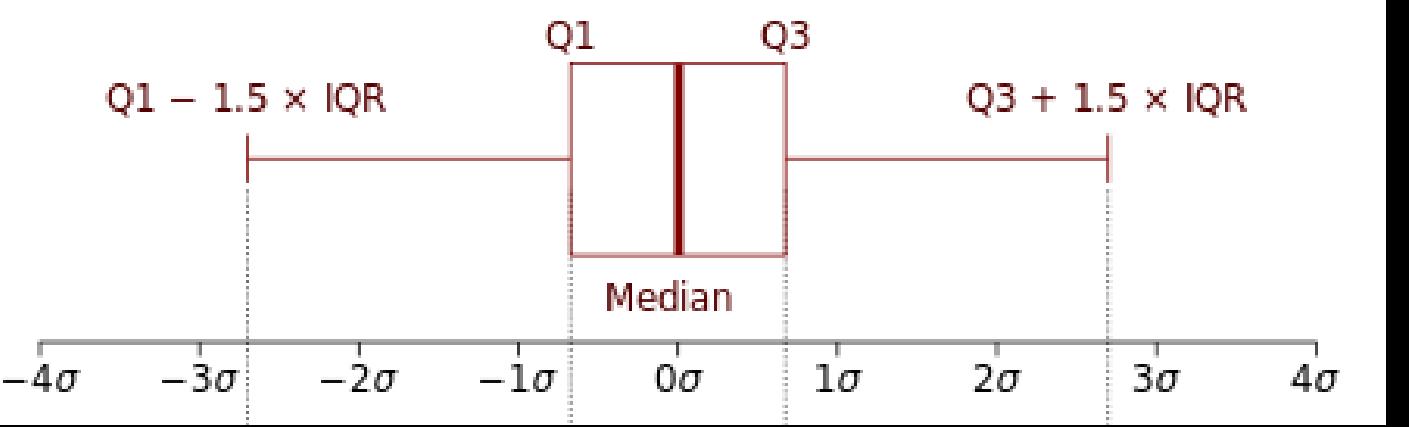


# INTERVALOS DE ATAQUES POR HORA, AÑO Y DIA SEMANAL



# ESTUDIO DE LA CORRELACION



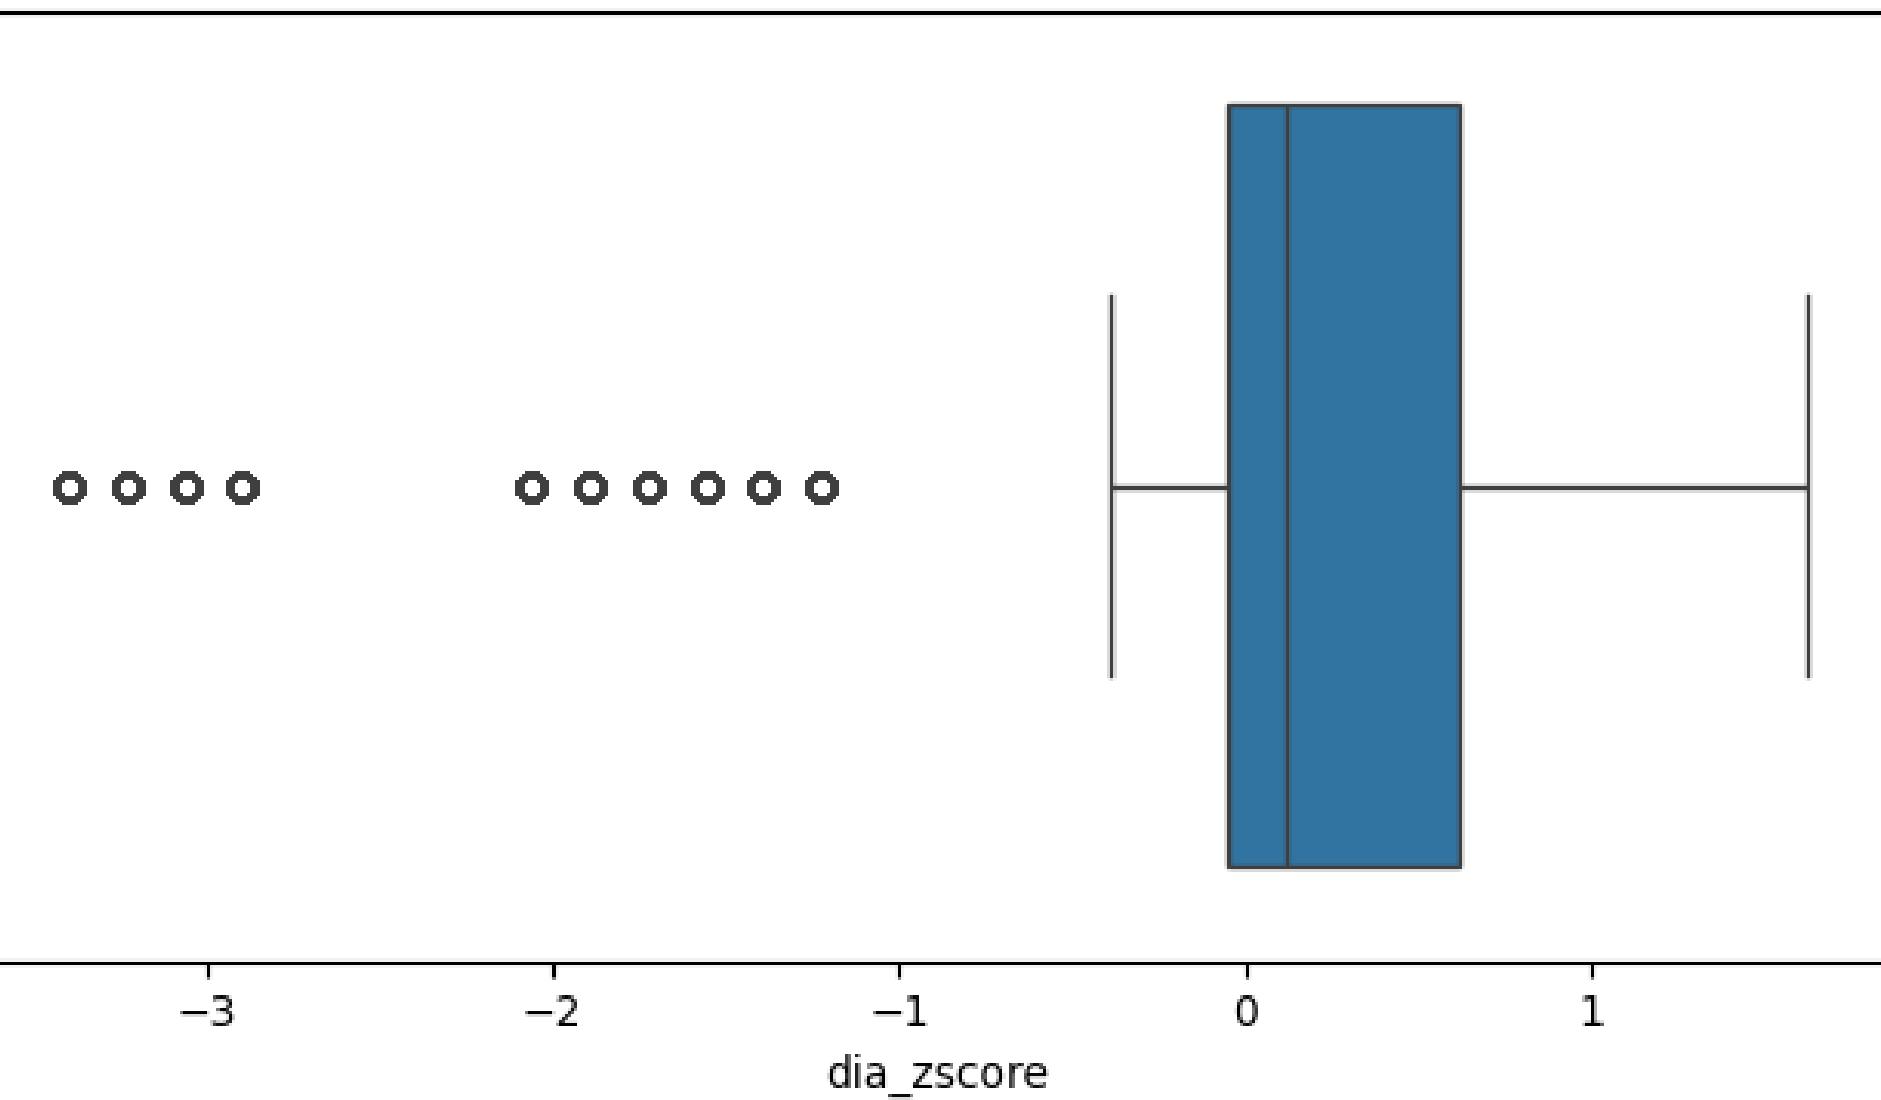


# BUSQUEDA Y ELIMINACION DE OUTLIERS



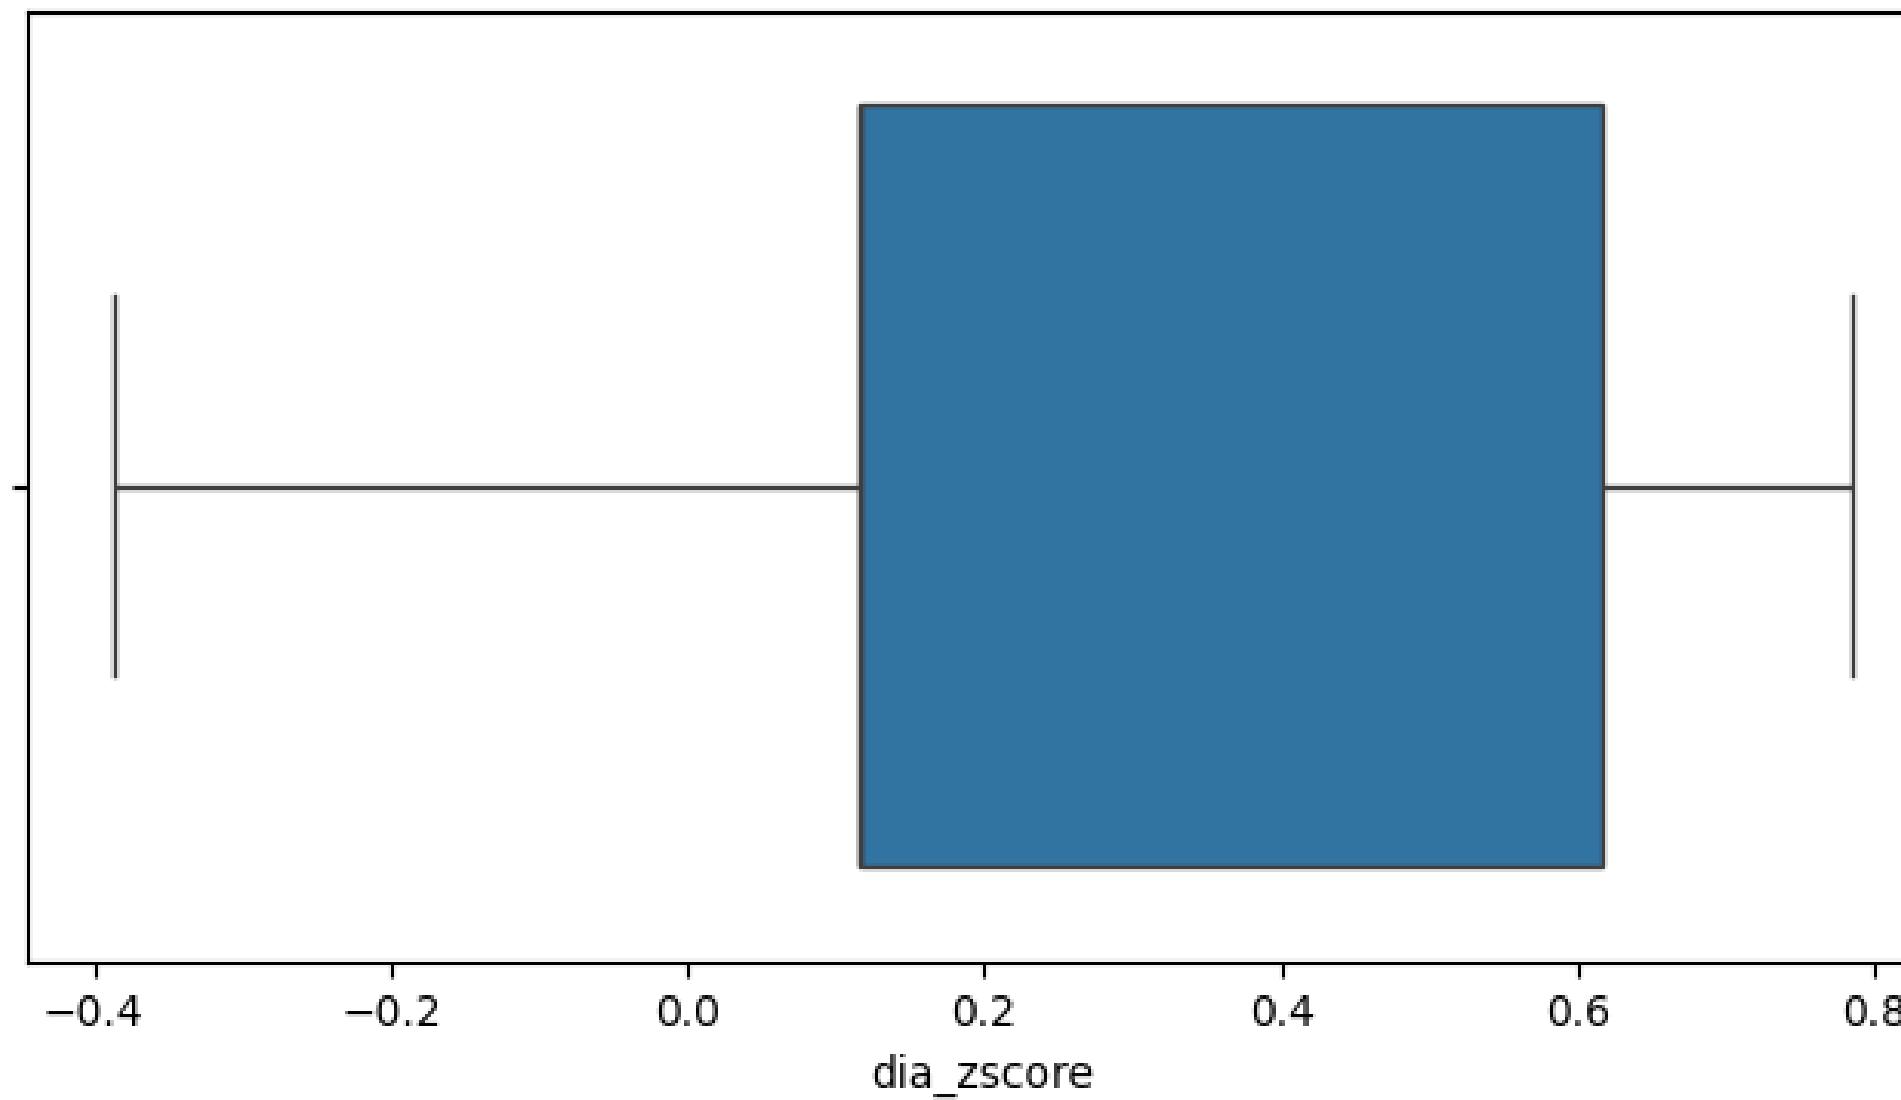
## ANTES DE ELIMINACION DE OUTLIERS

Boxplot de dia\_zscore



## TRAS LA ELIMINACION DE OUTLIERS

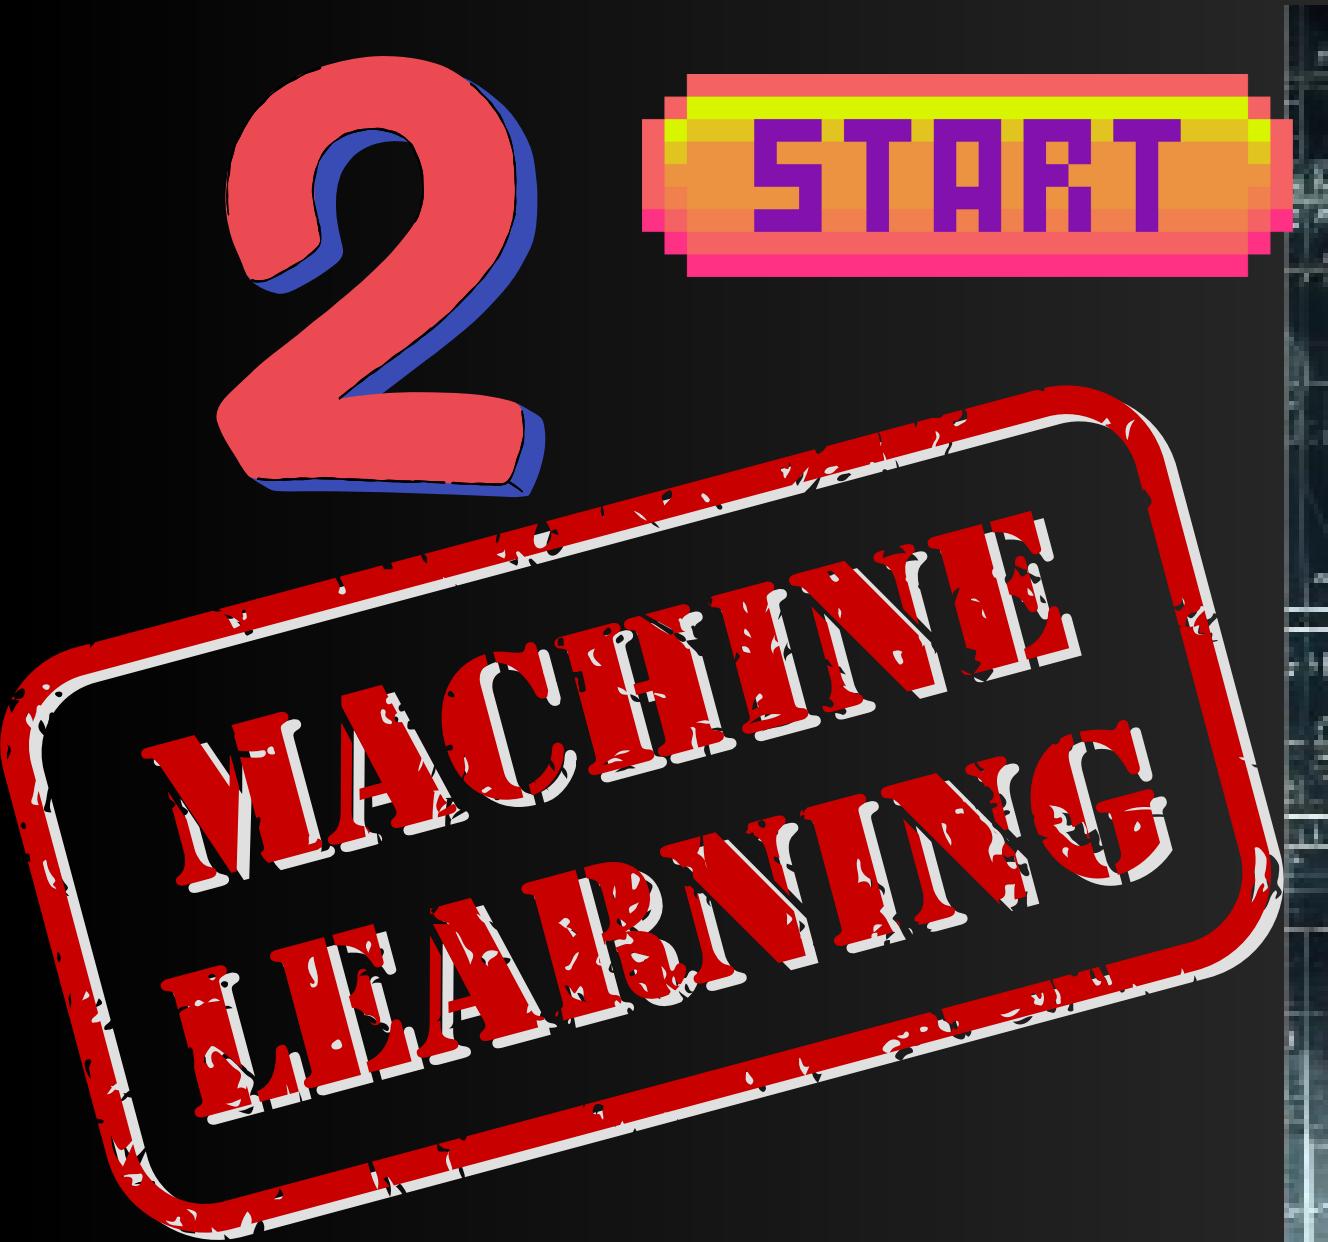
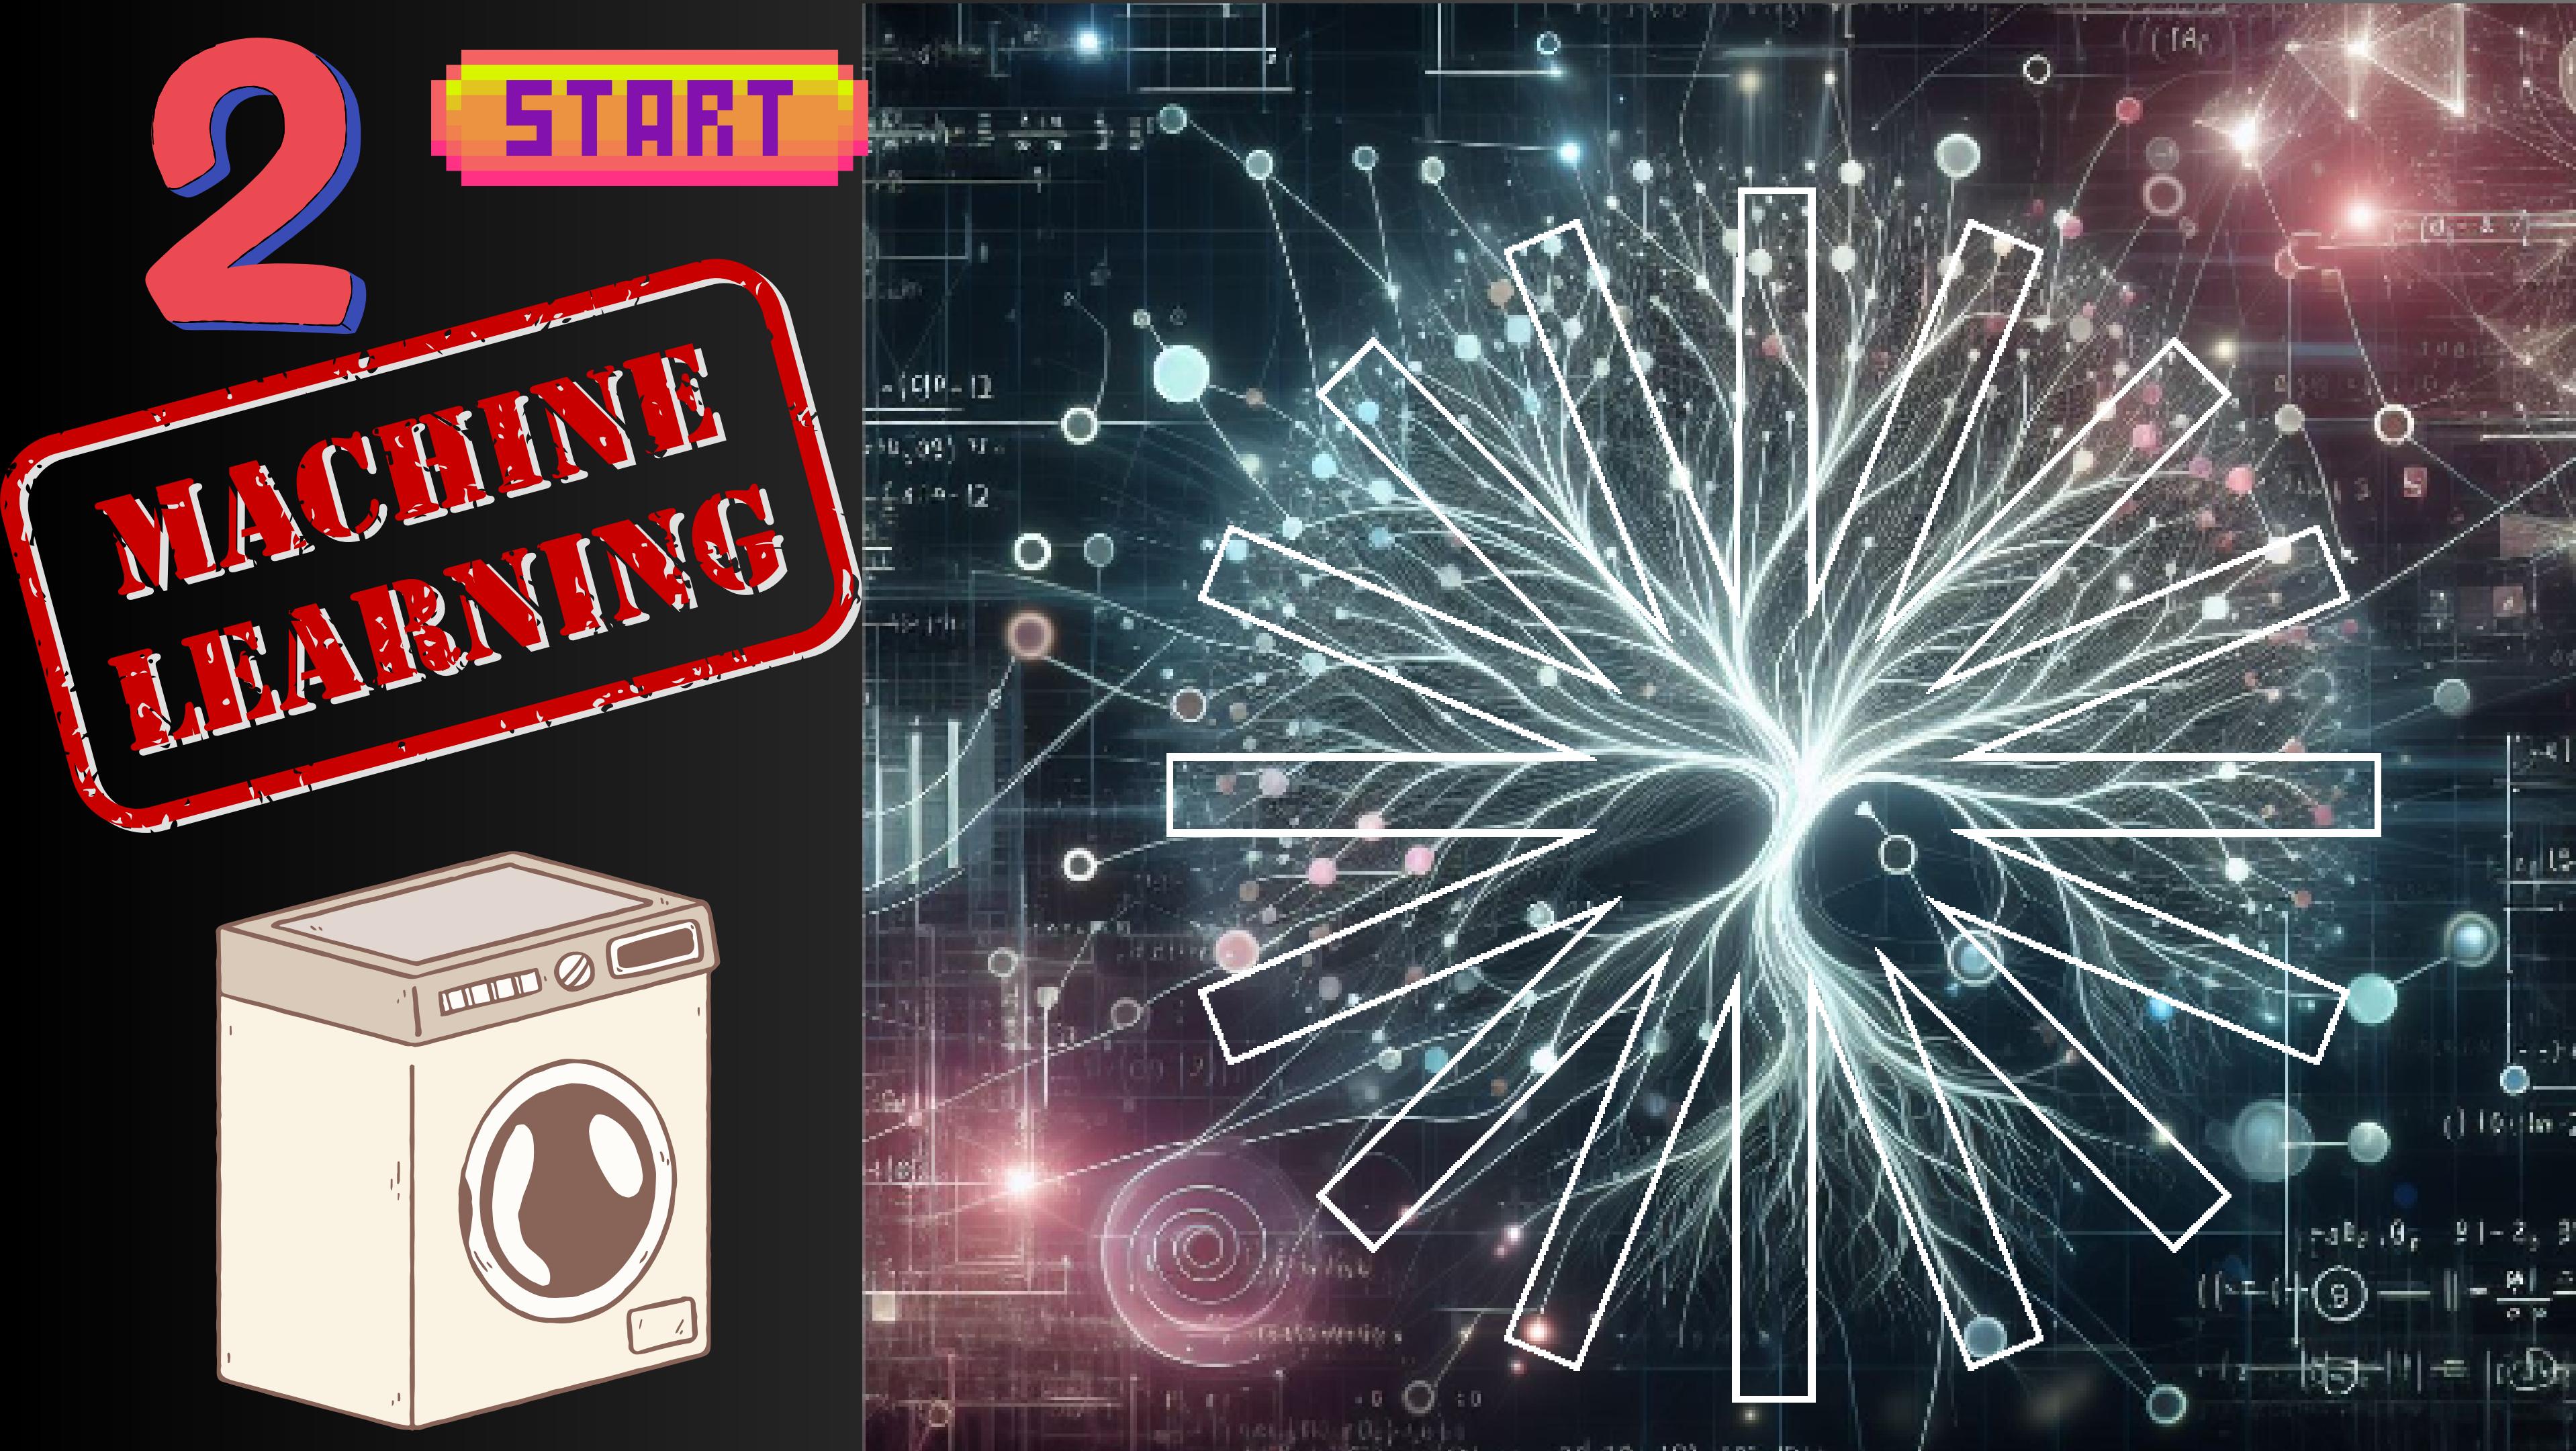
Boxplot de dia\_zscore



# FIN DEL EDA Y CUADRO DE COLUMNAS FINALES

The end

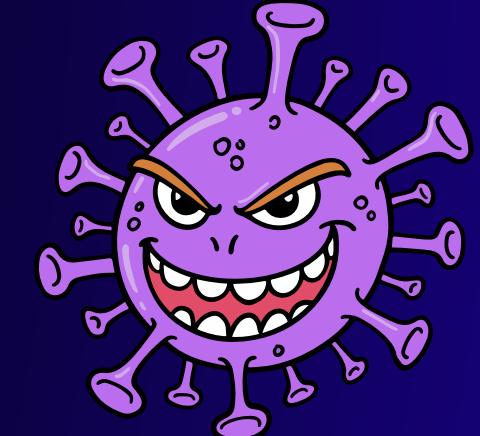
Columna:	Descripción:	Tipo:	Orden de entrada:	Información posterior a la etiqueta "label"
id.resp_h	Identificador único del host de destino	Categórico	1	No
proto	Protocolo utilizado (TCP, UDP, ICMP, etc.)	Categórico	2	No
conn_state	Estado de la conexión (SYN, ACK, FIN, etc.)	Categórico	3	No
history	Historial de tráfico relacionado con esta conexión o sesión, pudiendo obtener patrones de actividad	Categórico	4	No
id.orig_p	Mes de la conexión	Numérico	5	No
orig_pkts	Número de paquetes enviados por el host de origen	Numérico	6	No
resp_pkts	Número de paquetes recibidos por el host de destino	Numérico	7	No
anyo	Año de la conexión	Numérico	8	No
hora	Hora de la conexión	Numérico	9	No
Target	Etiqueta que indica si la conexión es normal o anómala	Categórica	10	Target



# 1.- MODELO ML RANDOMFOREST CON 115 COLUMNAS

## PROCESO TRANSFORMACION

## RESULTADO ALL



60%



```
# Definir el modelo Random Forest
model = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42, n_jobs=-1)

# Entrenar el modelo
with parallel_backend('threading'):
    model.fit(X_train, y_train)

# Evaluar el modelo
y_pred = model.predict(X_val)
accuracy = accuracy_score(y_val, y_pred)

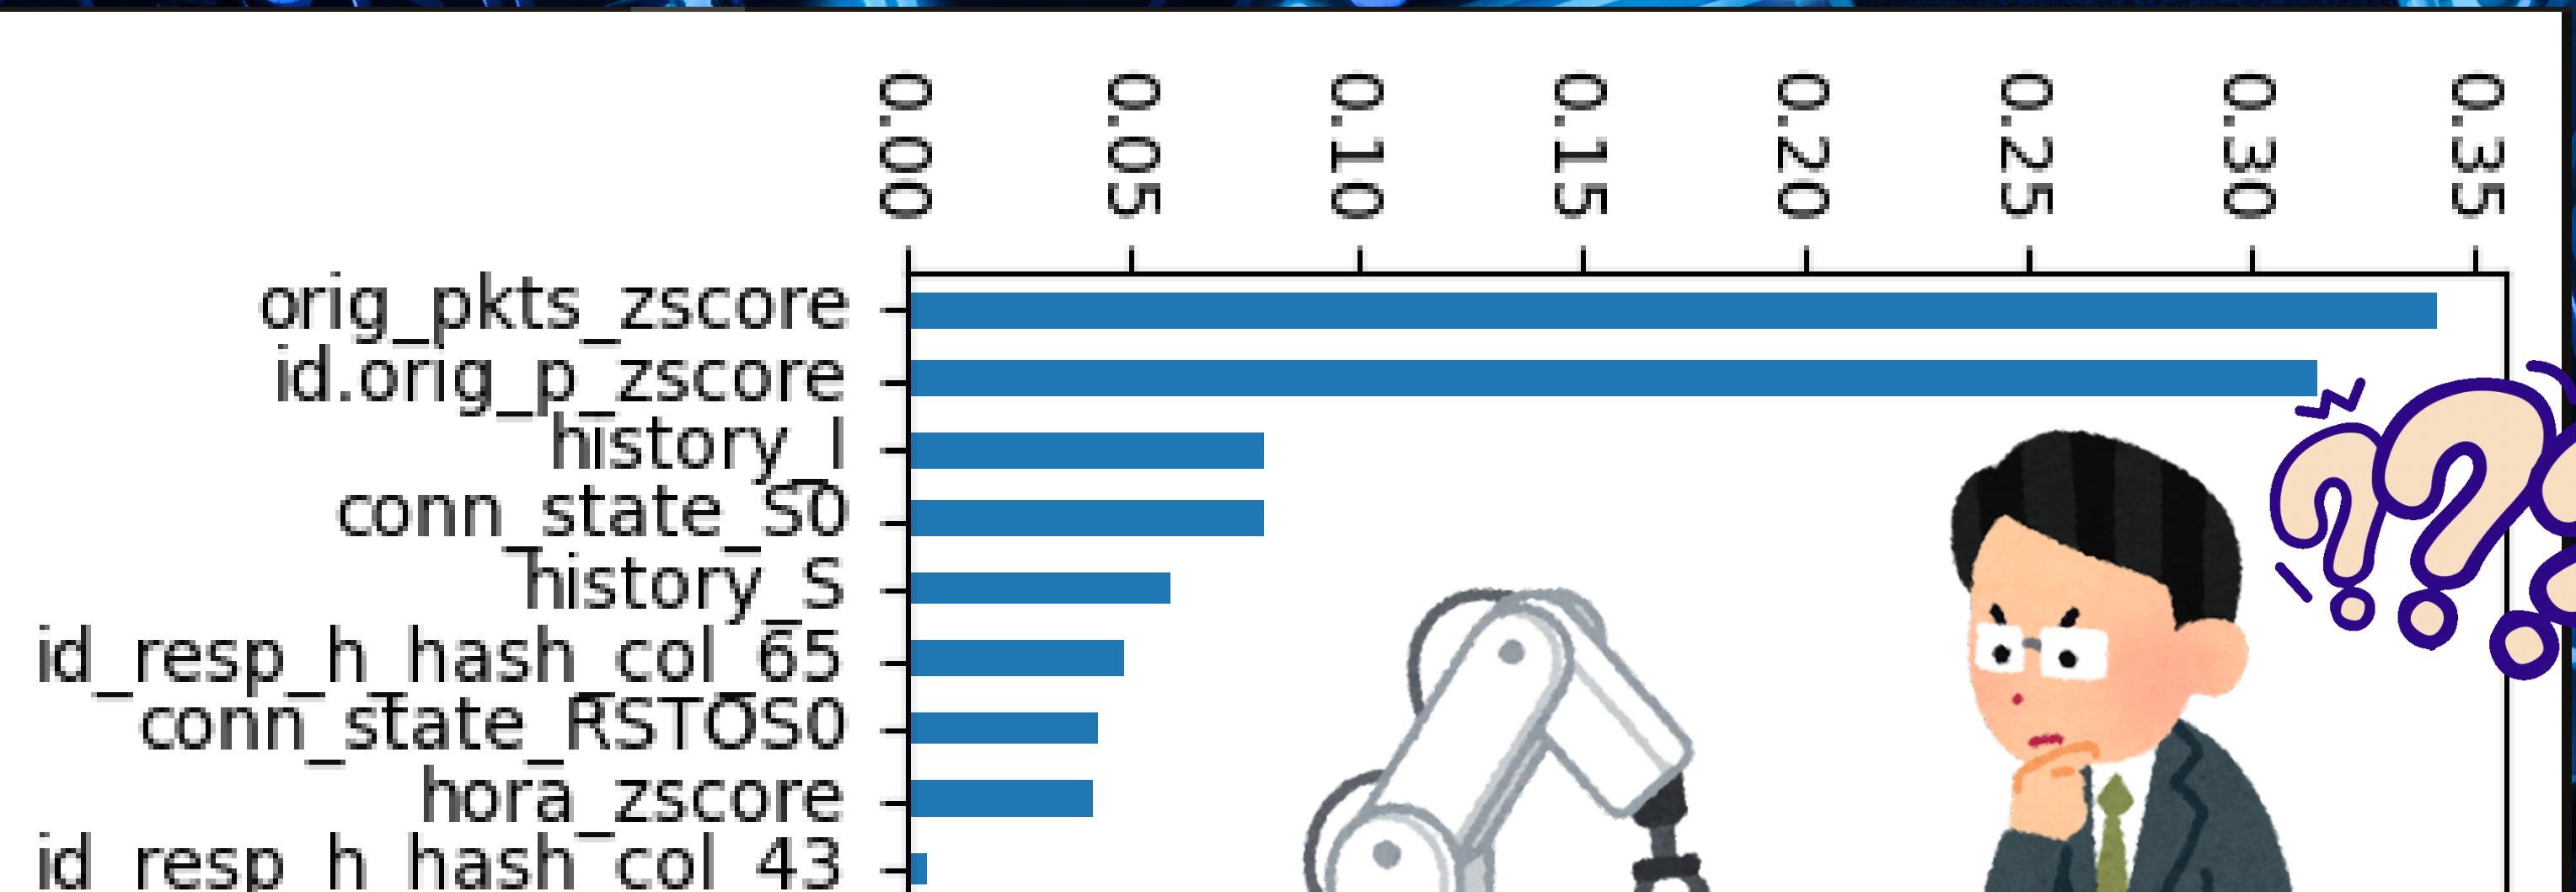
print(classification_report(y_pred, y_val))
```

WOW



	precision	recall	f1-score	support
0	0.58	1.00	0.74	808056
1	1.00	0.74	0.85	2228272
accuracy			0.81	3036328
macro avg	0.79	0.87	0.79	3036328
weighted avg	0.89	0.81	0.82	3036328

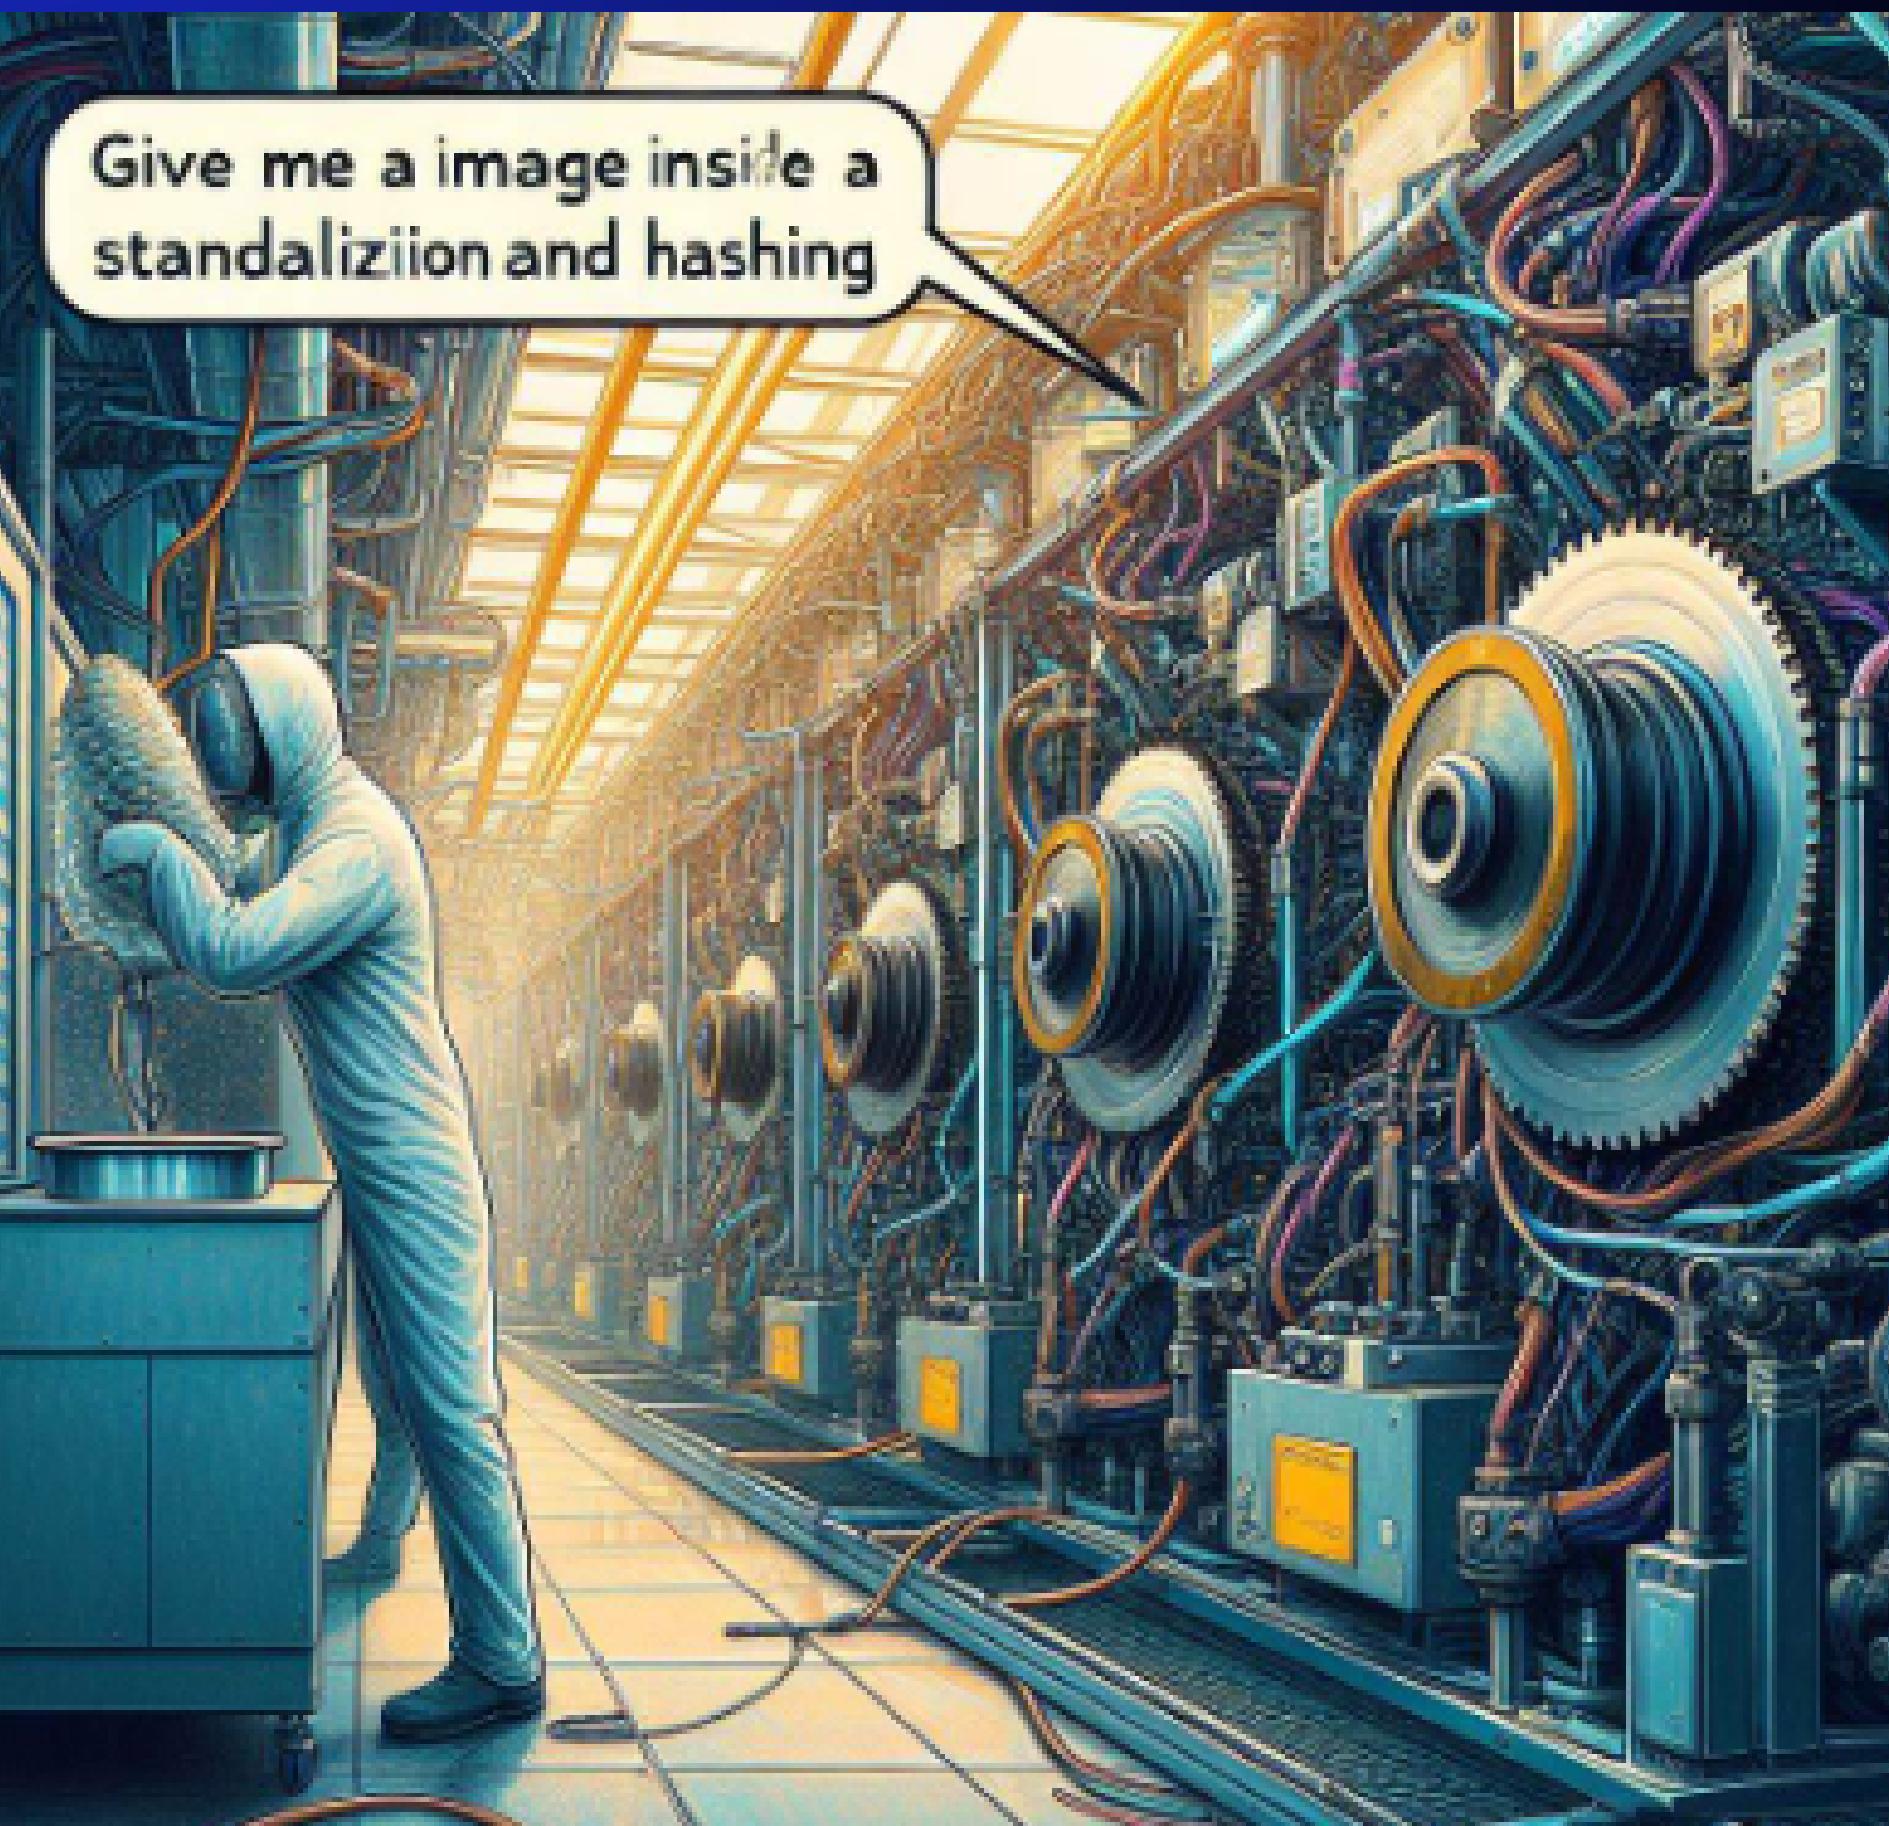
# CARACT. MÁS INFLUYENTES DURANTE EL R.FOREST.



....DE 115 COLUMNAS....

## 2.- MODELO ML LOGISTICREGRESSION CON 115 COLUMNAS

### PROCESO TRANSFORMACION



### RESULTADO

```
# Definir el modelo logistic con C para regular el sobreajuste
model = LogisticRegression(C=10, solver="sag")

# Entrenar el modelo
model.fit(X_train, y_train)

# Evaluar el modelo
y_pred = model.predict(X_val)
accuracy = accuracy_score(y_val, y_pred)

print(classification_report(y_pred, y_val))
```



	precision	recall	f1-score	support
0.0	0.57	0.70	0.62	1124927
1.0	0.79	0.69	0.73	1911401
accuracy			0.69	3036328
macro avg	0.68	0.69	0.68	3036328
weighted avg	0.71	0.69	0.69	3036328

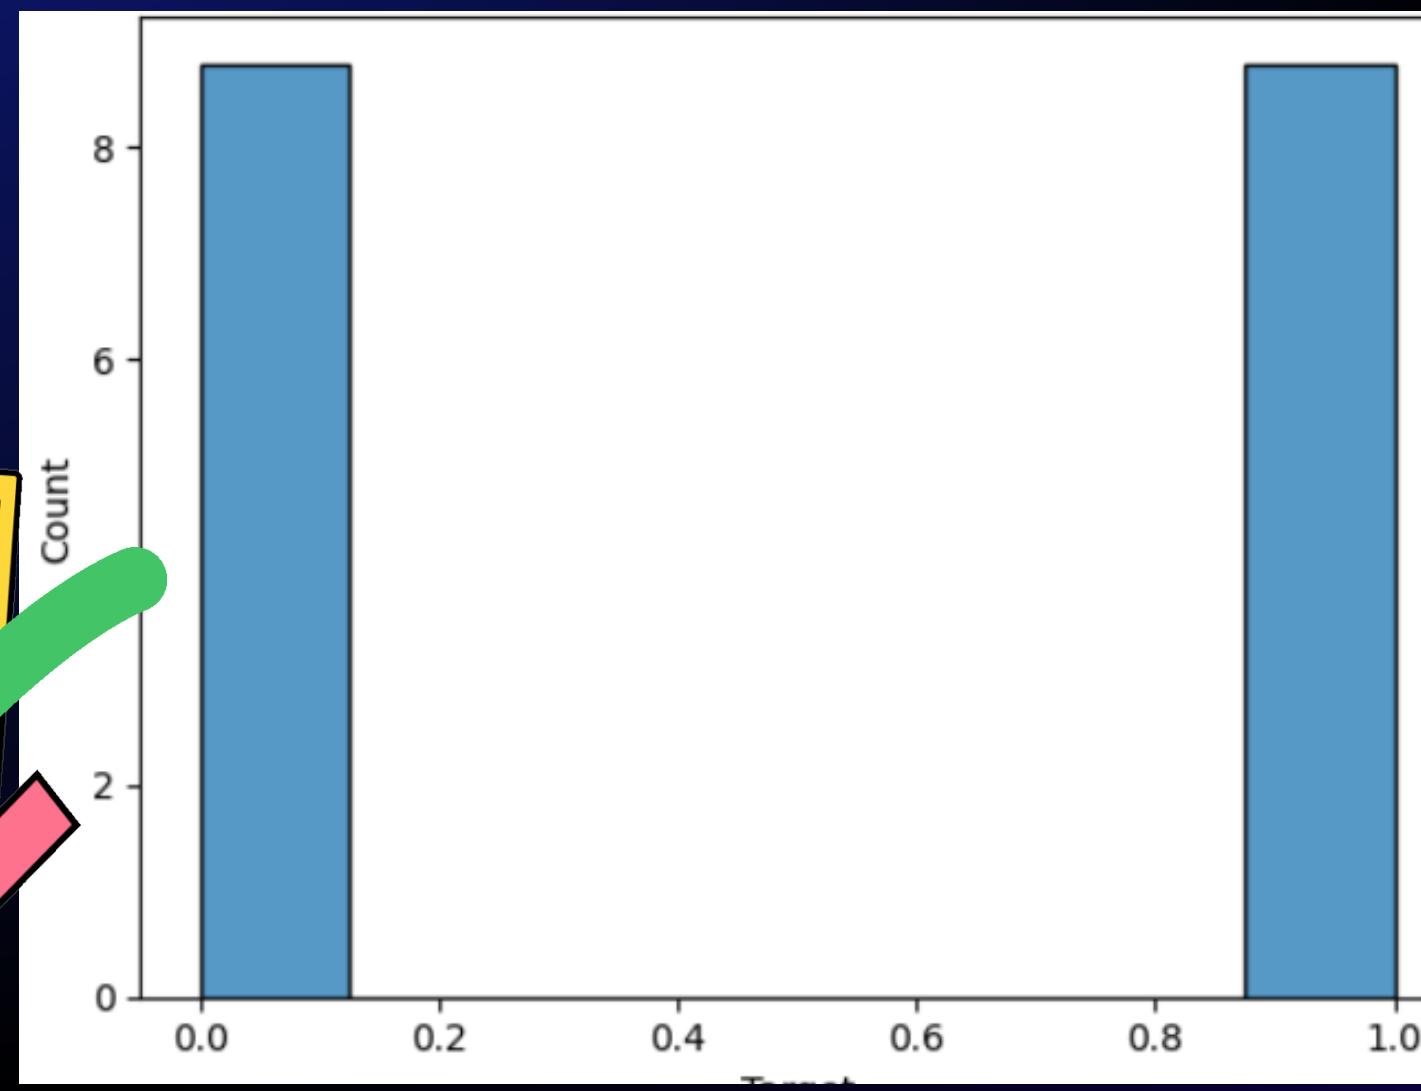
# 2.-SUBMUESTREO DEL DATASET Y CLASE MINORITARIA

## RESAMPLED

## MODELO LOGISTIC REGRESSION

```
# Aplica submuestro, para reducir el número de objetos
X=df_out.drop(["Target"], axis=1)
y=df_out["Target"]

rus = RandomUnderSampler(random_state=42)
X_resampled, y_resampled = rus.fit_resample(X, y)
```



```
# Definir el modelo
model = LogisticRegression(solver='sag', max_iter=100)

# Entrenar el modelo
with parallel_backend('threading'):
    model.fit(X_train, y_train)

# Predecir
y_pred = model.predict(X_val)

# Evaluar el modelo
print(classification_report(y_val, y_pred))
```

✓ 18m 1.1s

```
c:\Users\victo\anaconda3\envs\tf-gpu\lib\site-packages\sklearn\metrics\classification.py:51: UserWarning: Fitting estimator without sampling
warnings.warn("Fitting estimator without sampling")
```

	precision	recall	f1-score	support
0.0	0.73	0.62	0.67	1385714
1.0	0.67	0.77	0.72	1383801
accuracy				0.69 2769515
macro avg	0.70	0.69	0.69	2769515
weighted avg	0.70	0.69	0.69	2769515

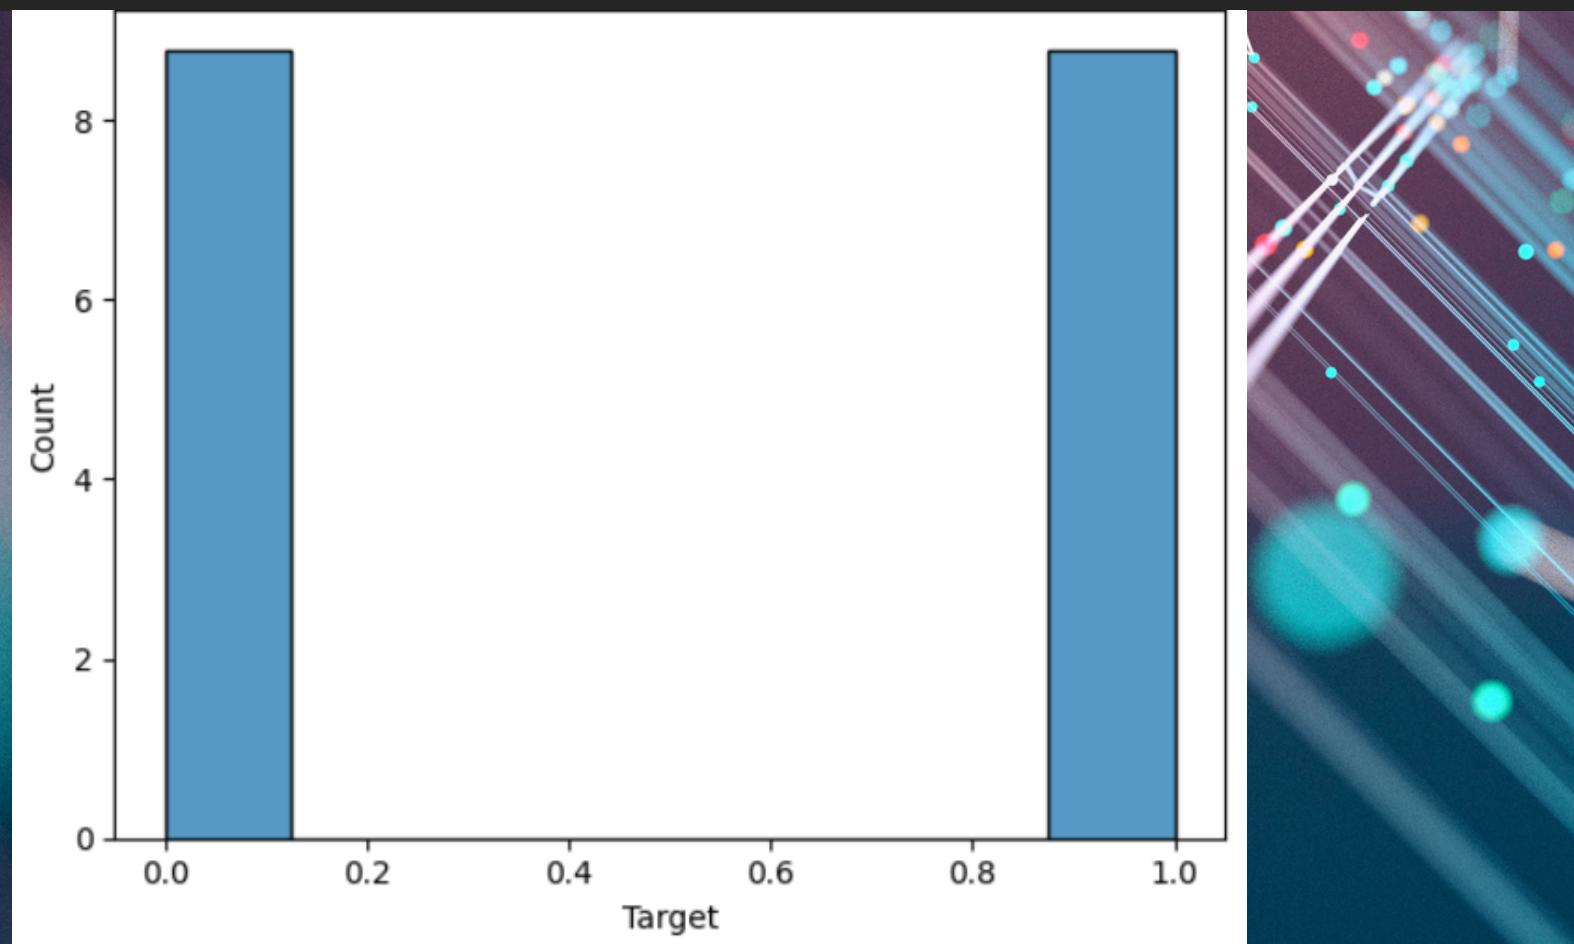


# 3.-SUBMUESTREO DEL DATASET Y CLASE MINORITARIA

## RESAMPLED

```
# Aplica submuestro, para reducir el número de observaciones
X=df_out.drop(["Target"], axis=1)
y=df_out["Target"]

rus = RandomUnderSampler(random_state=42)
X_resampled, y_resampled = rus.fit_resample(X, y)
```



## MODELO RANDOMFOREST

```
model_rf = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)

# Entrenar el modelo
# Usar parallel_backend para aprovechar la multiprocesamiento
with parallel_backend('threading'):
    model_rf.fit(X_train, y_train)

# Predecir
y_pred = model_rf.predict(X_val)

# Evaluar el modelo
print(classification_report(y_val, y_pred))
```

✓ 3m 45.0s

	precision	recall	f1-score	support
0.0	0.98	0.60	0.75	1385714
1.0	0.71	0.99	0.83	1383801
accuracy			0.80	2769515
macro avg	0.85	0.80	0.79	2769515
weighted avg	0.85	0.80	0.79	2769515



# APLICACION DE TECNICAS DE DEEP-LEARNING



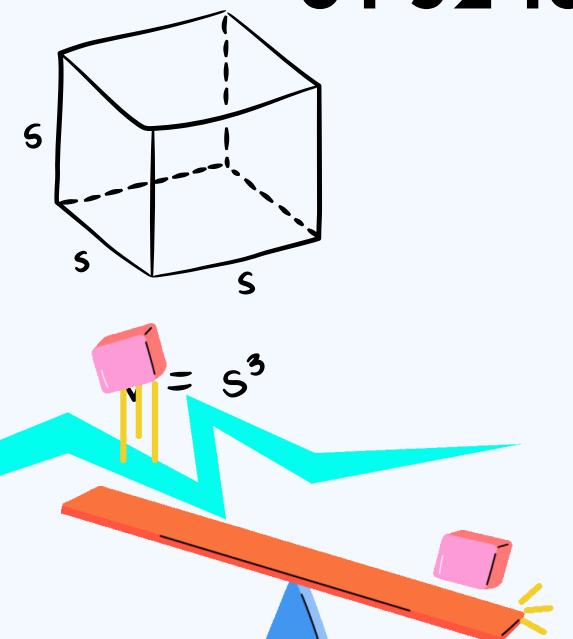
# MI MODELO DE RED NEURONAL 1

CAPAS OCULTAS

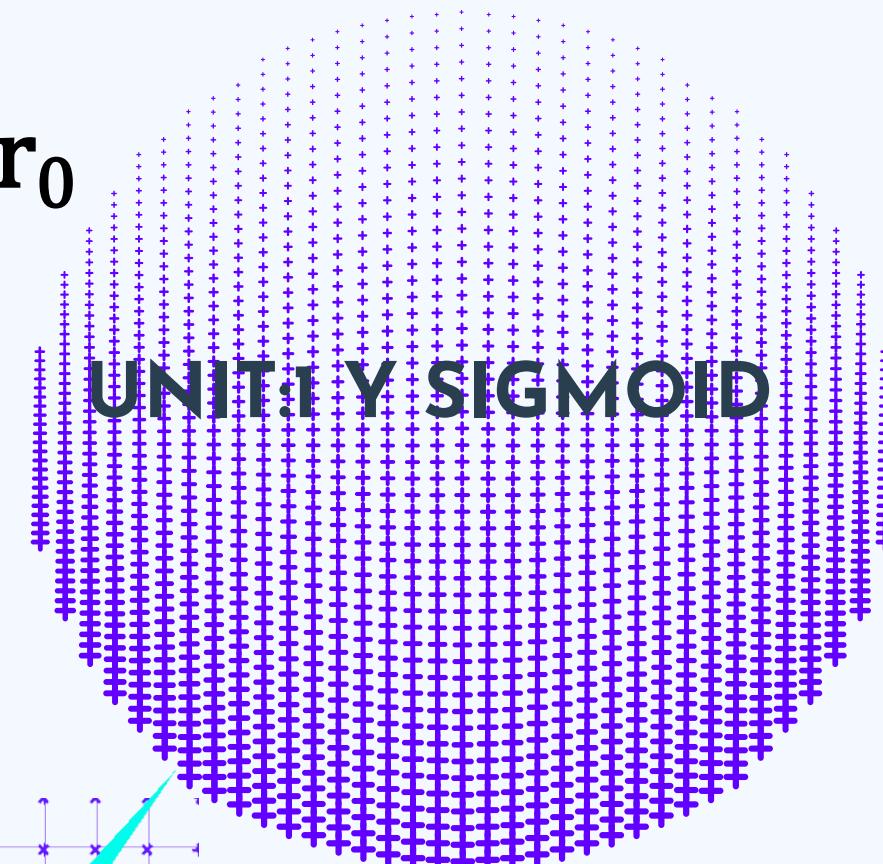
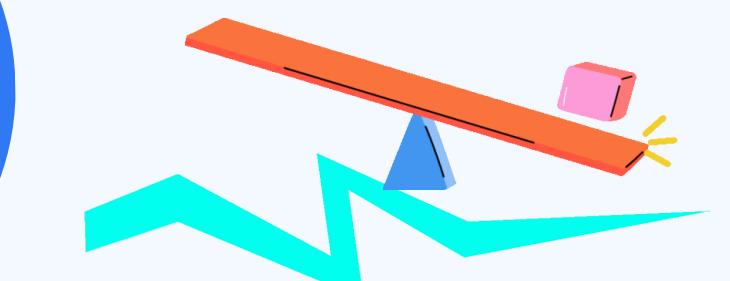
64-32-16-8 Y L2 1 CAPA Y 2 DENSA

CAPA DE SALIDA

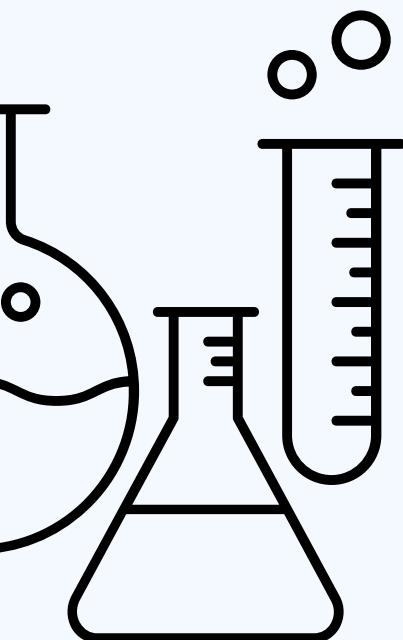
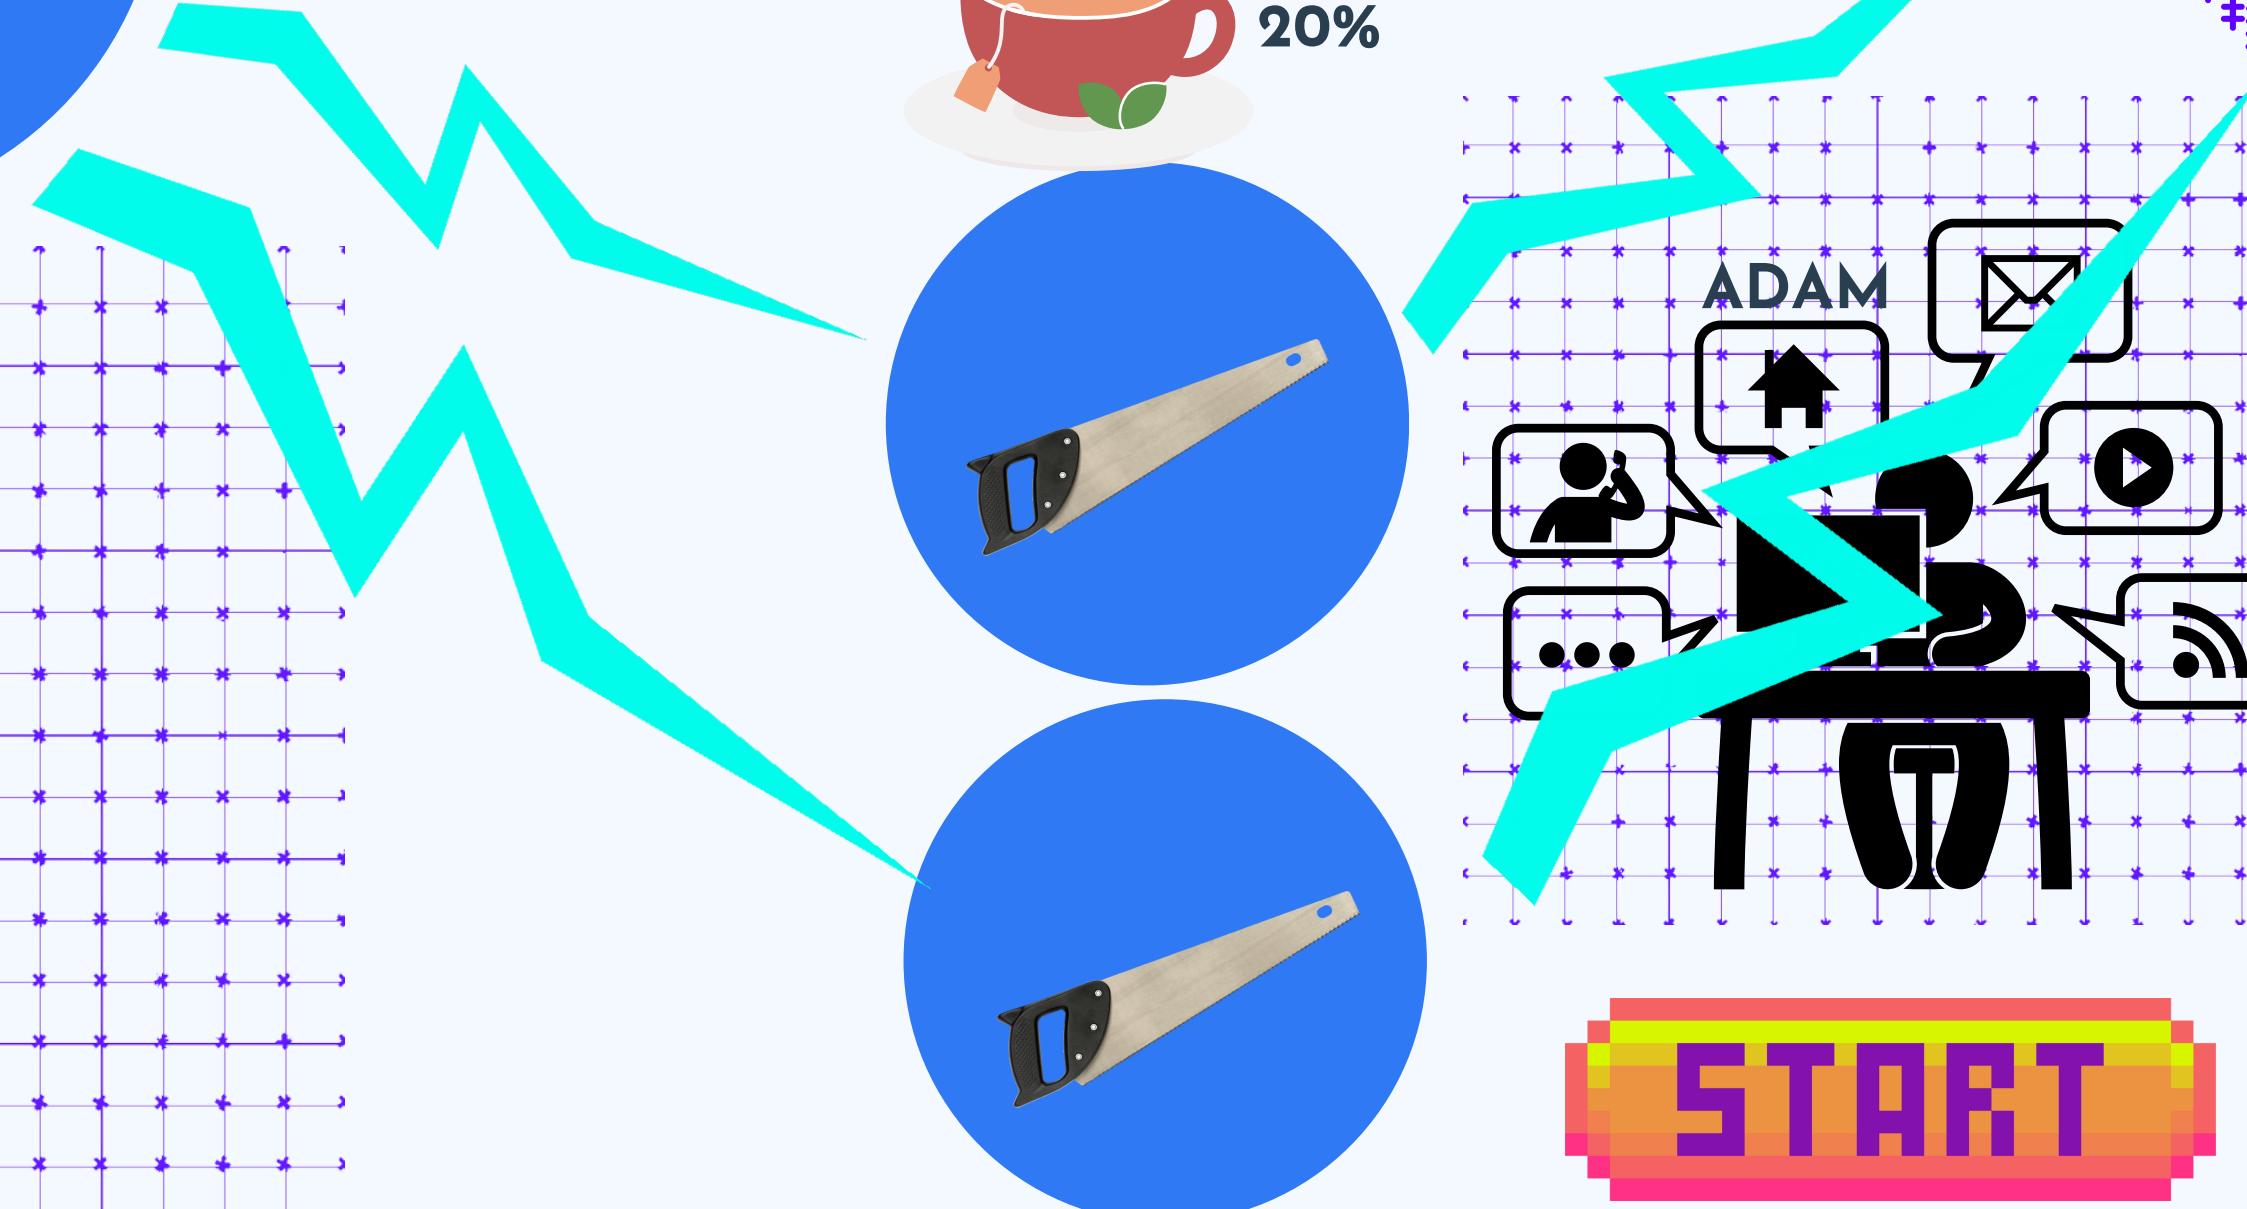
CAPA DE ENTRADA  
64 UNITS



$$r = \frac{1}{2}at^2 + v_0t + r_0$$



EARLY\_STOPPING



REDUCE\_LR

```

modelo = tf.keras.Sequential()
# Definir las 4 entradas
modelo.add(tf.keras.layers.Dense(units=64, activation='relu', input_shape=(115,), kernel_regularizer=l2(0.01)))
modelo.add(tf.keras.layers.BatchNormalization())#mejora la tasa de aprendizaje, evitando sobreajuste, ya que tb regulariza la funcion

# Capas ocultas
modelo.add(tf.keras.layers.Dense(units=32, activation="relu", input_dim=115))#relu, la mas comun, rápida y eficiente
modelo.add(tf.keras.layers.BatchNormalization())#mejora la tasa de aprendizaje, evitando sobreajuste, ya que tb regulariza la funcion

modelo.add(Dropout(0.2))

modelo.add(tf.keras.layers.Dense(units=16, activation="elu", kernel_regularizer=l2(0.01)))
modelo.add(tf.keras.layers.BatchNormalization())

modelo.add(tf.keras.layers.Dense(units=8, activation='elu'))#eLU mejora la convergencia y la velocidad de entrenamiento
modelo.add(tf.keras.layers.BatchNormalization())

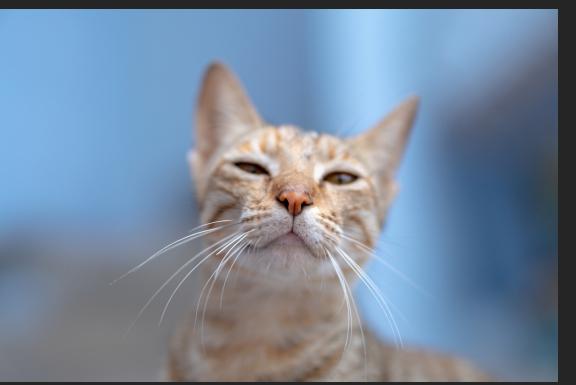
# Capa de salida
modelo.add(tf.keras.layers.Dense(units=1, activation='sigmoid')) # sigmoide para clasificación binaria
#optimizador
optimizador=Adam( learning_rate=0.01, # rebajo el learning rate por ser mas adecuado para un ajuste fino aunque tarde mas
                  beta_1=0.6, # influye en el gradiente pasado (1 mas al gradiente y 0 nada)
                  beta_2=0.4, #influye en los cuadrados de los gradientes pasados(1 mas a la varianza y 0 nada)
                  epsilon=1e-09, #es para prevenir divisiones entre 0
                  amsgrad=False)#es una variante de Adam y ayuda con la convergencia de ambos a evitar oscilaciones de los pesos durante la optimización"""
# Compilar el modelo
modelo.compile(optimizer='adam',
               loss='binary_crossentropy',
               metrics=['accuracy', 'Precision', 'Recall'])
tf.keras.metrics.AUC(name='auc')

#detiene el entrenamiento si la metrica no mejora
early_stopping_callbacks = tf.keras.callbacks.EarlyStopping(patience=5, restore_best_weights=True)

#reduce la tasa de aprendizaje cuando el rendimiento del conj. val no mejora
reduce_lr_callbacks = ReduceLROnPlateau(monitor='val_loss', patience=2, verbose=1)
#Guarda el modelo cuando mejora la métrica de validación
filepath = r'D:\Cursos\REPOSITORIOS\DATASET\malware_total\Logs\modelo.h5'
monitor = 'val_loss'
checkpoint_callbacks = ModelCheckpoint(filepath=filepath, monitor=monitor, verbose=1, save_best_only=True)
#visionado_tensorboard
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)

callbacks_list = [early_stopping_callbacks, reduce_lr_callbacks, checkpoint_callbacks, tensorboard_callback]
# Entrenar el modelo
historia = modelo.fit(X_train_a, y_train_a, batch_size=32, epochs=10,verbose=1,validation_split=0.1,callbacks=callbacks_list)

```



# modelo de red neuronal 1

## Adam y desbalanceada

# PRUEBA DE VALIDACION

```
# Evaluar el modelo en el conjunto de prueba  
loss, accuracy, precision, recall = modelo.evaluate(X_val_a, y_val_a)  
print("Loss:", loss)  
print("Accuracy:", accuracy)  
print("precision:", precision)  
print("recall:", recall)
```



## RESULTADO ADAM DESBALANCEADO

Loss: 0.42553311586380005  
Accuracy: 0.8055193424224854  
precision: 0.7378716468811035

recall: 0.9966943264007568

94886/94886 [=====] - 75s 787u

```
c:\Users\victo\anaconda3\envs\tf-gpu\lib\site-packages\  
    _warn_prf(average, modifier, msg_start, len(result))  
c:\Users\victo\anaconda3\envs\tf-gpu\lib\site-packages\  
    _warn_prf(average, modifier, msg_start, len(result))  
        precision      recall      f1-score     support
```

0.0	0.46	1.00	0.63	1384008
1.0	0.00	0.00	0.00	1652320

accuracy			0.46	3036328
----------	--	--	------	---------

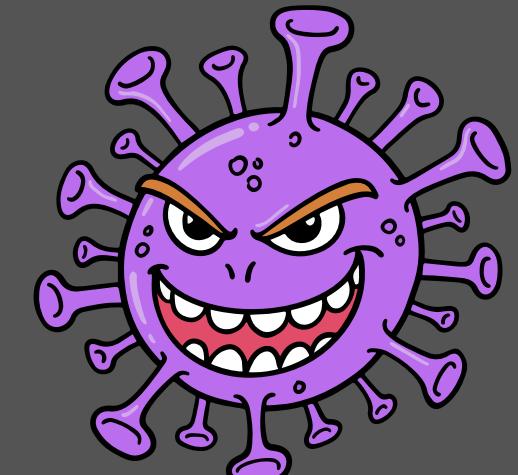
macro avg	0.23	0.50	0.31	3036328
-----------	------	------	------	---------

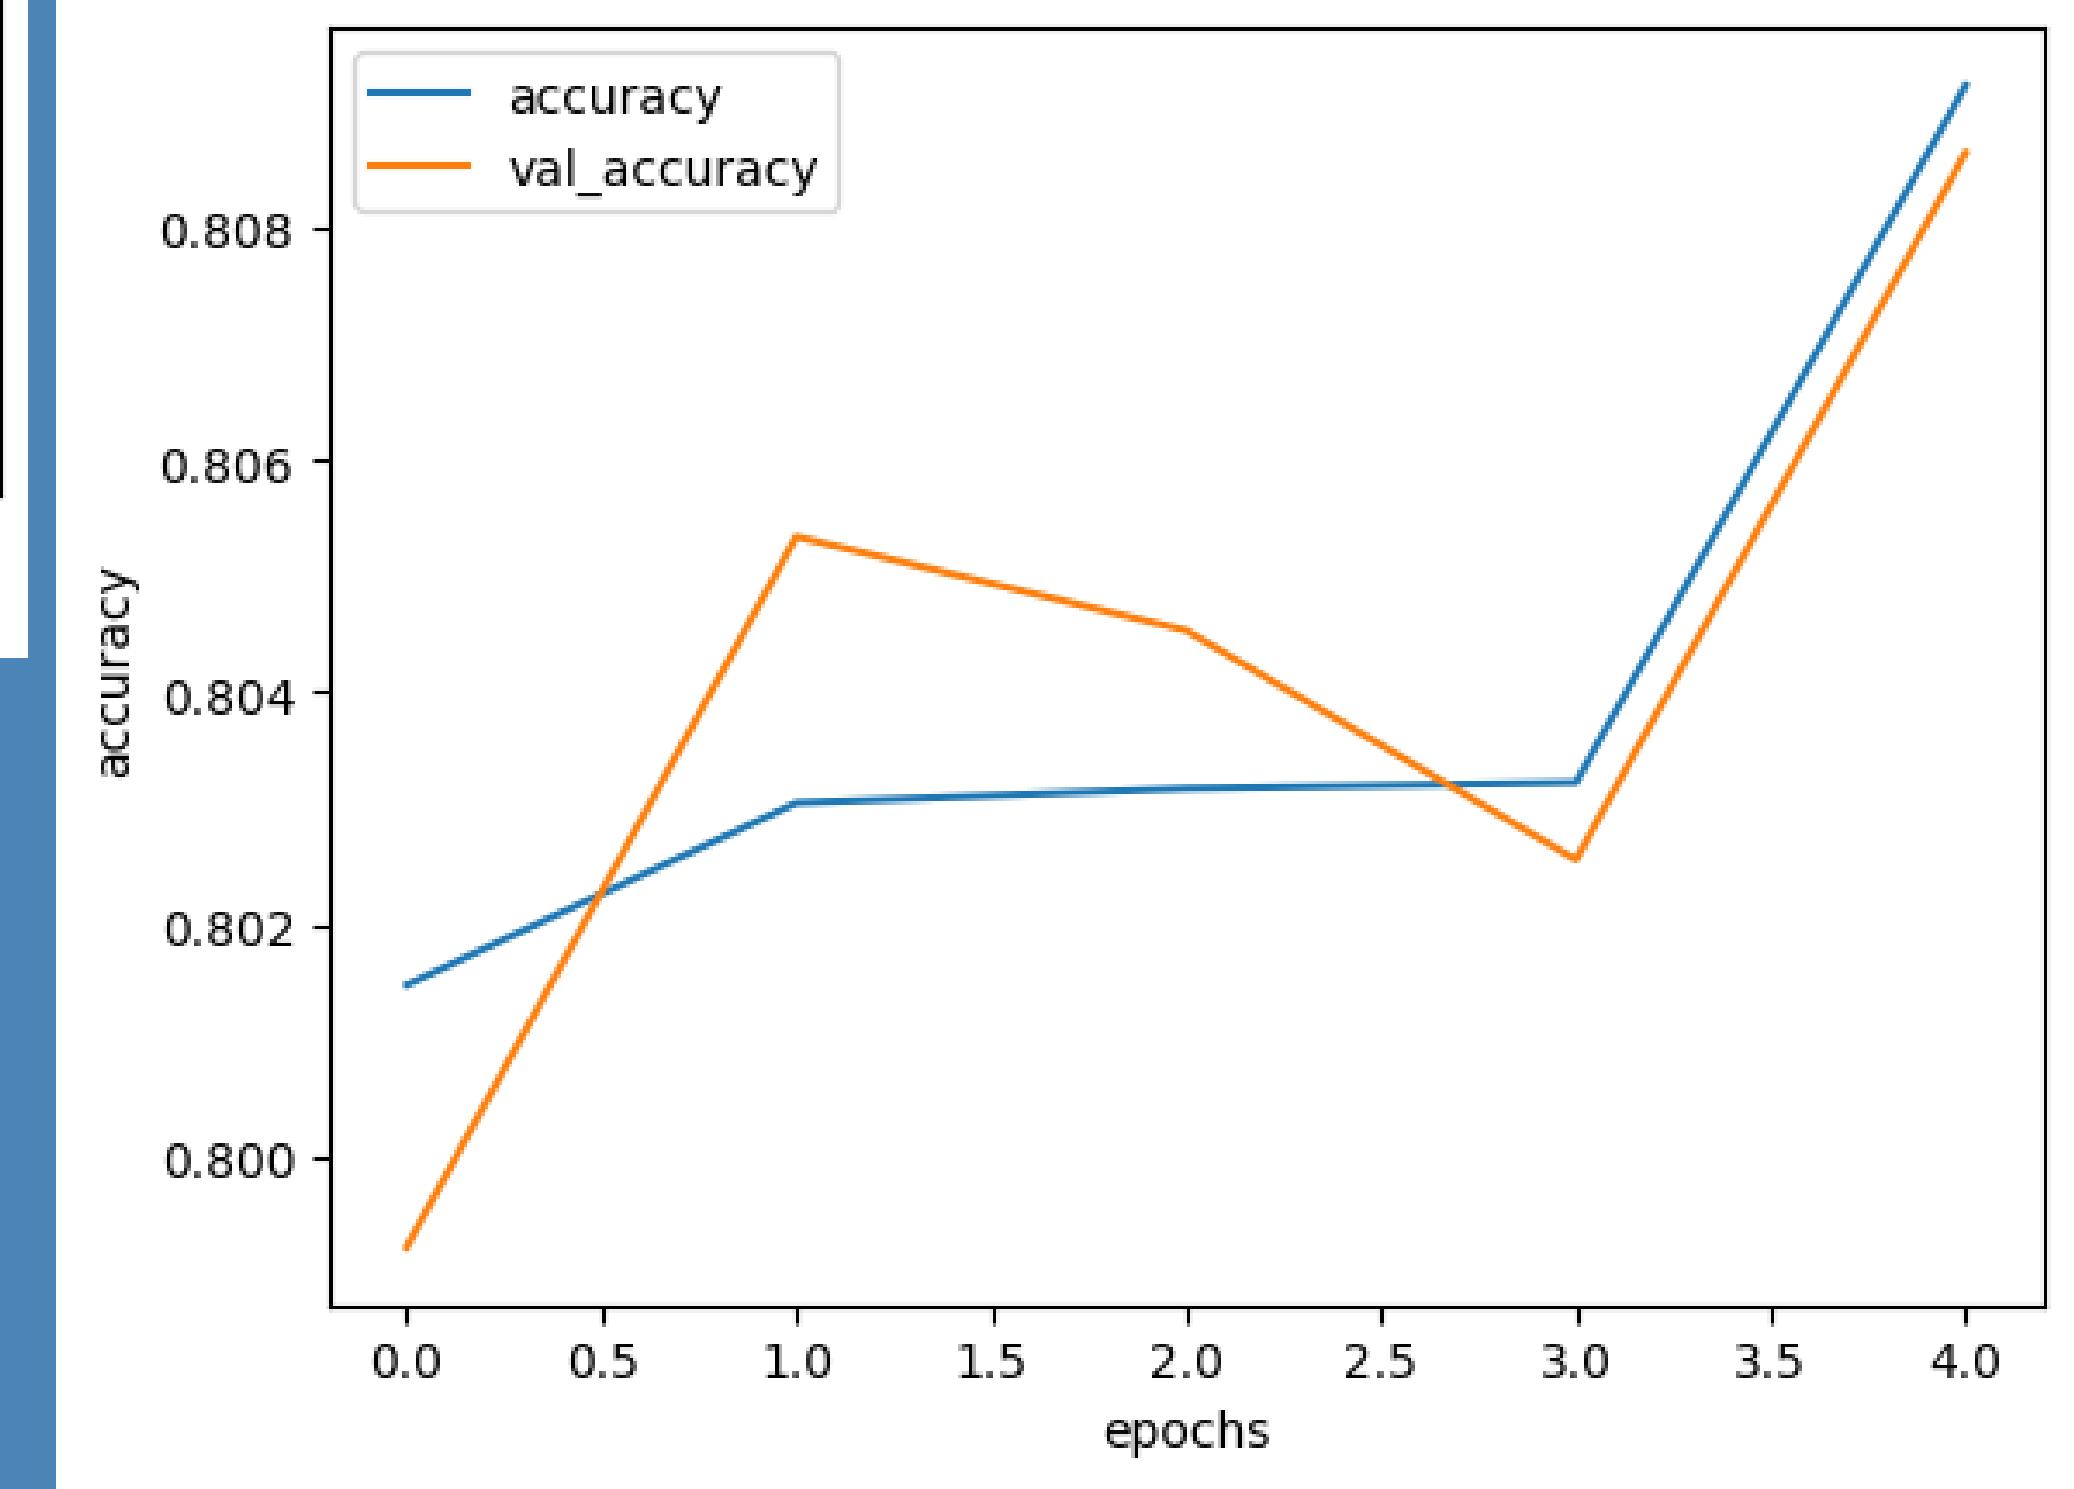
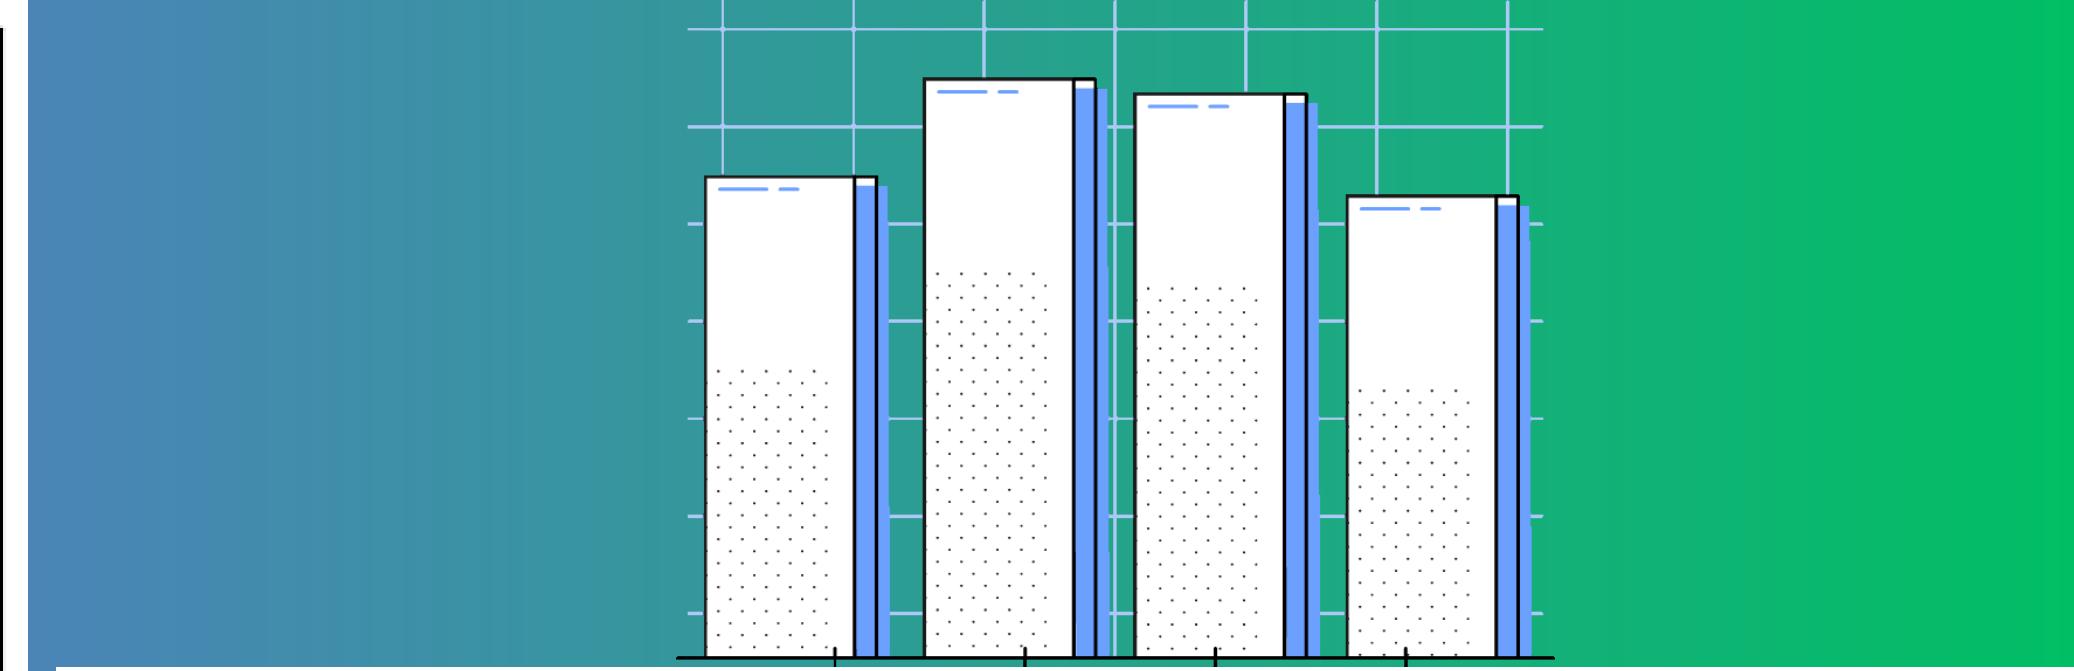
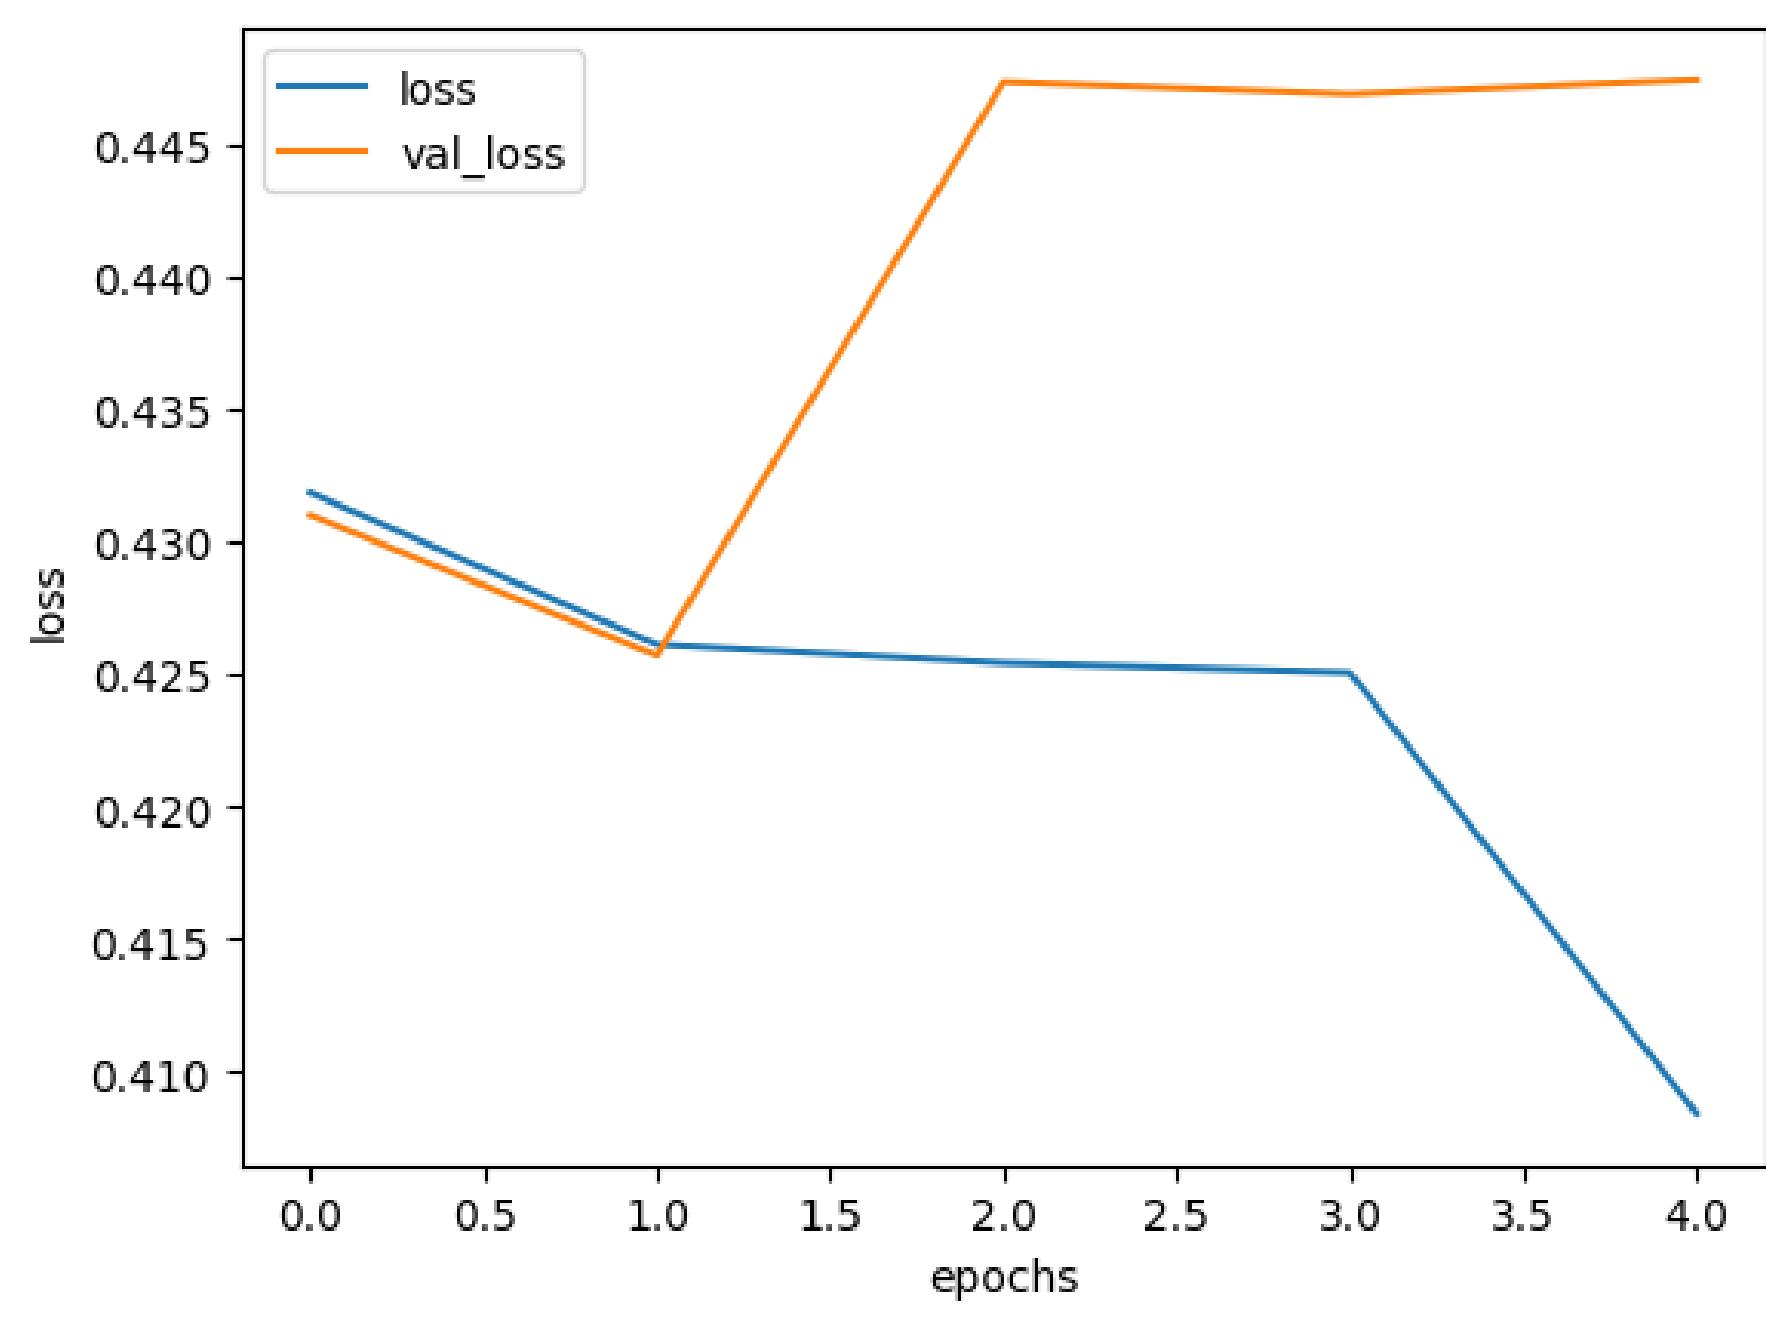
weighted avg	0.21	0.46	0.29	3036328
--------------	------	------	------	---------

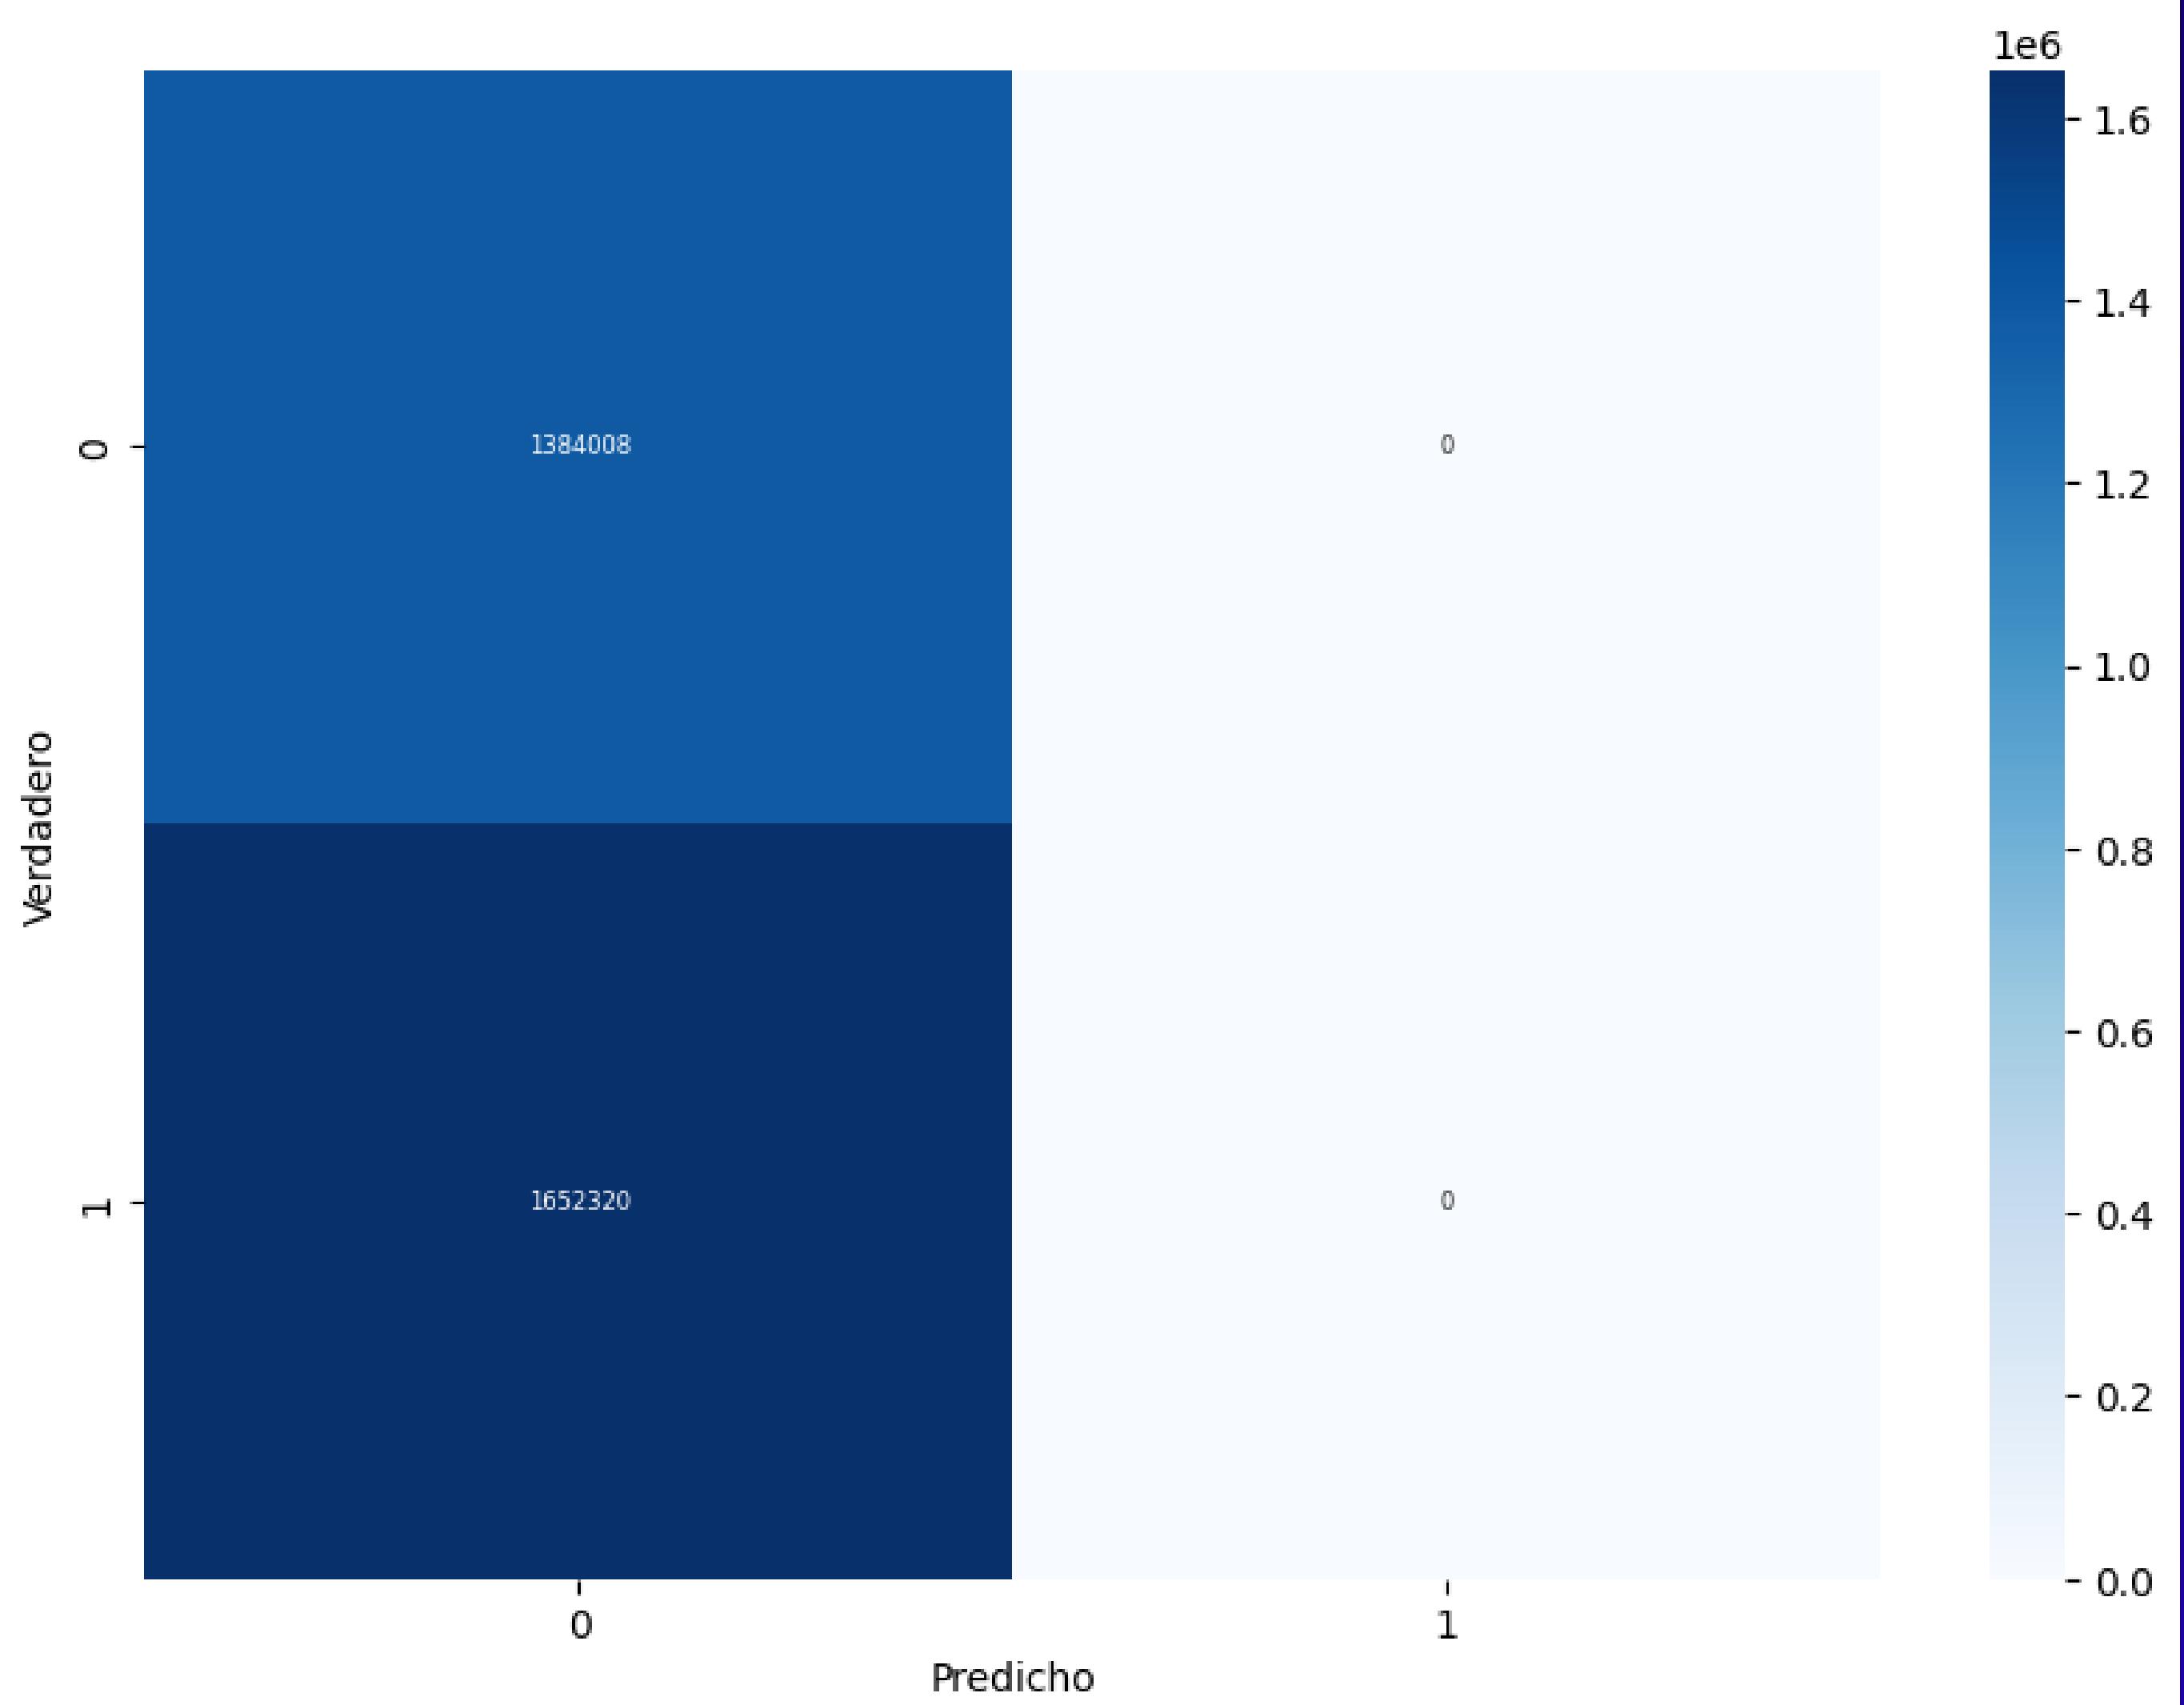
```
# Predicción  
y_pred = modelo.predict(X_val_a)  
y_pred_classes = np.argmax(y_pred, axis=1)
```

```
# Métricas de clasificación  
print(classification_report(y_val_a, y_pred_classes))
```

un 46% de Malware  
acertado de forma  
certera







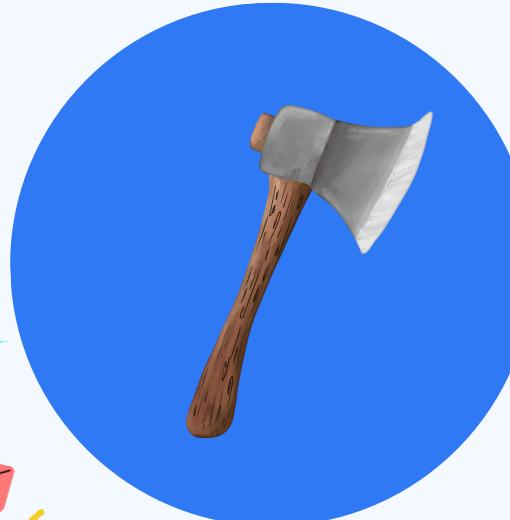
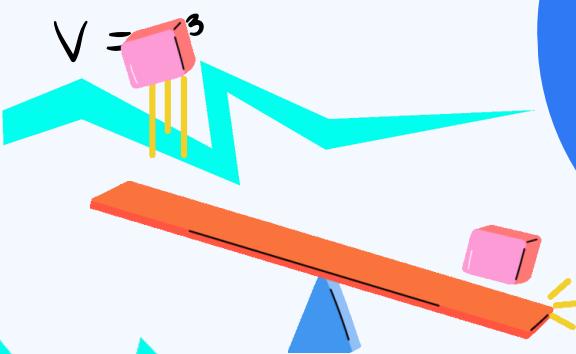
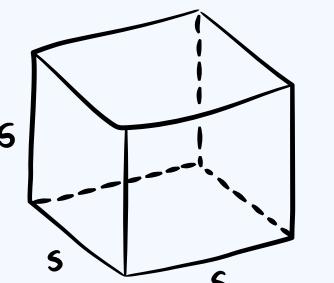
# MI MODELO DE RED NEURONAL 2

CAPA DE SALIDA

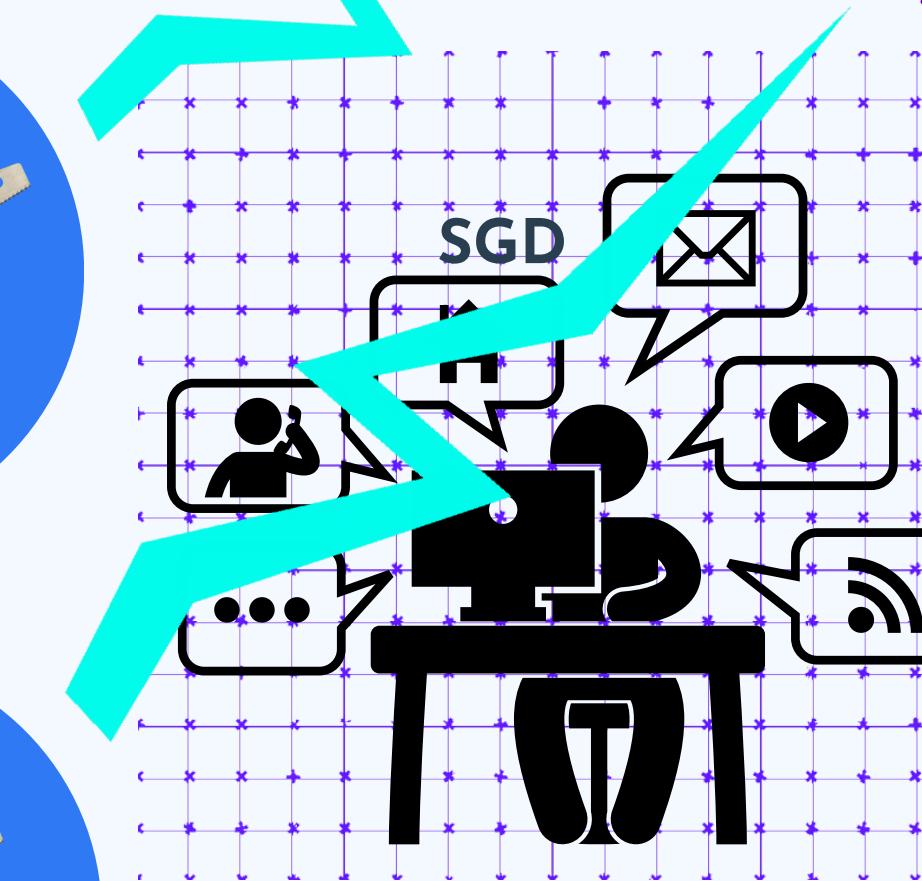
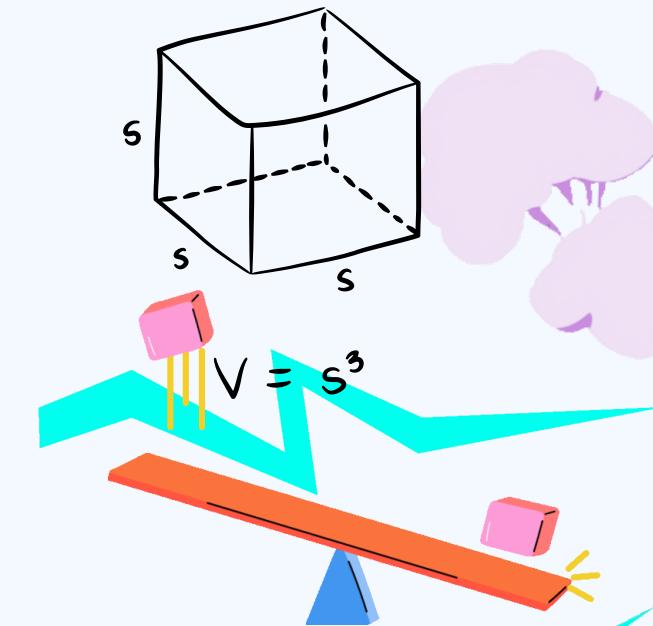
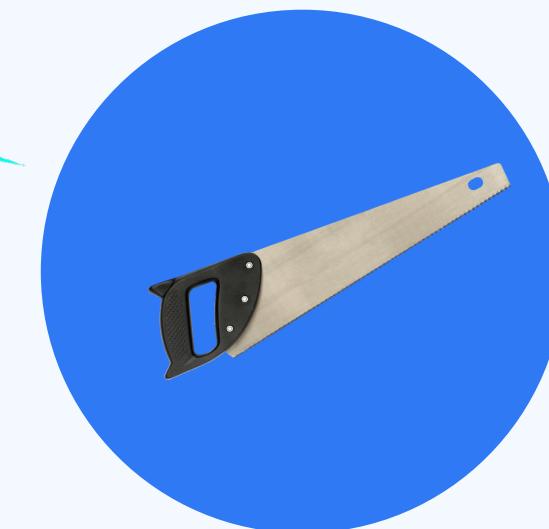
CAPA DE ENTRADA



CAPAS OCULTAS

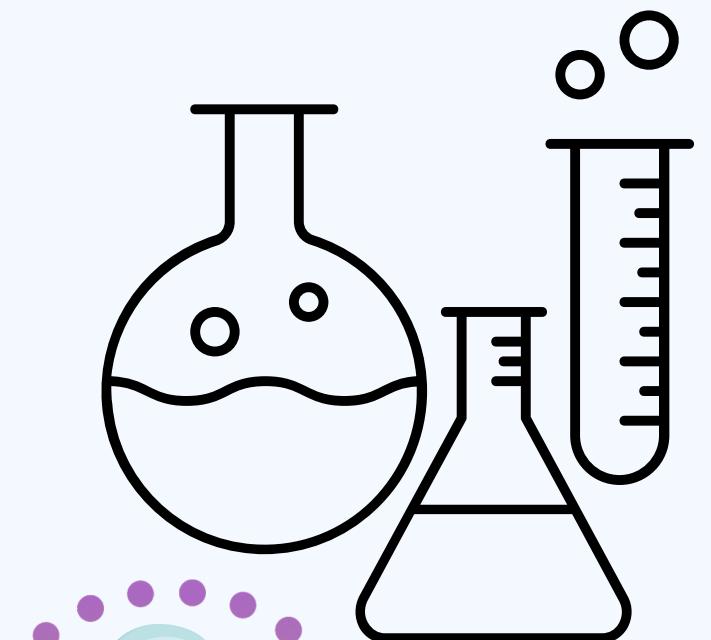


TODAS LAS CAPAS  
TIENEN 64 UNITS



UNIT 1 Y SIGMOID

EARLY\_STOPPING



REDUCE\_LR



# segundo modelo de red neuronal

```
modelo = tf.keras.Sequential()
# Definir las 4 entradas

modelo.add(tf.keras.layers.Dense(units=64, activation='relu', input_shape=(115,), kernel_regularizer=l2(0.01)))
modelo.add(tf.keras.layers.BatchNormalization())#mejora la tasa de parendizaje, evitando sobreajuste, ya que tb regulariza la funcion

# Capas ocultas
modelo.add(tf.keras.layers.Dense(units=64, activation="relu", input_dim=115))#relu, la mas comun,rapida y eficiente
modelo.add(tf.keras.layers.BatchNormalization())#mejora la tasa de parendizaje, evitando sobreajuste, ya que tb regulariza la funcion

modelo.add(Dropout(0.2))

modelo.add(tf.keras.layers.Dense(units=64, activation="elu", kernel_regularizer=l2(0.01)))
modelo.add(tf.keras.layers.BatchNormalization())

modelo.add(tf.keras.layers.Dense(units=64, activation='elu'))#eLU mejora la convergencia y la velocidad de entrenamiento
modelo.add(tf.keras.layers.BatchNormalization())

# Capa de salida
modelo.add(tf.keras.layers.Dense(units=1, activation='sigmoid')) # sigmoide para clasificacion binaria
#optimizador
optimizador= tf.keras.optimizers.SGD(learning_rate=0.005,
                                     momentum=0.9, #aporta inercia al proceso de optimizacion, acelerando la convergencia
                                     nesterov=True)# en True actualiza momentum demanera anticipada al gradiente, mejorando la estabilidad

# Compilar el modelo
modelo.compile(optimizer='SGD',
               loss='binary_crossentropy',
               metrics=[ 'accuracy', 'Precision', 'Recall'])
tf.keras.metrics.AUC(name='auc')

#detiene el entrenamiento si la metrica no mejora
early_stopping_callbacks = tf.keras.callbacks.EarlyStopping(patience=5,restore_best_weights=True)

#reduce la tasa de aprendizaje cuando el rendimiento del conj. val no mejora
reduce_lr_callbacks = ReduceLROnPlateau(monitor='val_loss', patience=2, verbose=1)

#Guarda el modelo cuando mejora la métrica de validación
filepath = r'D:\Cursos\REPOSITORIOS\DATASET\malware_total\Logs\modelo_SGD.h5'
monitor = 'val_Loss'
checkpoint_callbacks = ModelCheckpoint(filepath=filepath, monitor=monitor, verbose=1, save_best_only=True)

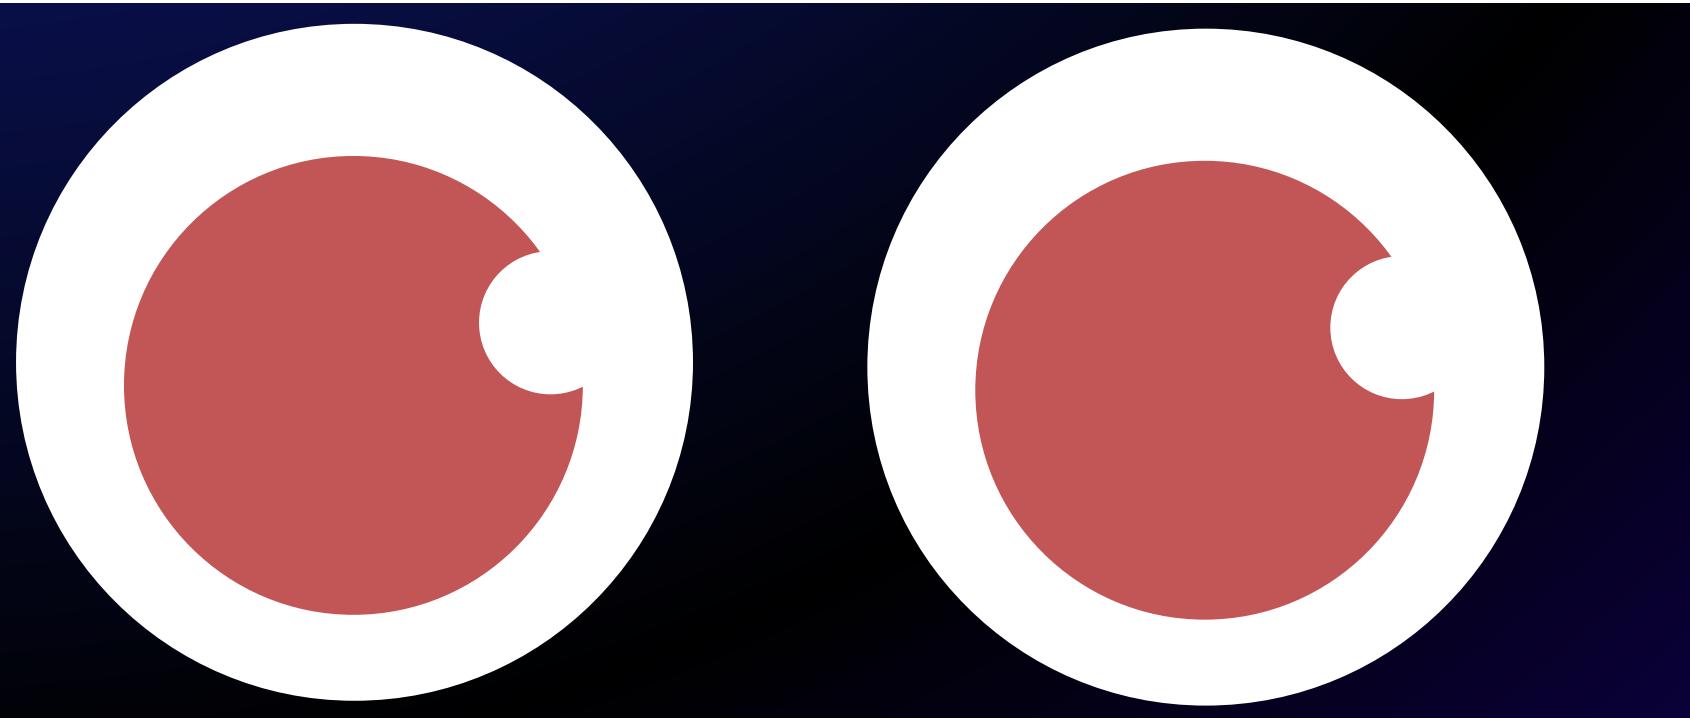
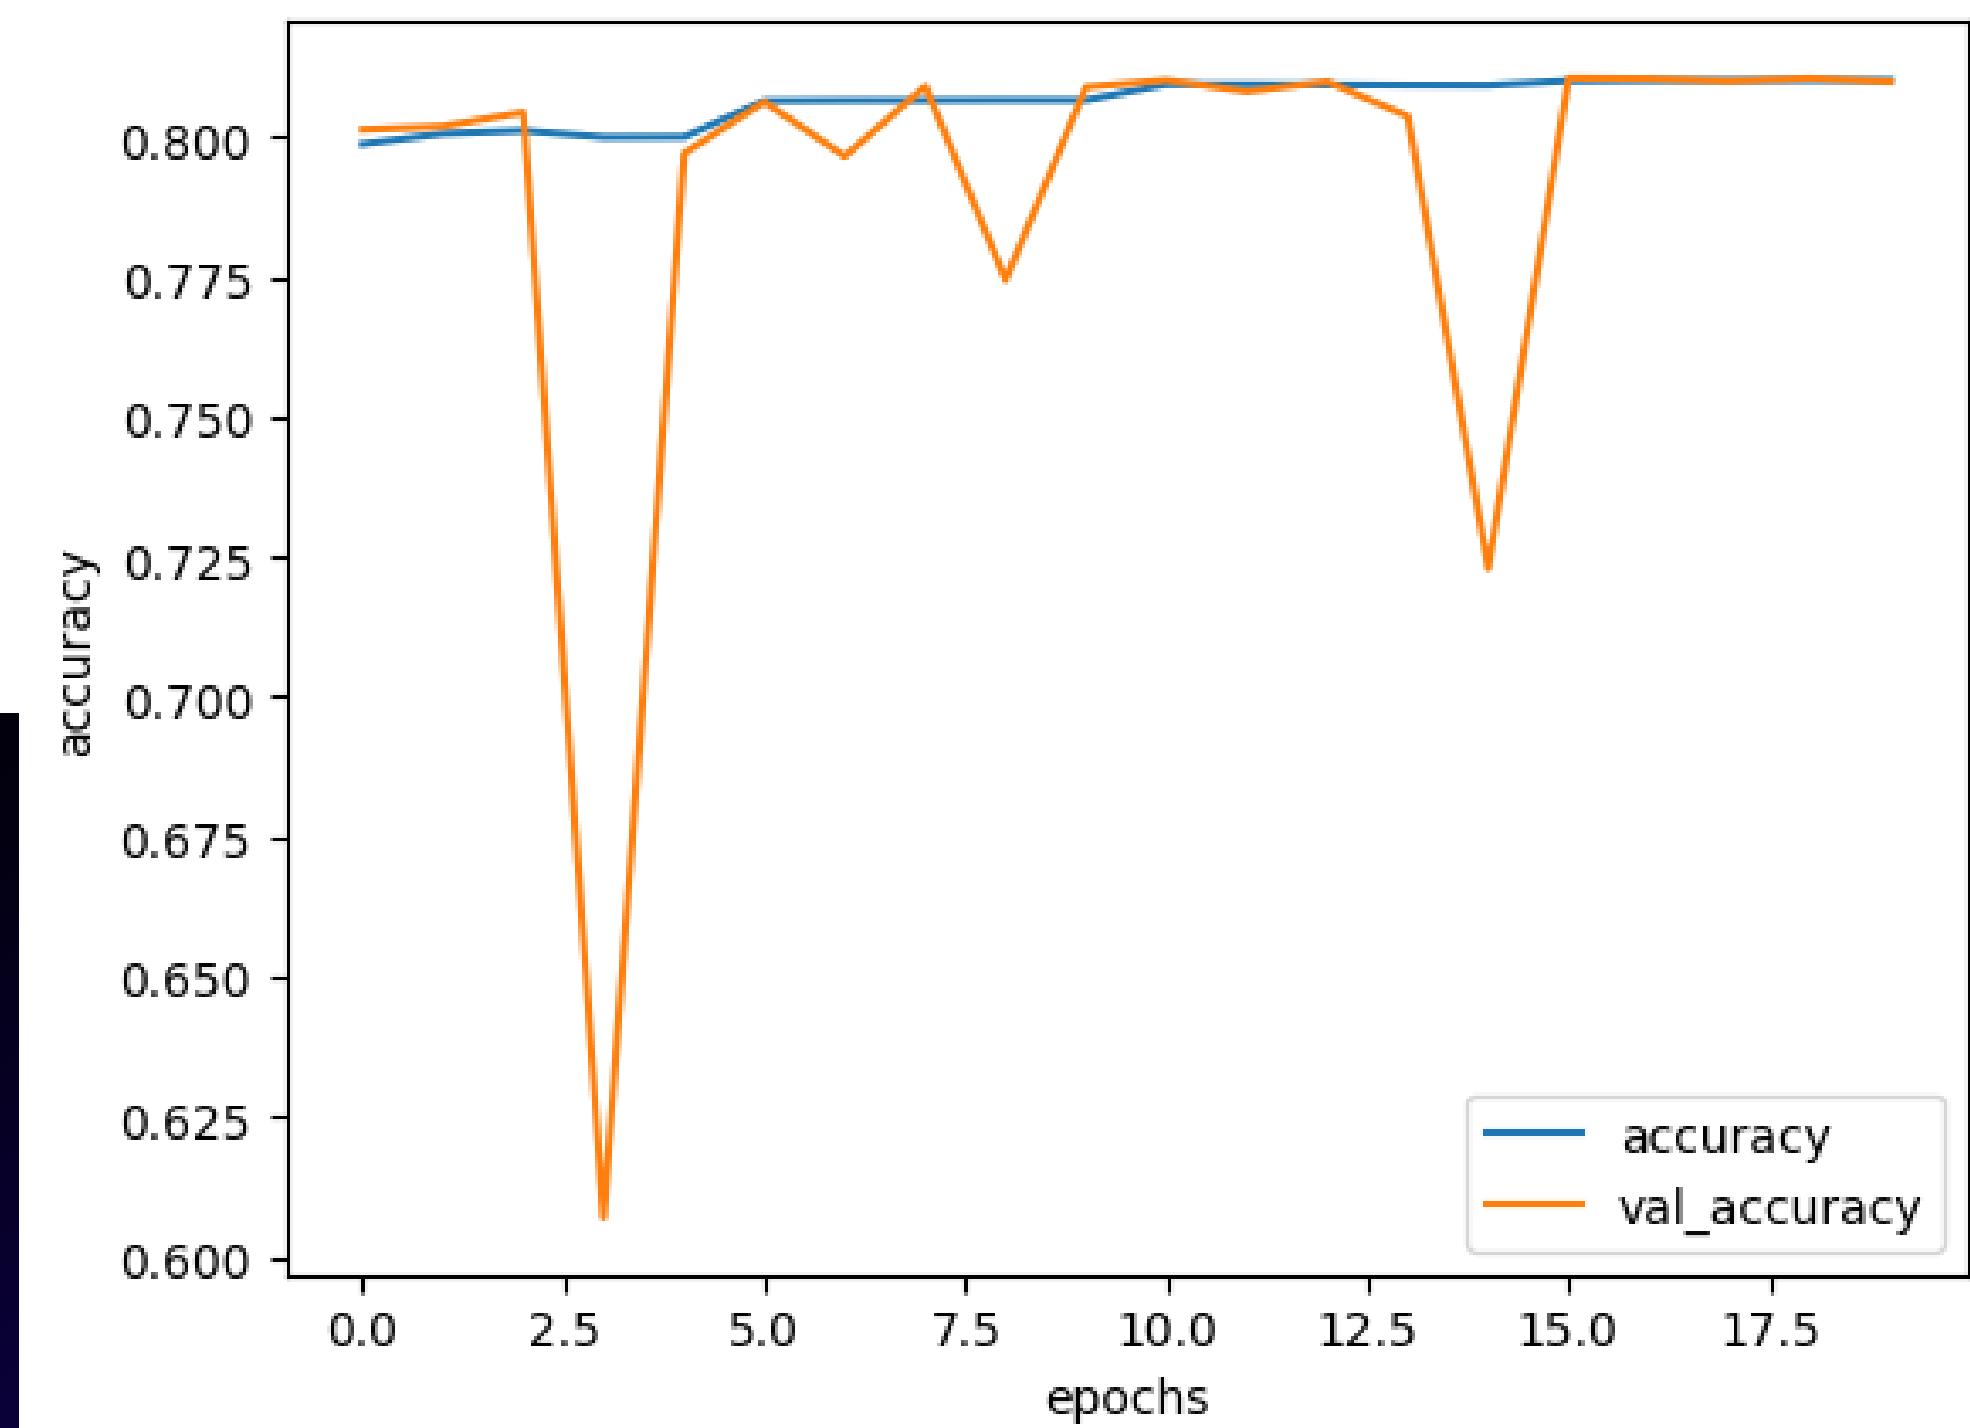
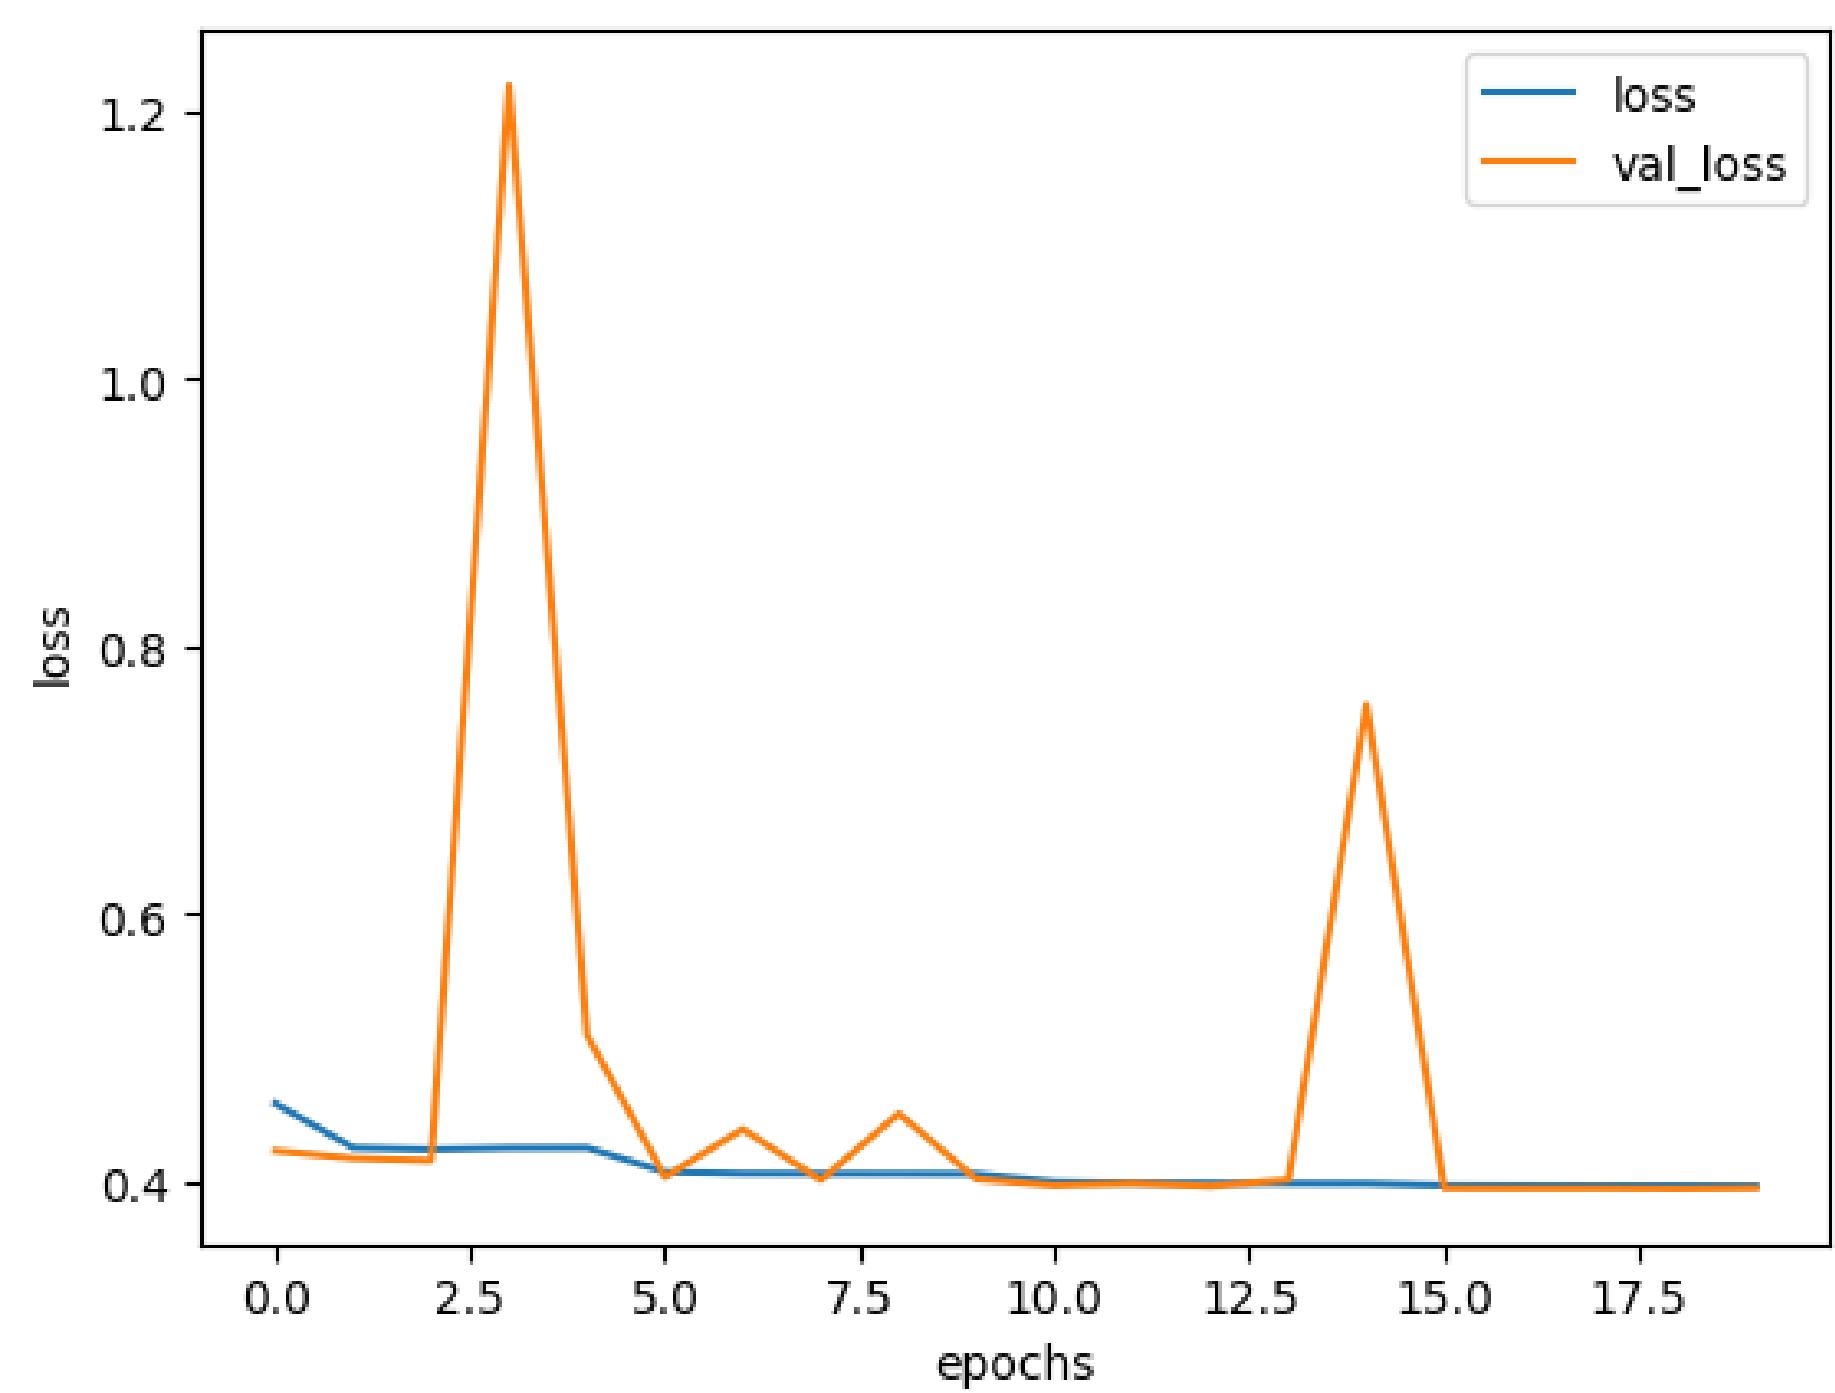
#visionado_tensorboard
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
```

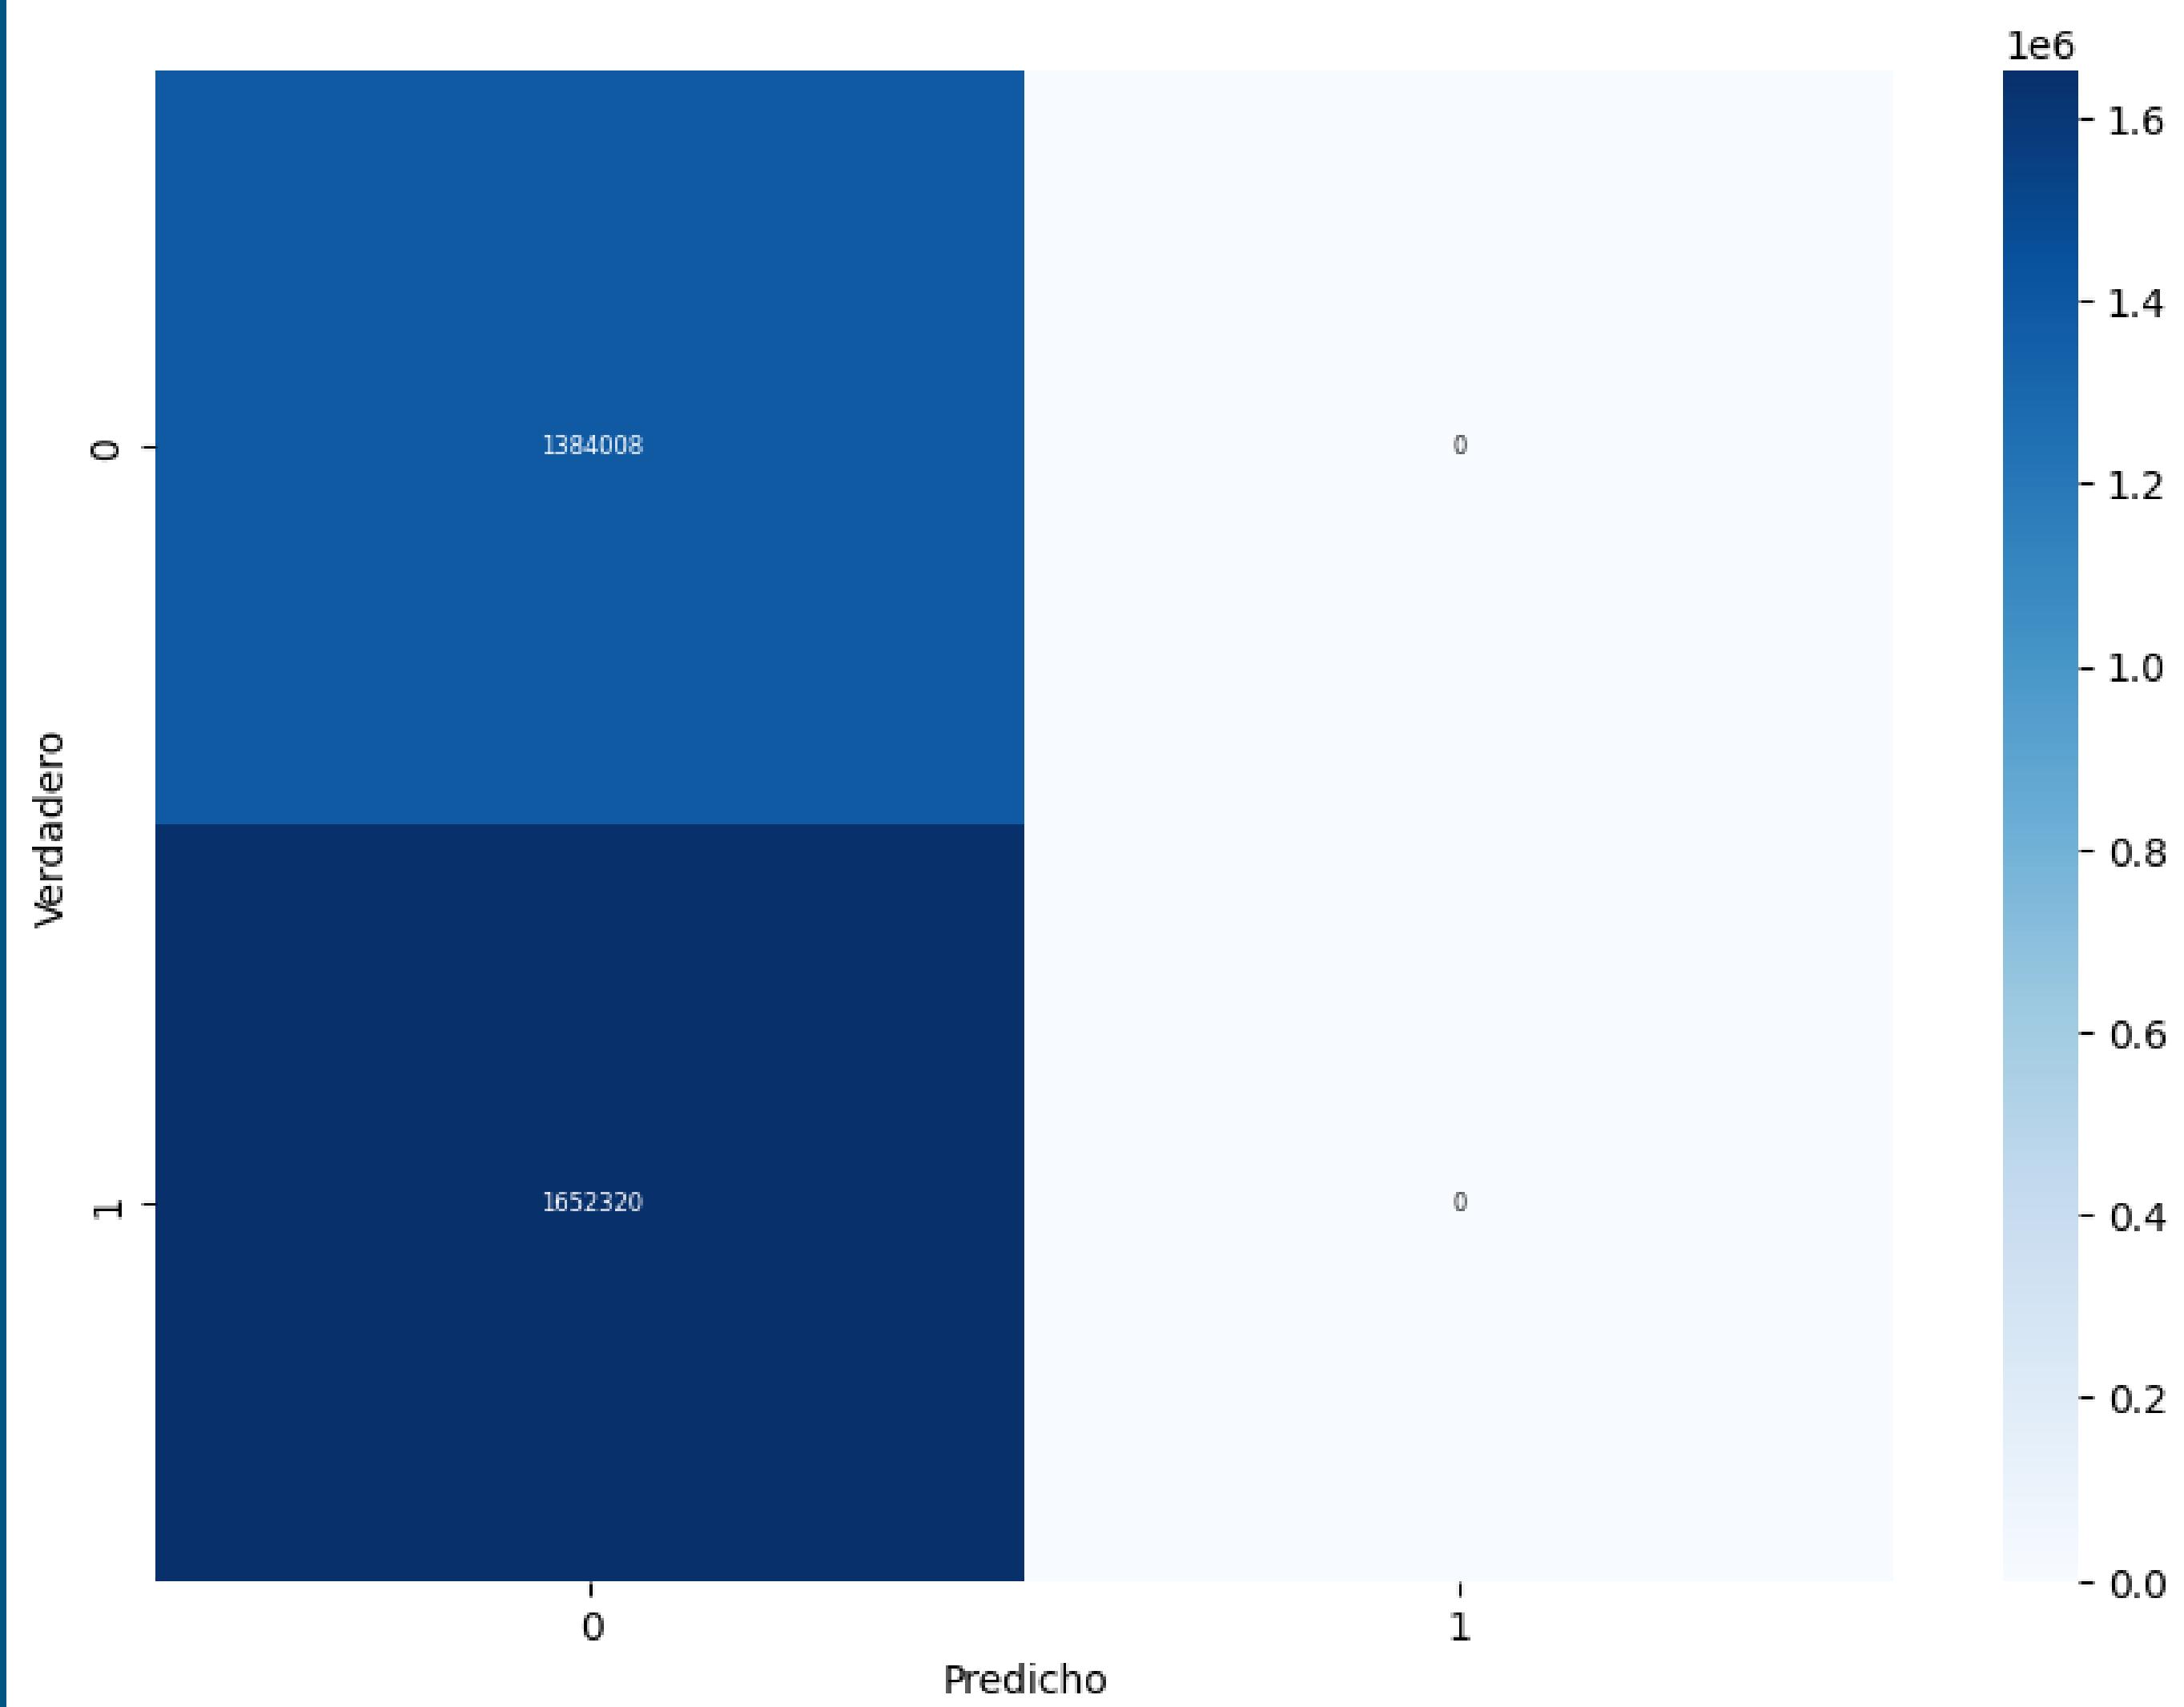
## modelo de red neuronal 2 SGD y desbalanceada

Igual que con optimizador Adam y con  
distinta configuracion de neuronas!

```
94886/94886 [=====] - 95s 997us
Loss: 0.40602222084999084
Accuracy: 0.8063921332359314
precision: 0.7384393215179443
recall: 0.9974405765533447
94886/94886 [=====] - 74s 777us
c:\Users\Victor\anaconda3\envs\tf-gpu\lib\site-packages\sklearn\metrics\classification.py:515: UserWarning: F-beta score requires precision and recall to have been computed. Need to call report_score() with arguments `labels` and `sample_weight` or `average` and `labels` or `sample_weight` before this.
  _warn_prf(average, modifier, f"{metric.capitalize()}") if metric != "accuracy"
c:\Users\Victor\anaconda3\envs\tf-gpu\lib\site-packages\sklearn\metrics\classification.py:515: UserWarning: F-beta score requires precision and recall to have been computed. Need to call report_score() with arguments `labels` and `sample_weight` or `average` and `labels` or `sample_weight` before this.
  _warn_prf(average, modifier, f"{metric.capitalize()}") if metric != "accuracy"
      precision      recall   f1-score   support
          0.0        0.46      1.00      0.63    1384396
          1.0        0.00      0.00      0.00    1651932

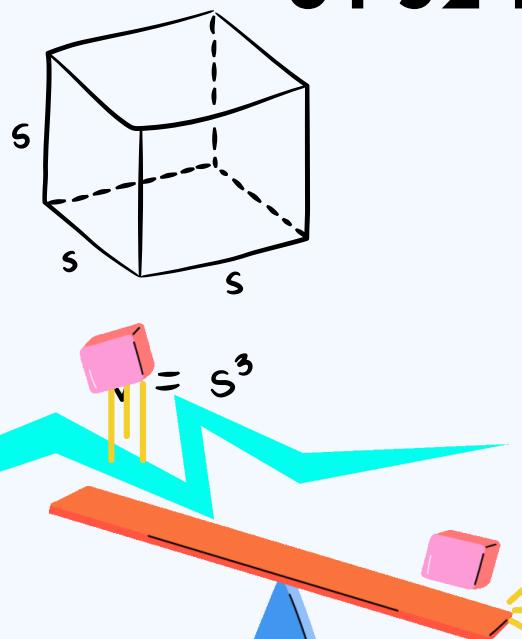
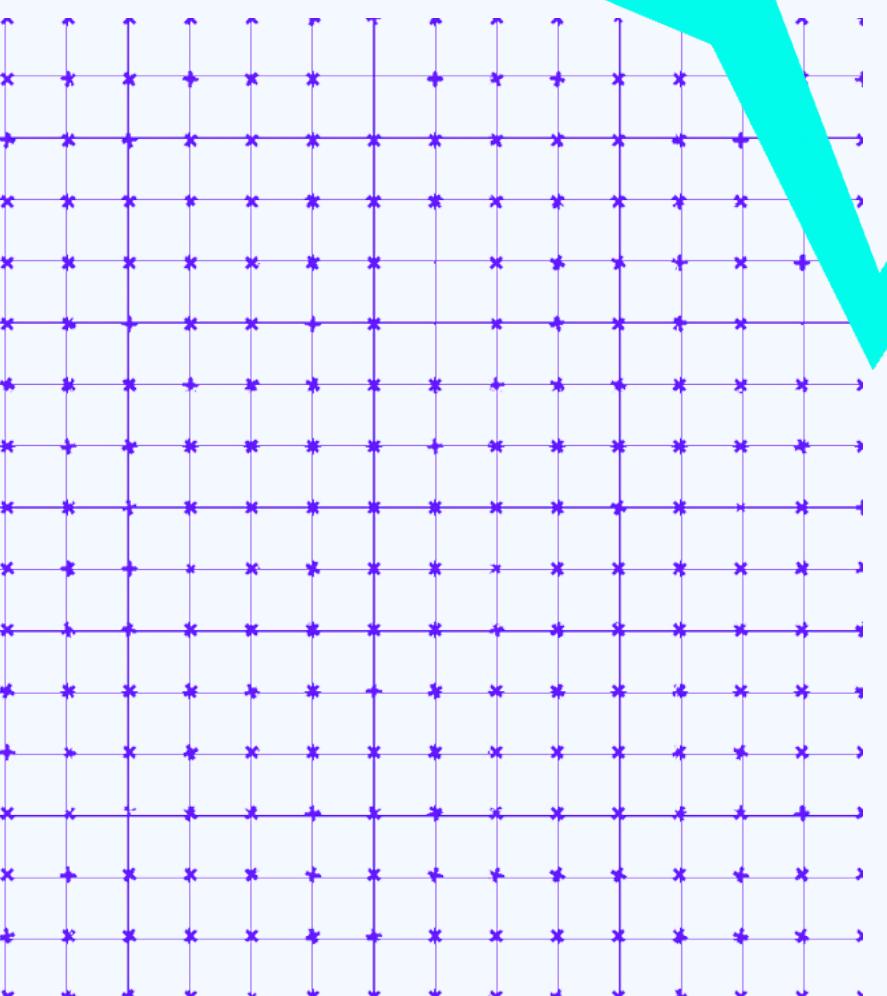
      accuracy      0.46    3036328
macro avg      0.23      0.50      0.31    3036328
weighted avg   0.21      0.46      0.29    3036328
```



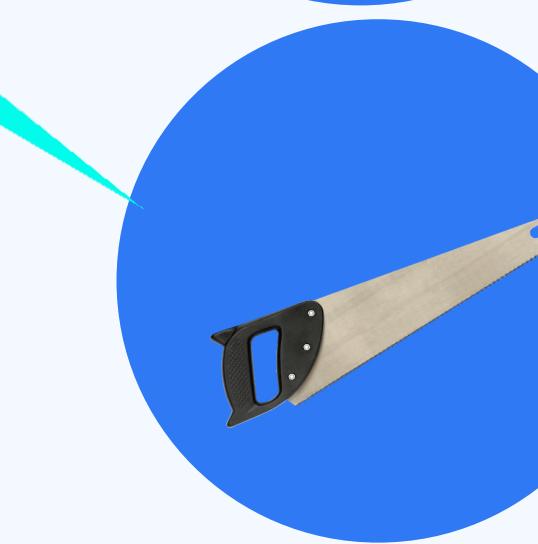


# MI MODELO DE RED NEURONAL RESAMPLING 1

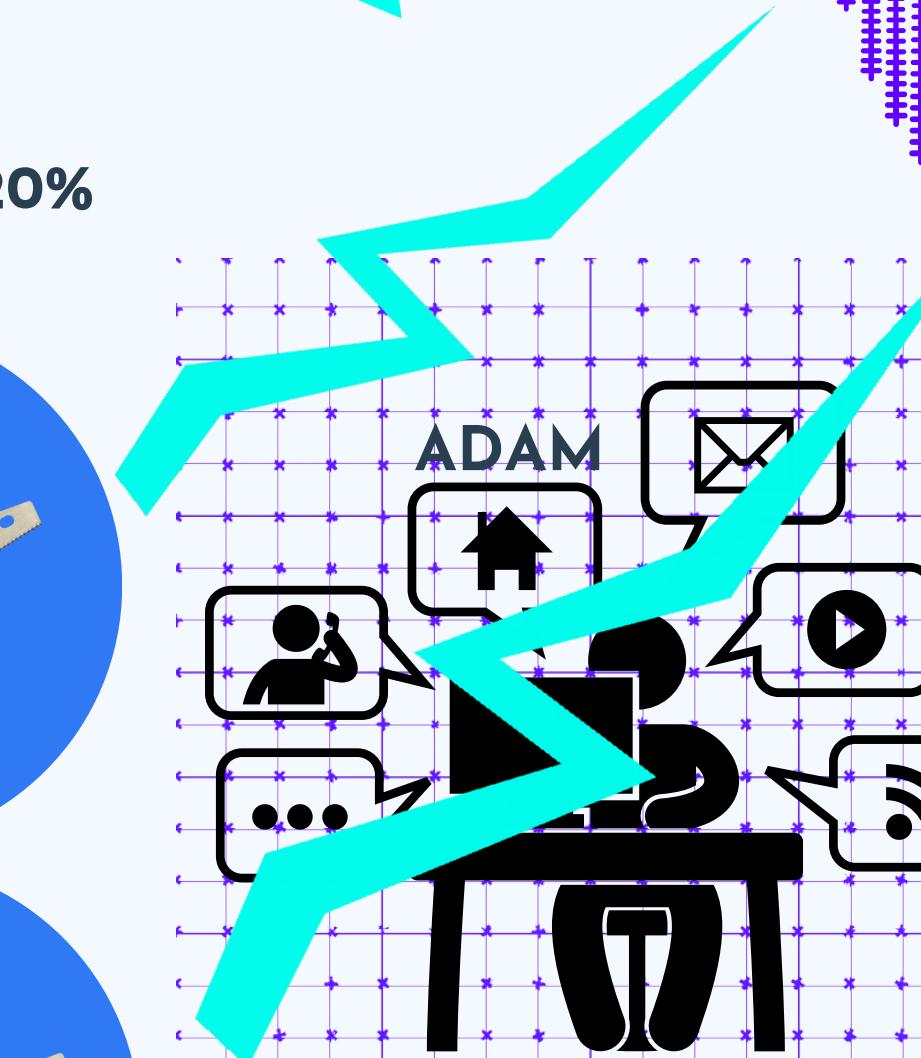
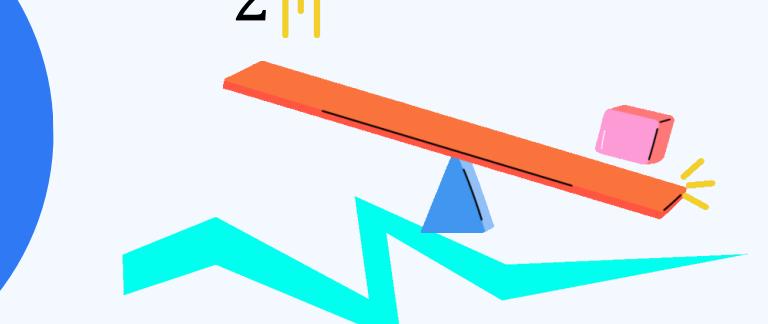
CAPA DE ENTRADA  
64 UNITS



CAPAS OCULTAS  
64-32-16-8 Y L2 1 CAPA Y 2 DENSA



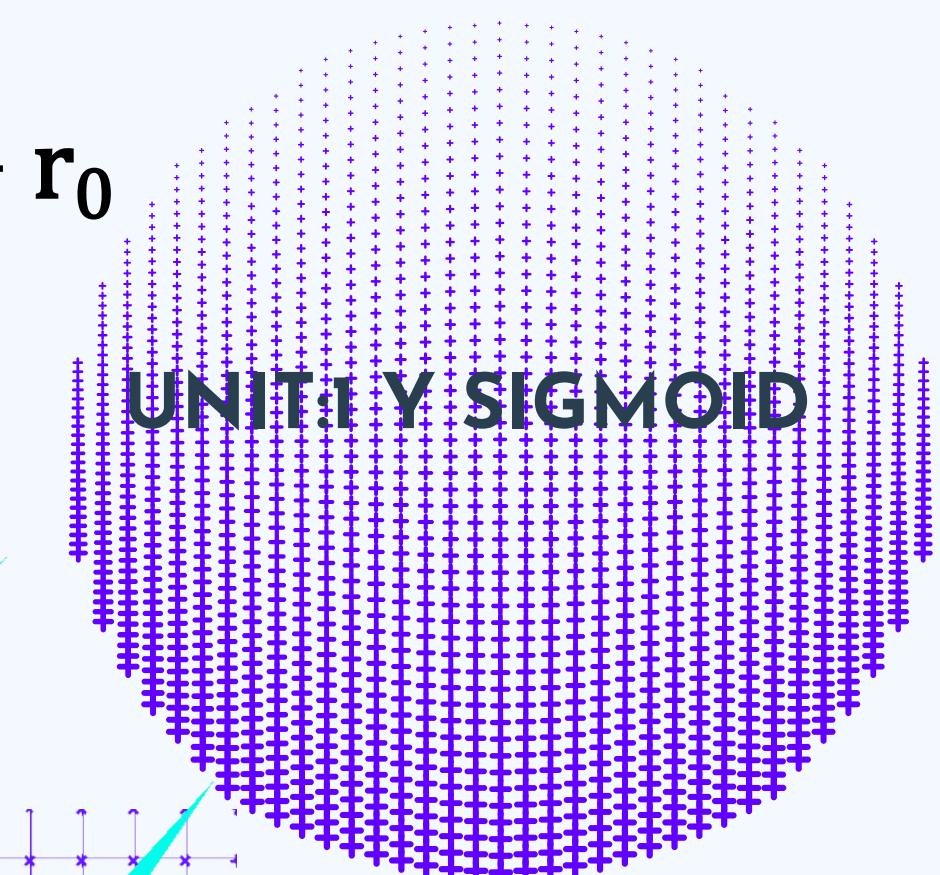
$$r = \frac{1}{2}at^2 + v_0t + r_0$$



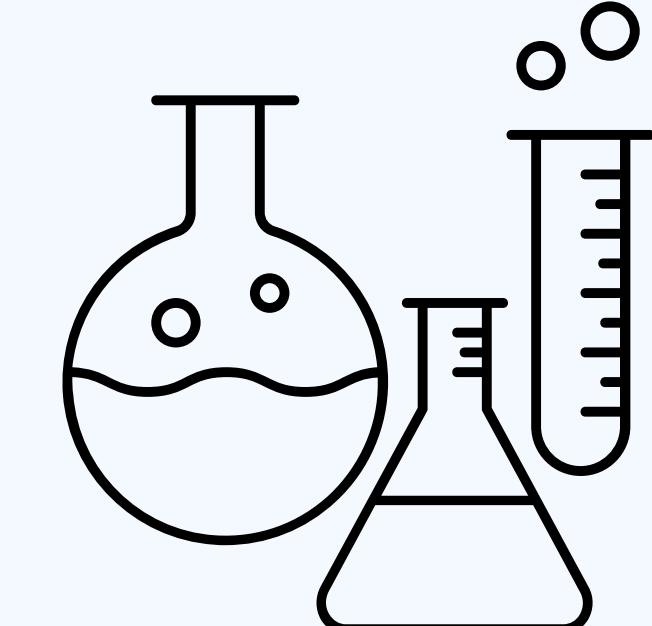
START

CAPAS OCULTAS

CAPA DE SALIDA



EARLY\_STOPPING



REDUCE\_LR



```
# Aplica submuestro, para reducir el número de observaciones de todas las clases menos la clase minoritaria
X=df_rn.drop(["Target"], axis=1)
y=df_rn["Target"]

rus = RandomUnderSampler(random_state=42)
X_resampled, y_resampled = rus.fit_resample(X, y)
```

## modelo de red neuronal 1

```
X_train, X_temp, y_train, y_temp = train_test_split(X_resampled, y_resampled, test_size=0.2, random_state=4
X_test, X_val, y_test, y_val = train_test_split(X_temp, y_temp, test_size=0.25, random_state=42)
```

Loss: 0.43704232573509216

Accuracy: 0.7845828533172607

precision: 0.7084441184997559

recall: 0.9675958752632141

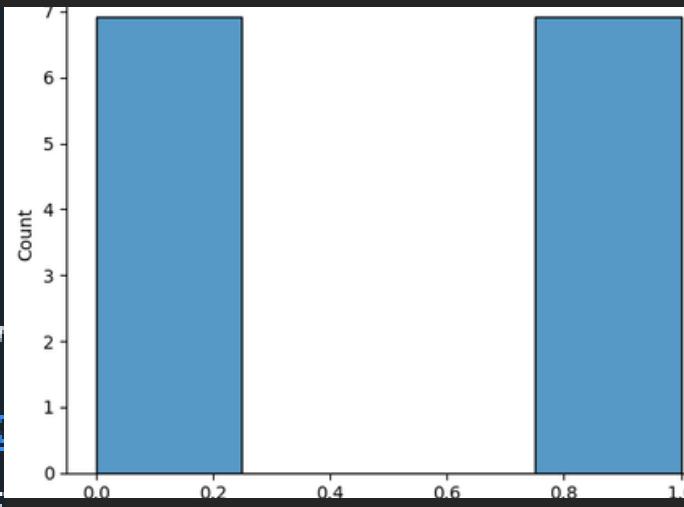
86548/86548 [=====] - 63s 731us/

```
c:\Users\icto\anaconda3\envs\tf-gpu\lib\site-packages\sk
    _warn_prf(average, modifier, f"{metric.capitalize()} is
c:\Users\icto\anaconda3\envs\tf-gpu\lib\site-packages\sk
    _warn_prf(average, modifier, f"{metric.capitalize()} is
```

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0.0	0.50	1.00	0.67	1384044
1.0	0.00	0.00	0.00	1385471

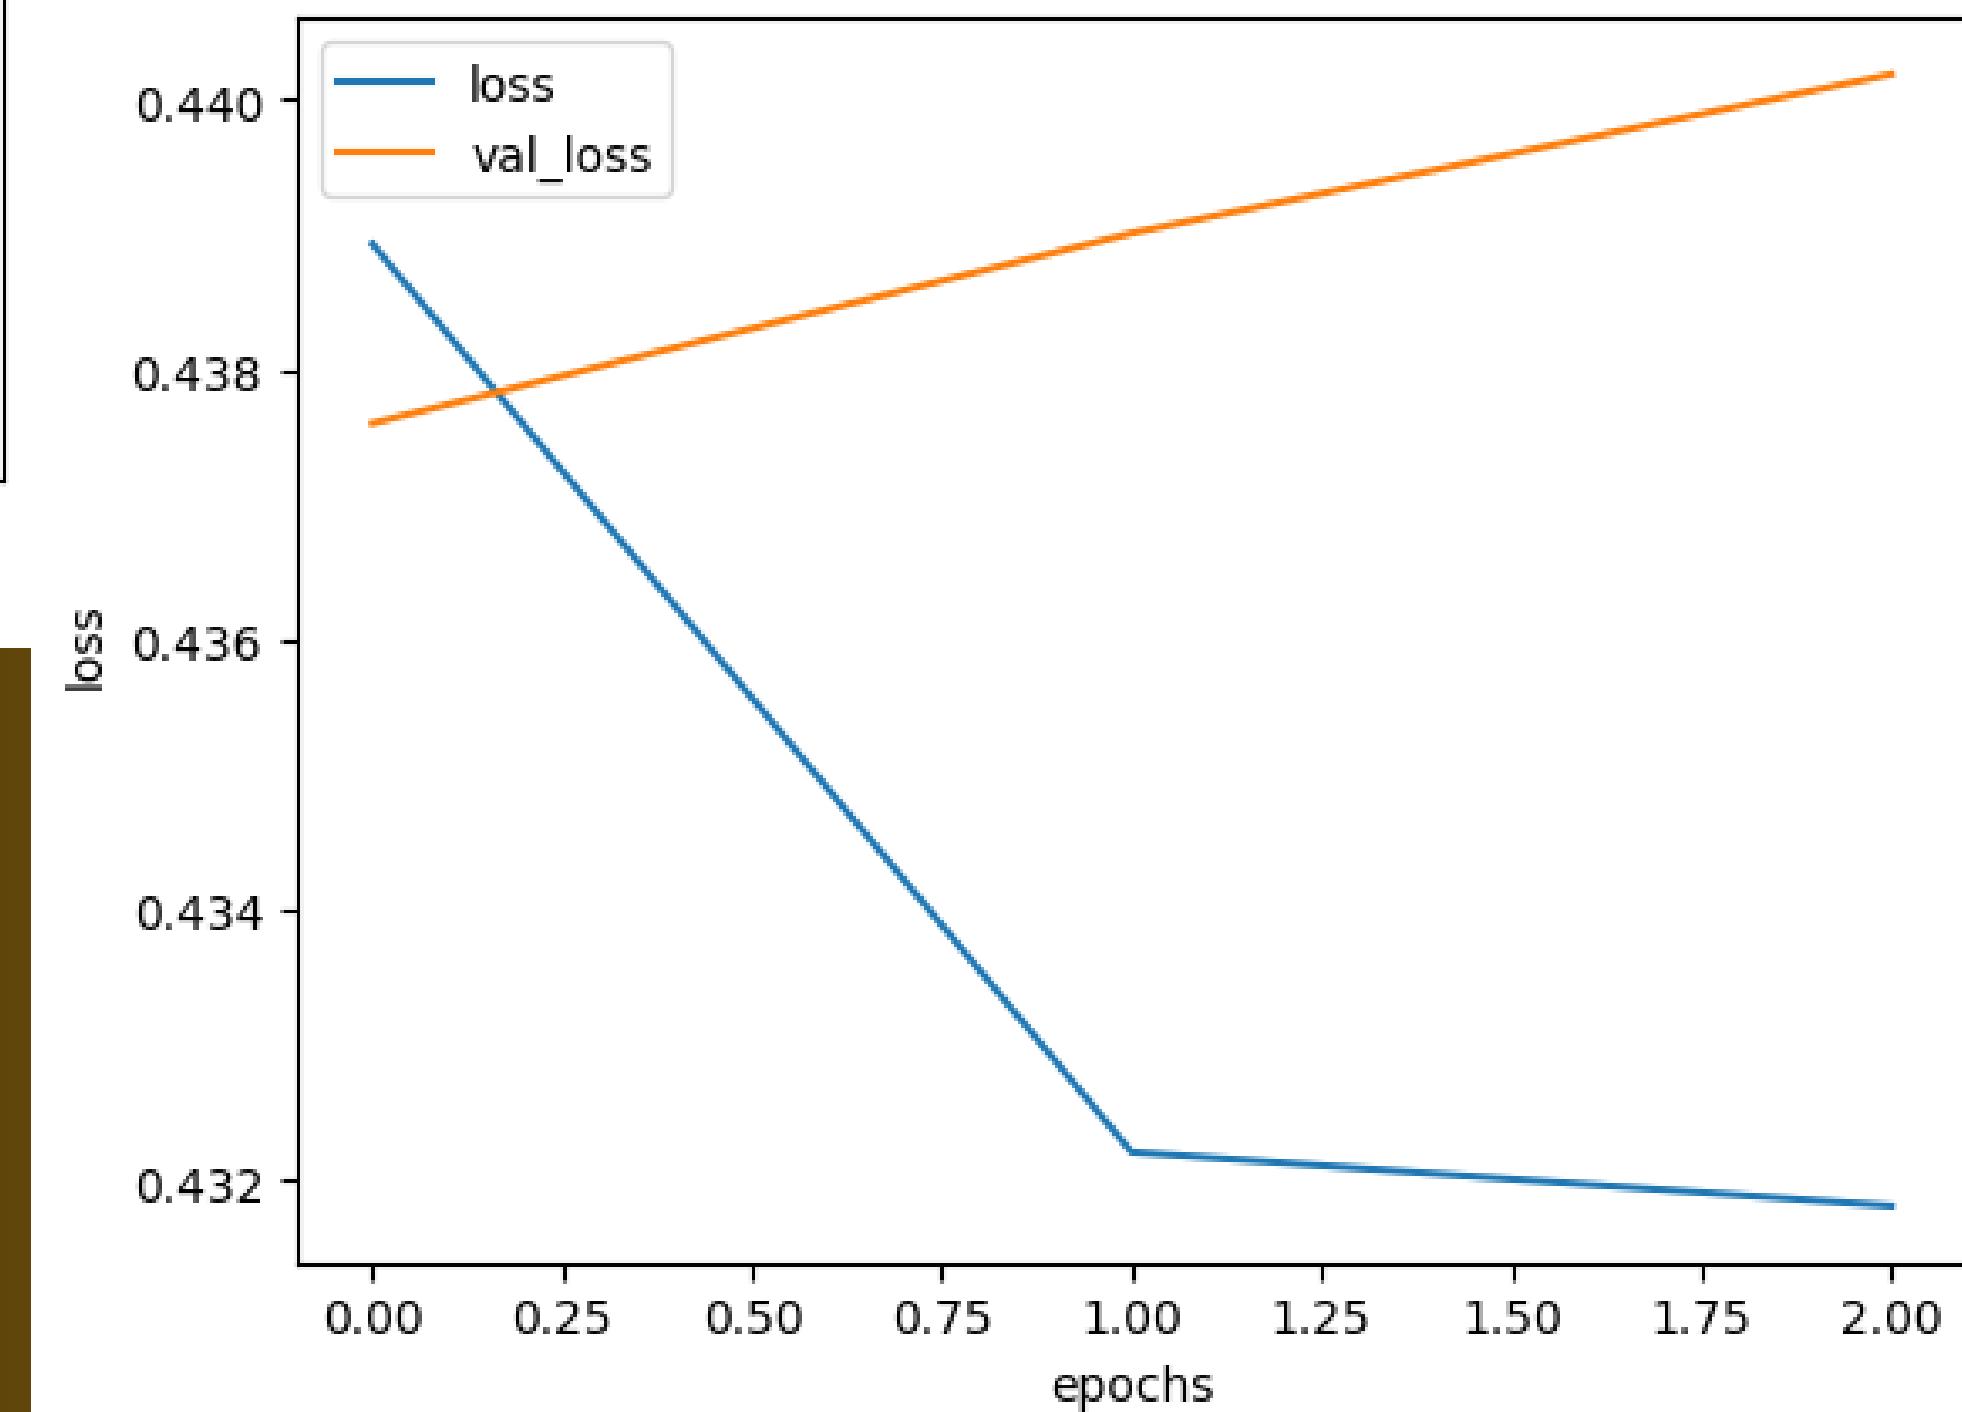
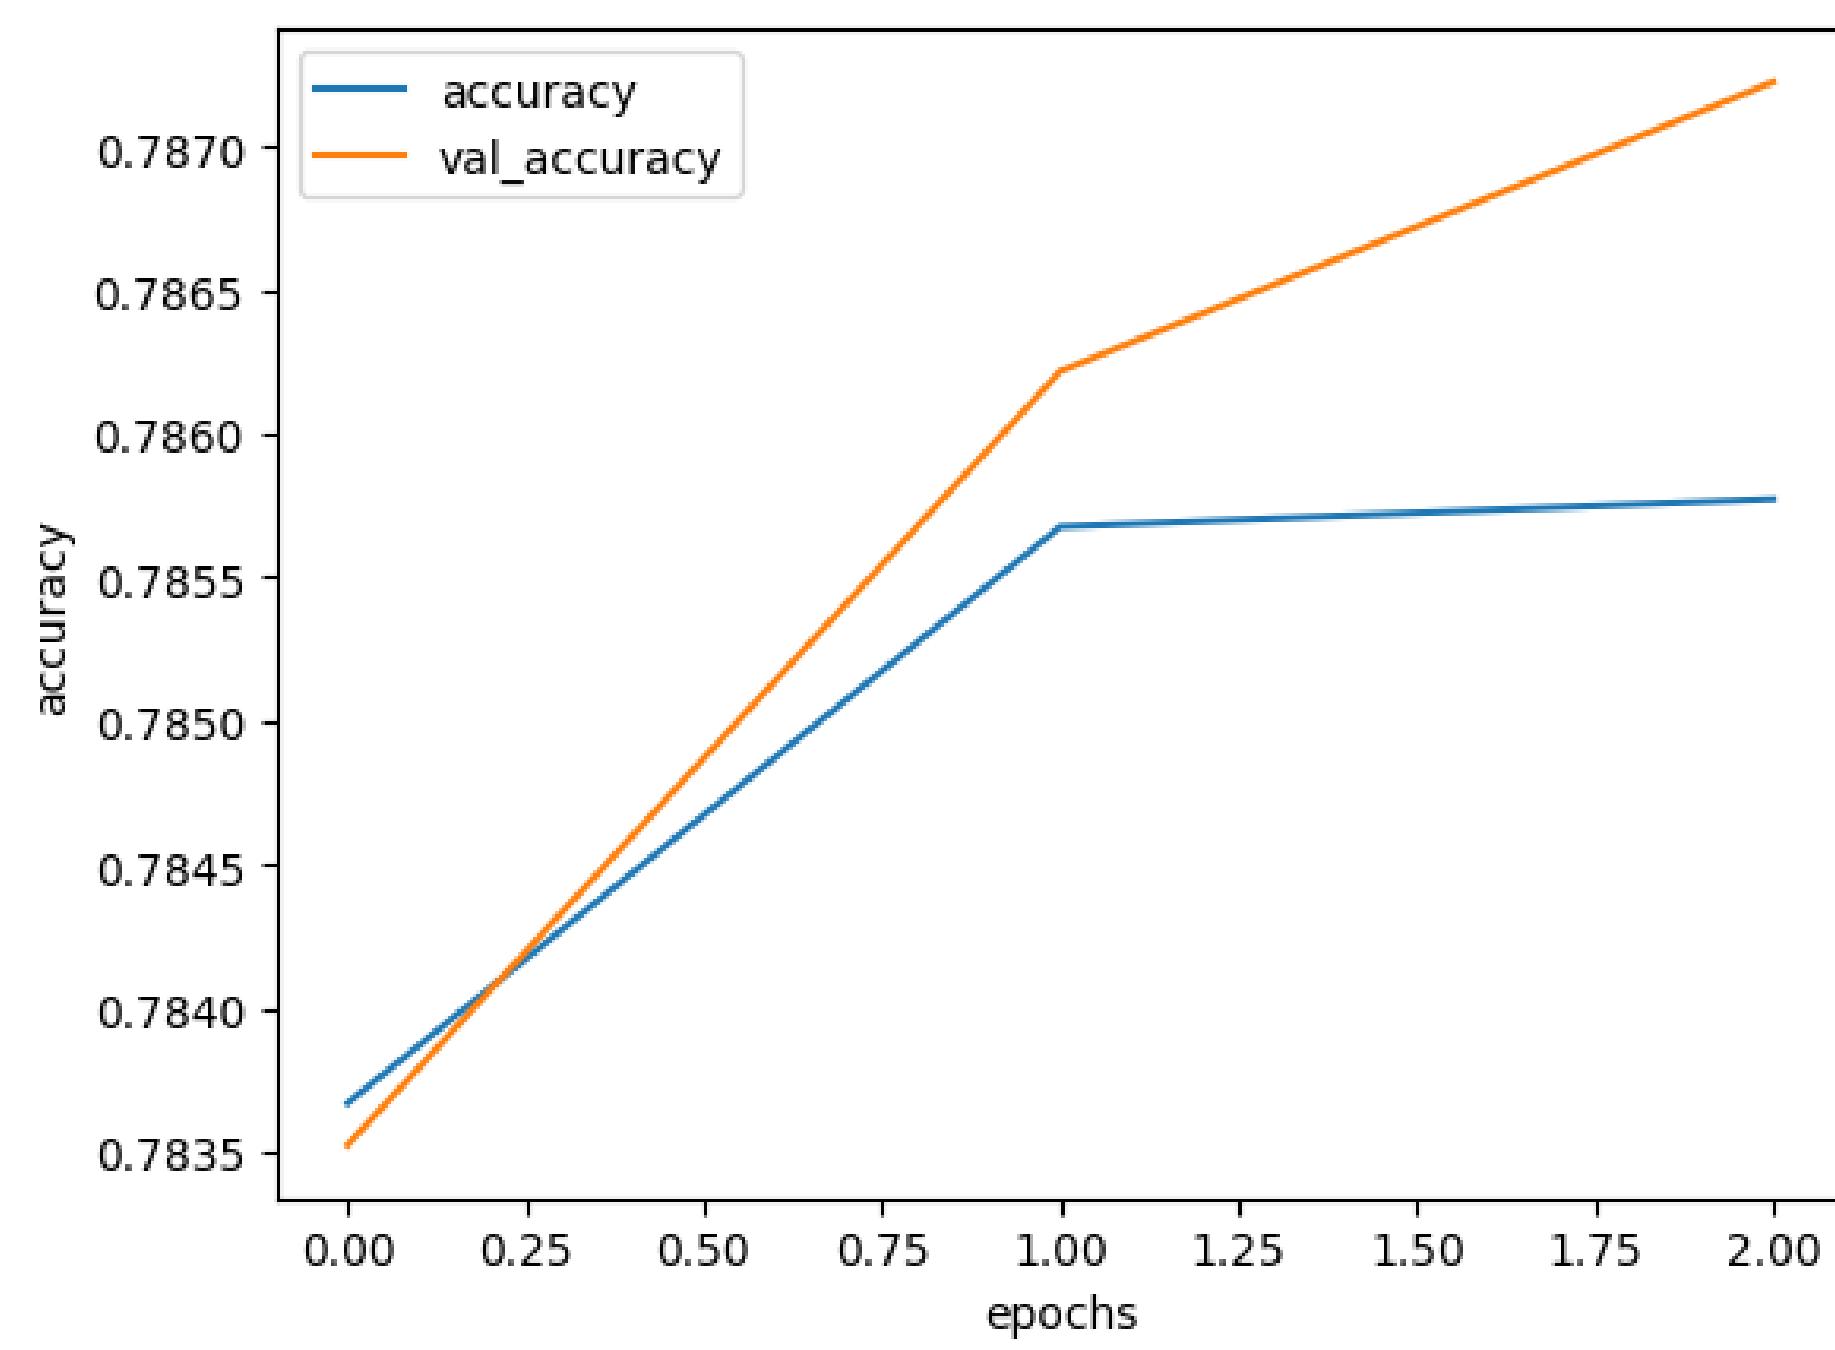
accuracy			0.50	2769515
macro avg	0.25	0.50	0.33	2769515
weighted avg	0.25	0.50	0.33	2769515

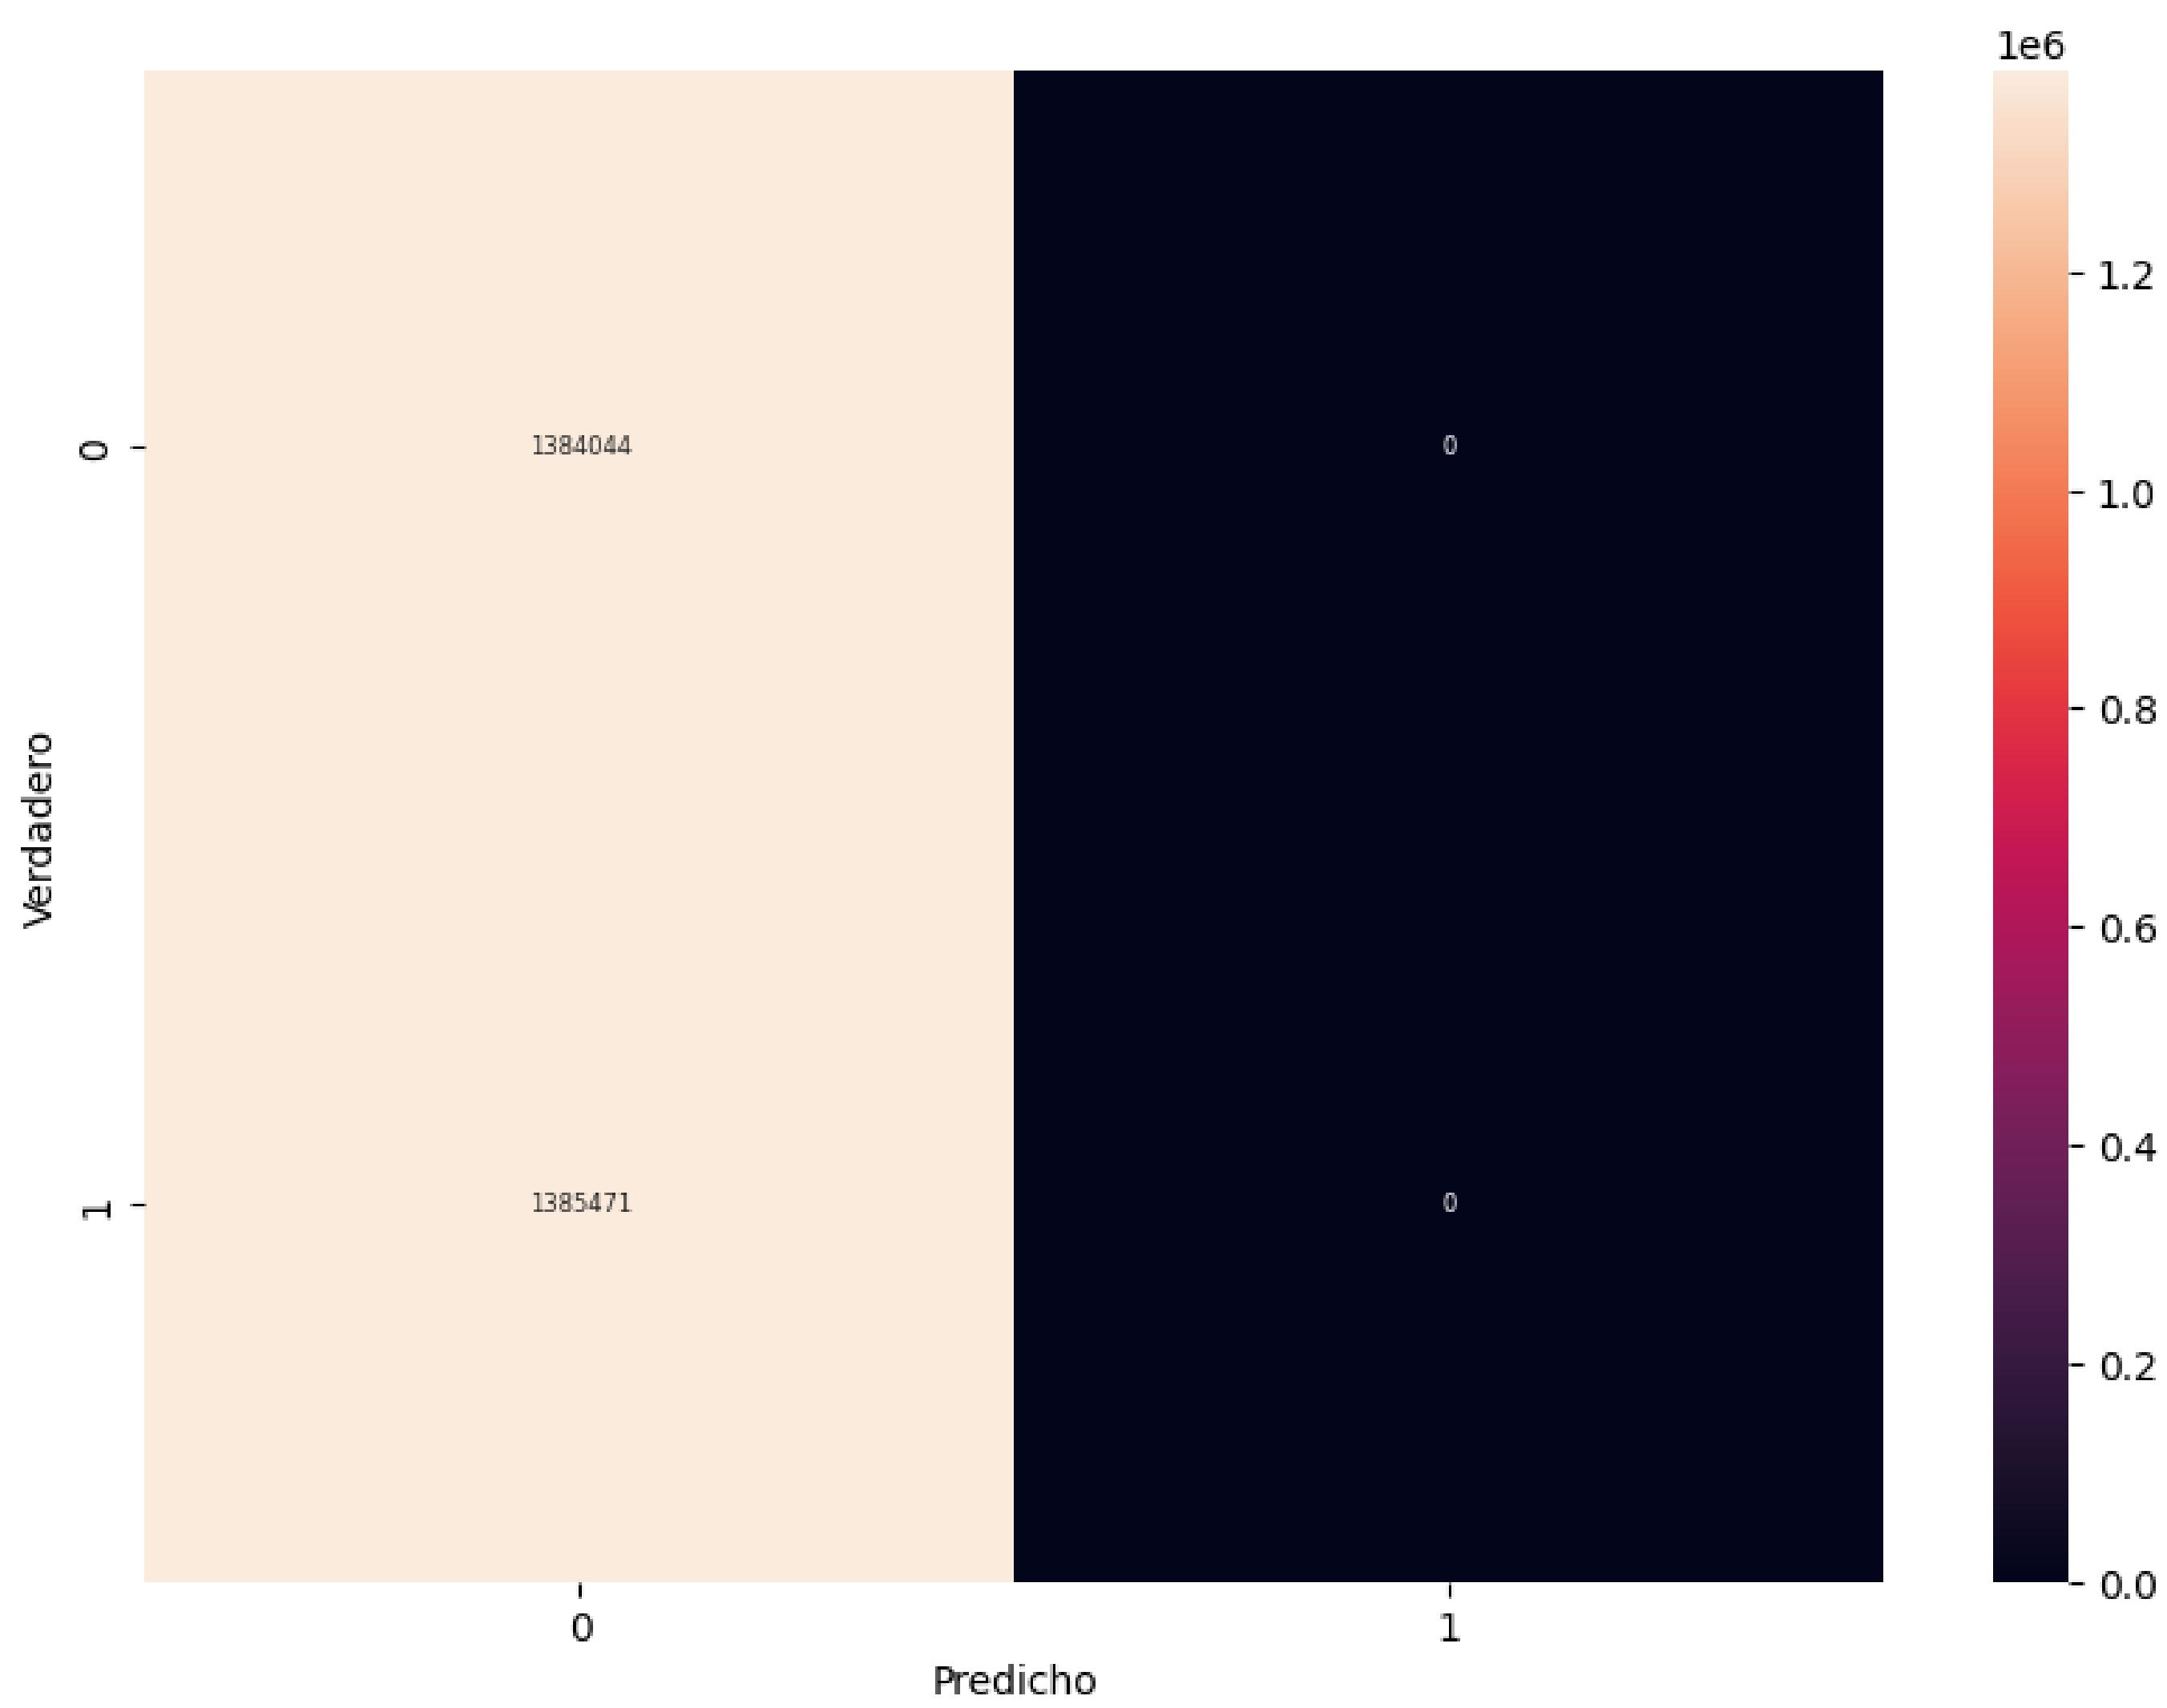


GAME  
ON

50%

ACIERTOS  
CERTEROS

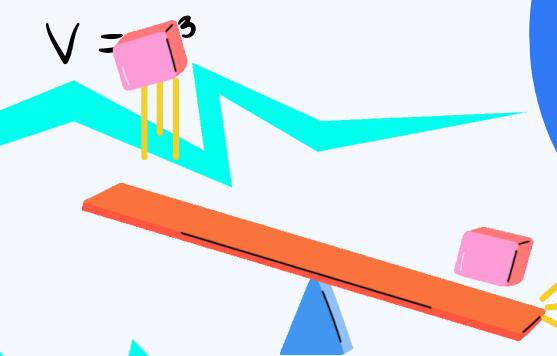
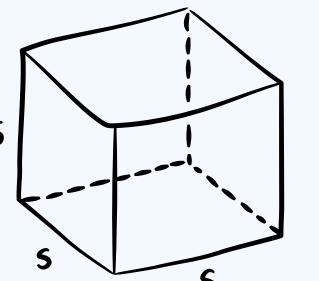




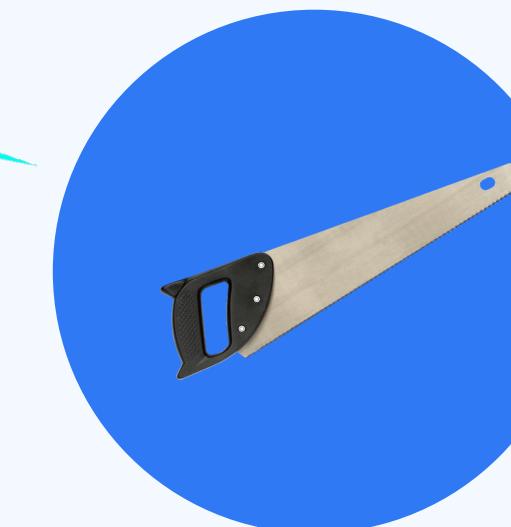
# MI MODELO DE RED NEURONAL RESAMPLING 2

CAPA DE SALIDA

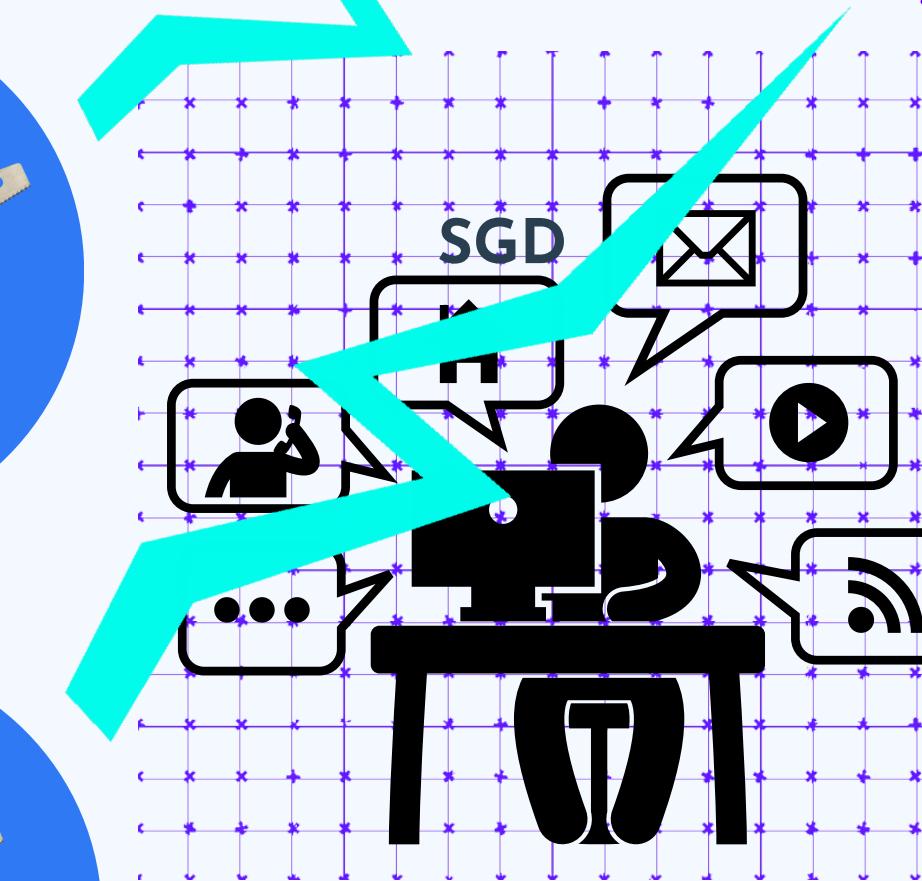
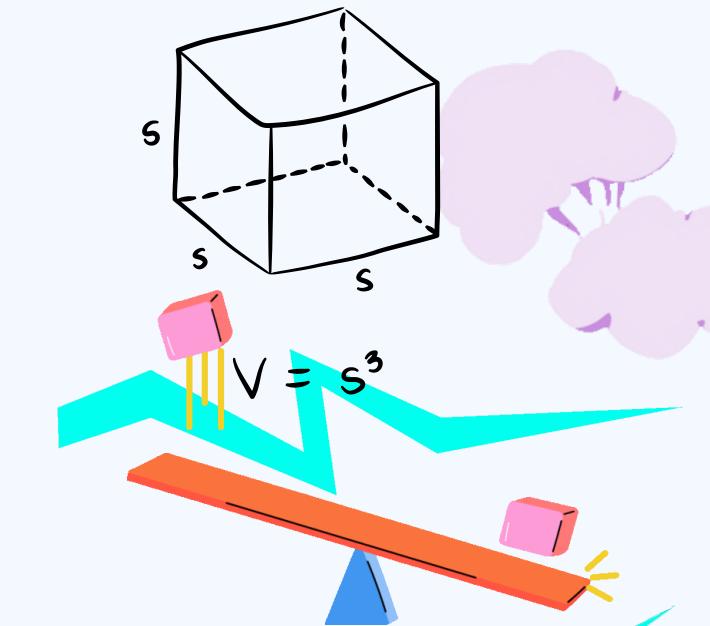
CAPA DE ENTRADA



CAPAS OCULTAS

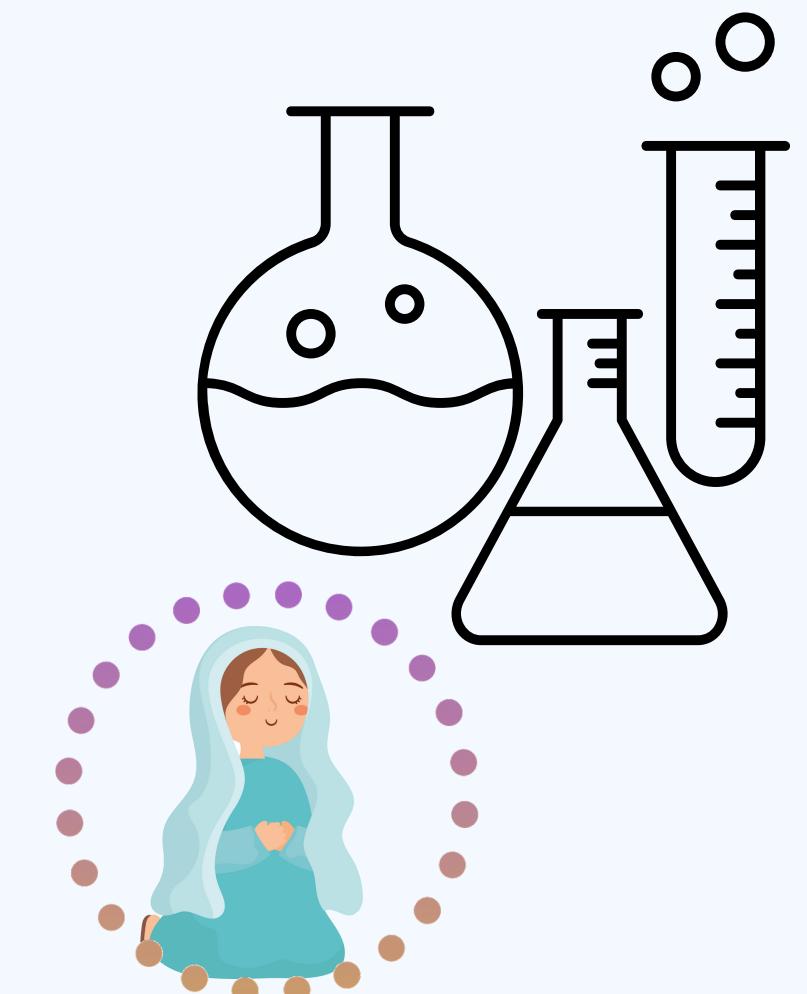


TODAS LAS CAPAS  
TIENEN 64 UNITS



UNIT 1 Y SIGMOID

EARLY\_STOPPING



Loss: 0.43704232573509216

Accuracy: 0.7845828533172607

precision: 0.7084441184997559

recall: 0.9675958752632141

86548/86548 [=====] - 63s 731us/step

c:\Users\vit\anaconda3\envs\tf-gpu\lib\site-packages\sklearn\

\_warn\_prf(average, modifier, f'{metric.capitalize()} is', len

c:\Users\vit\anaconda3\envs\tf-gpu\lib\site-packages\sklearn\

\_warn\_prf(average, modifier, f'{metric.capitalize()} is', len

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

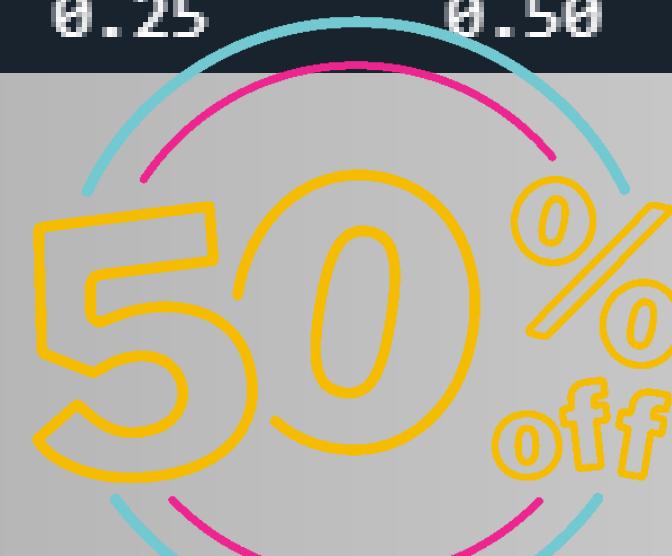
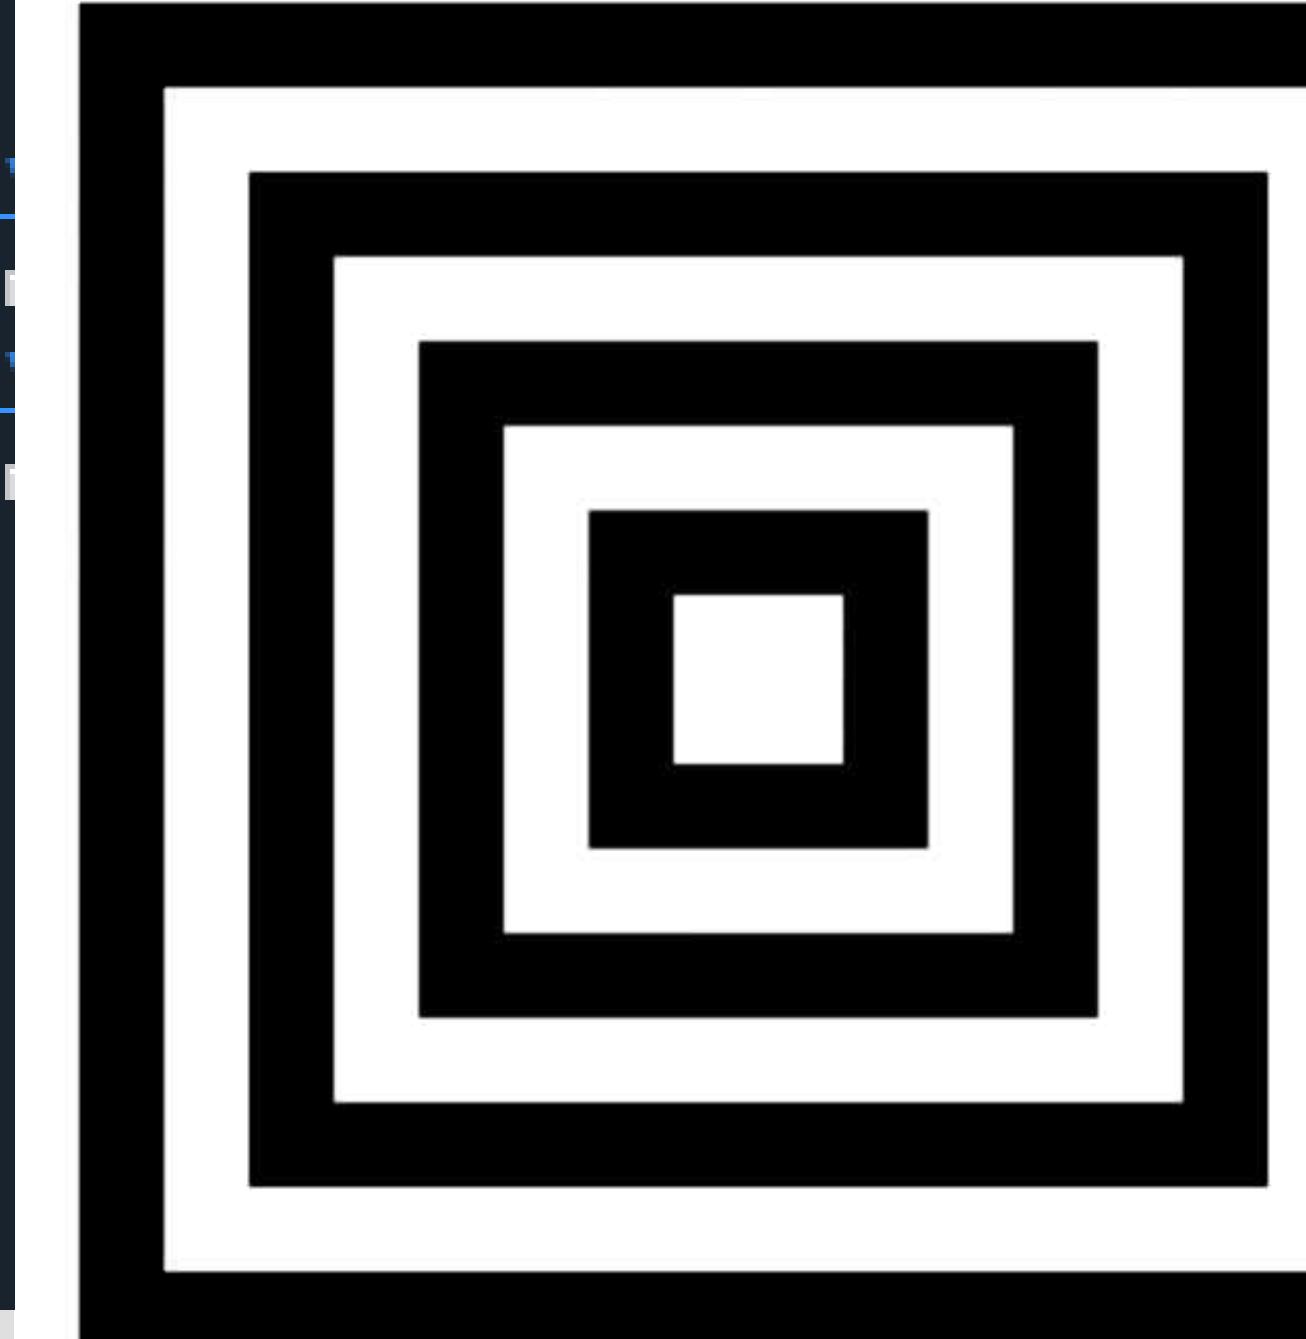
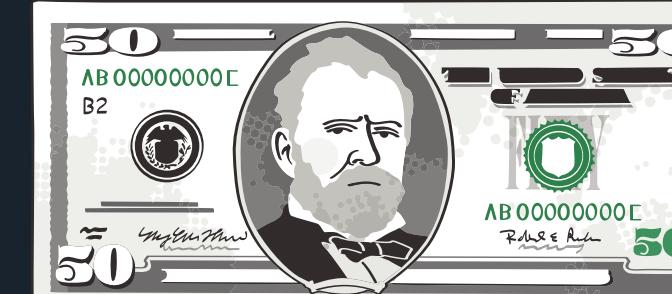
0.0	0.50	1.00	0.67	1384044
-----	------	------	------	---------

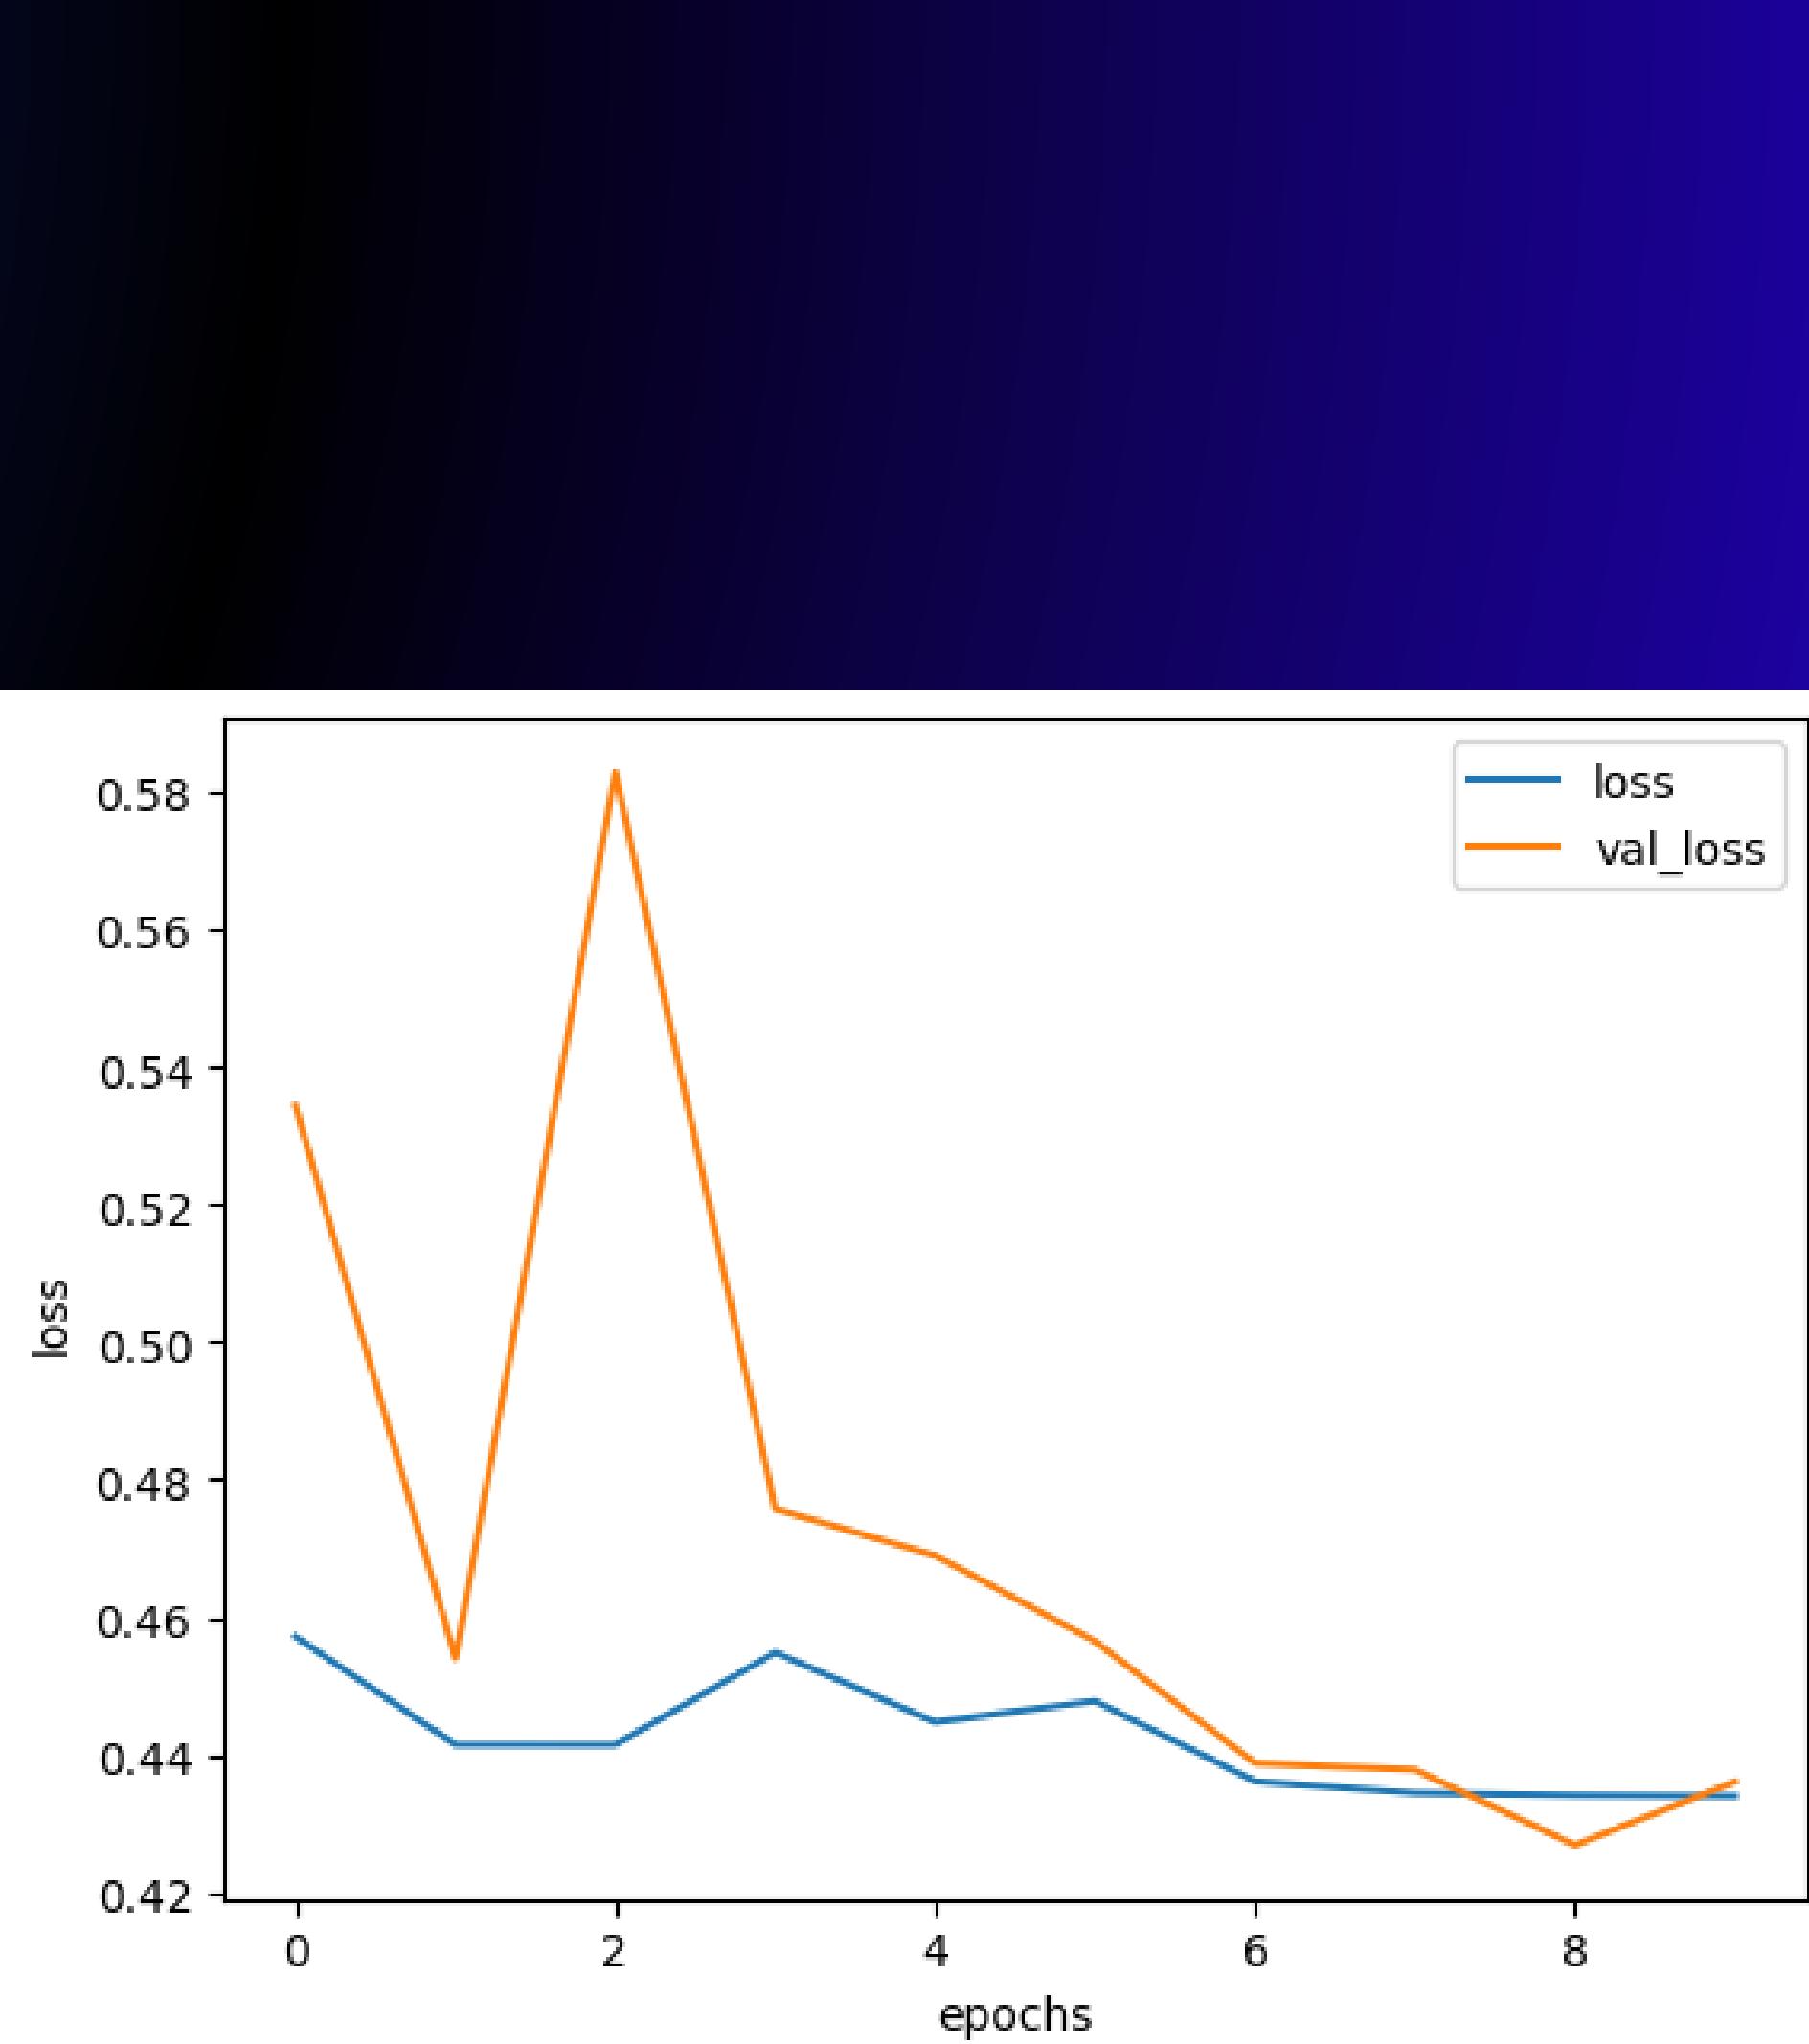
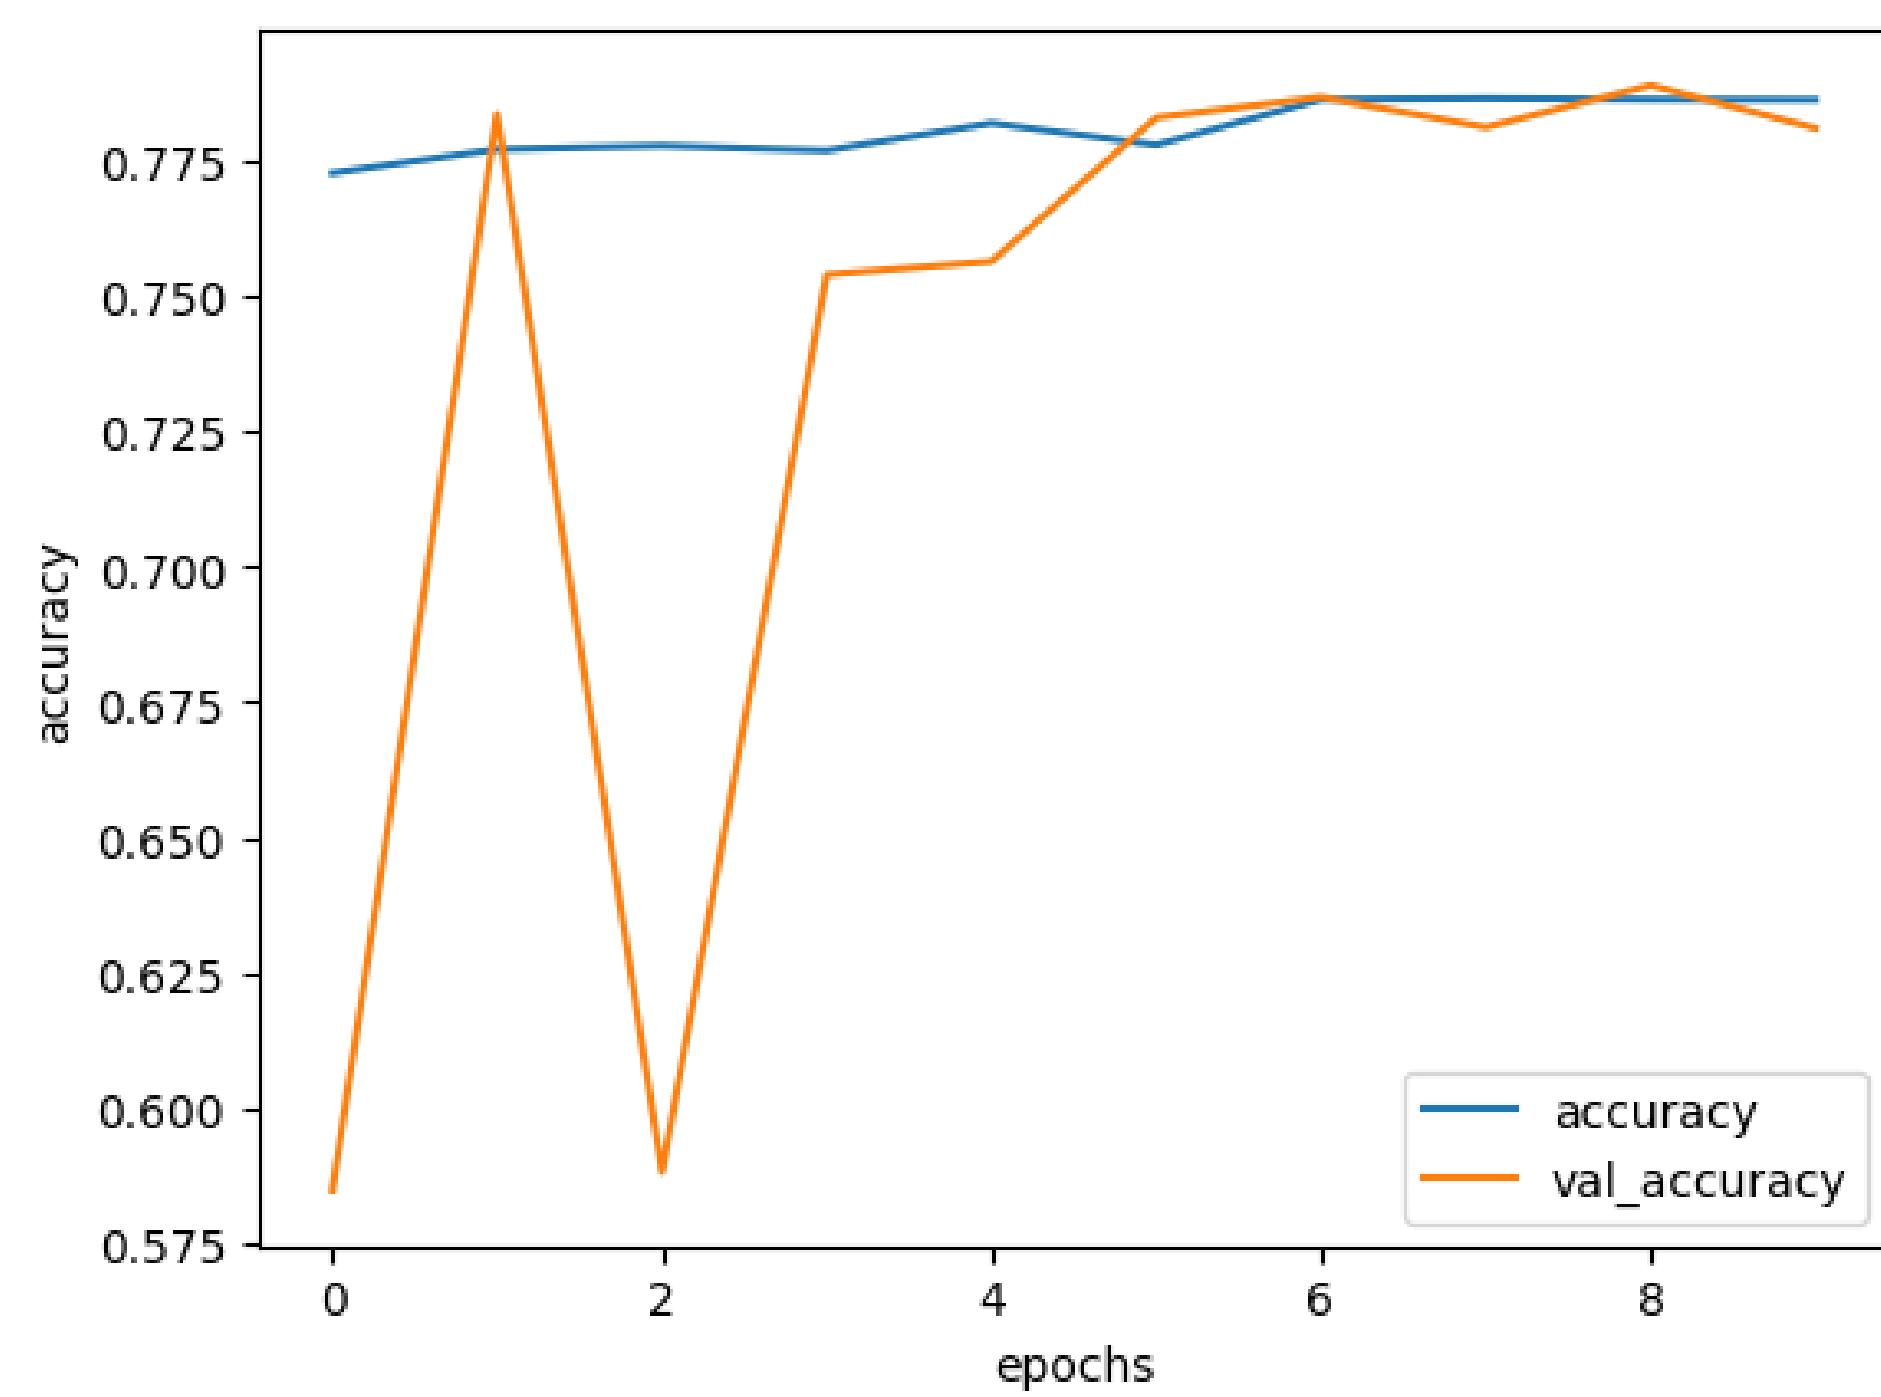
1.0	0.00	0.00	0.00	1385471
-----	------	------	------	---------

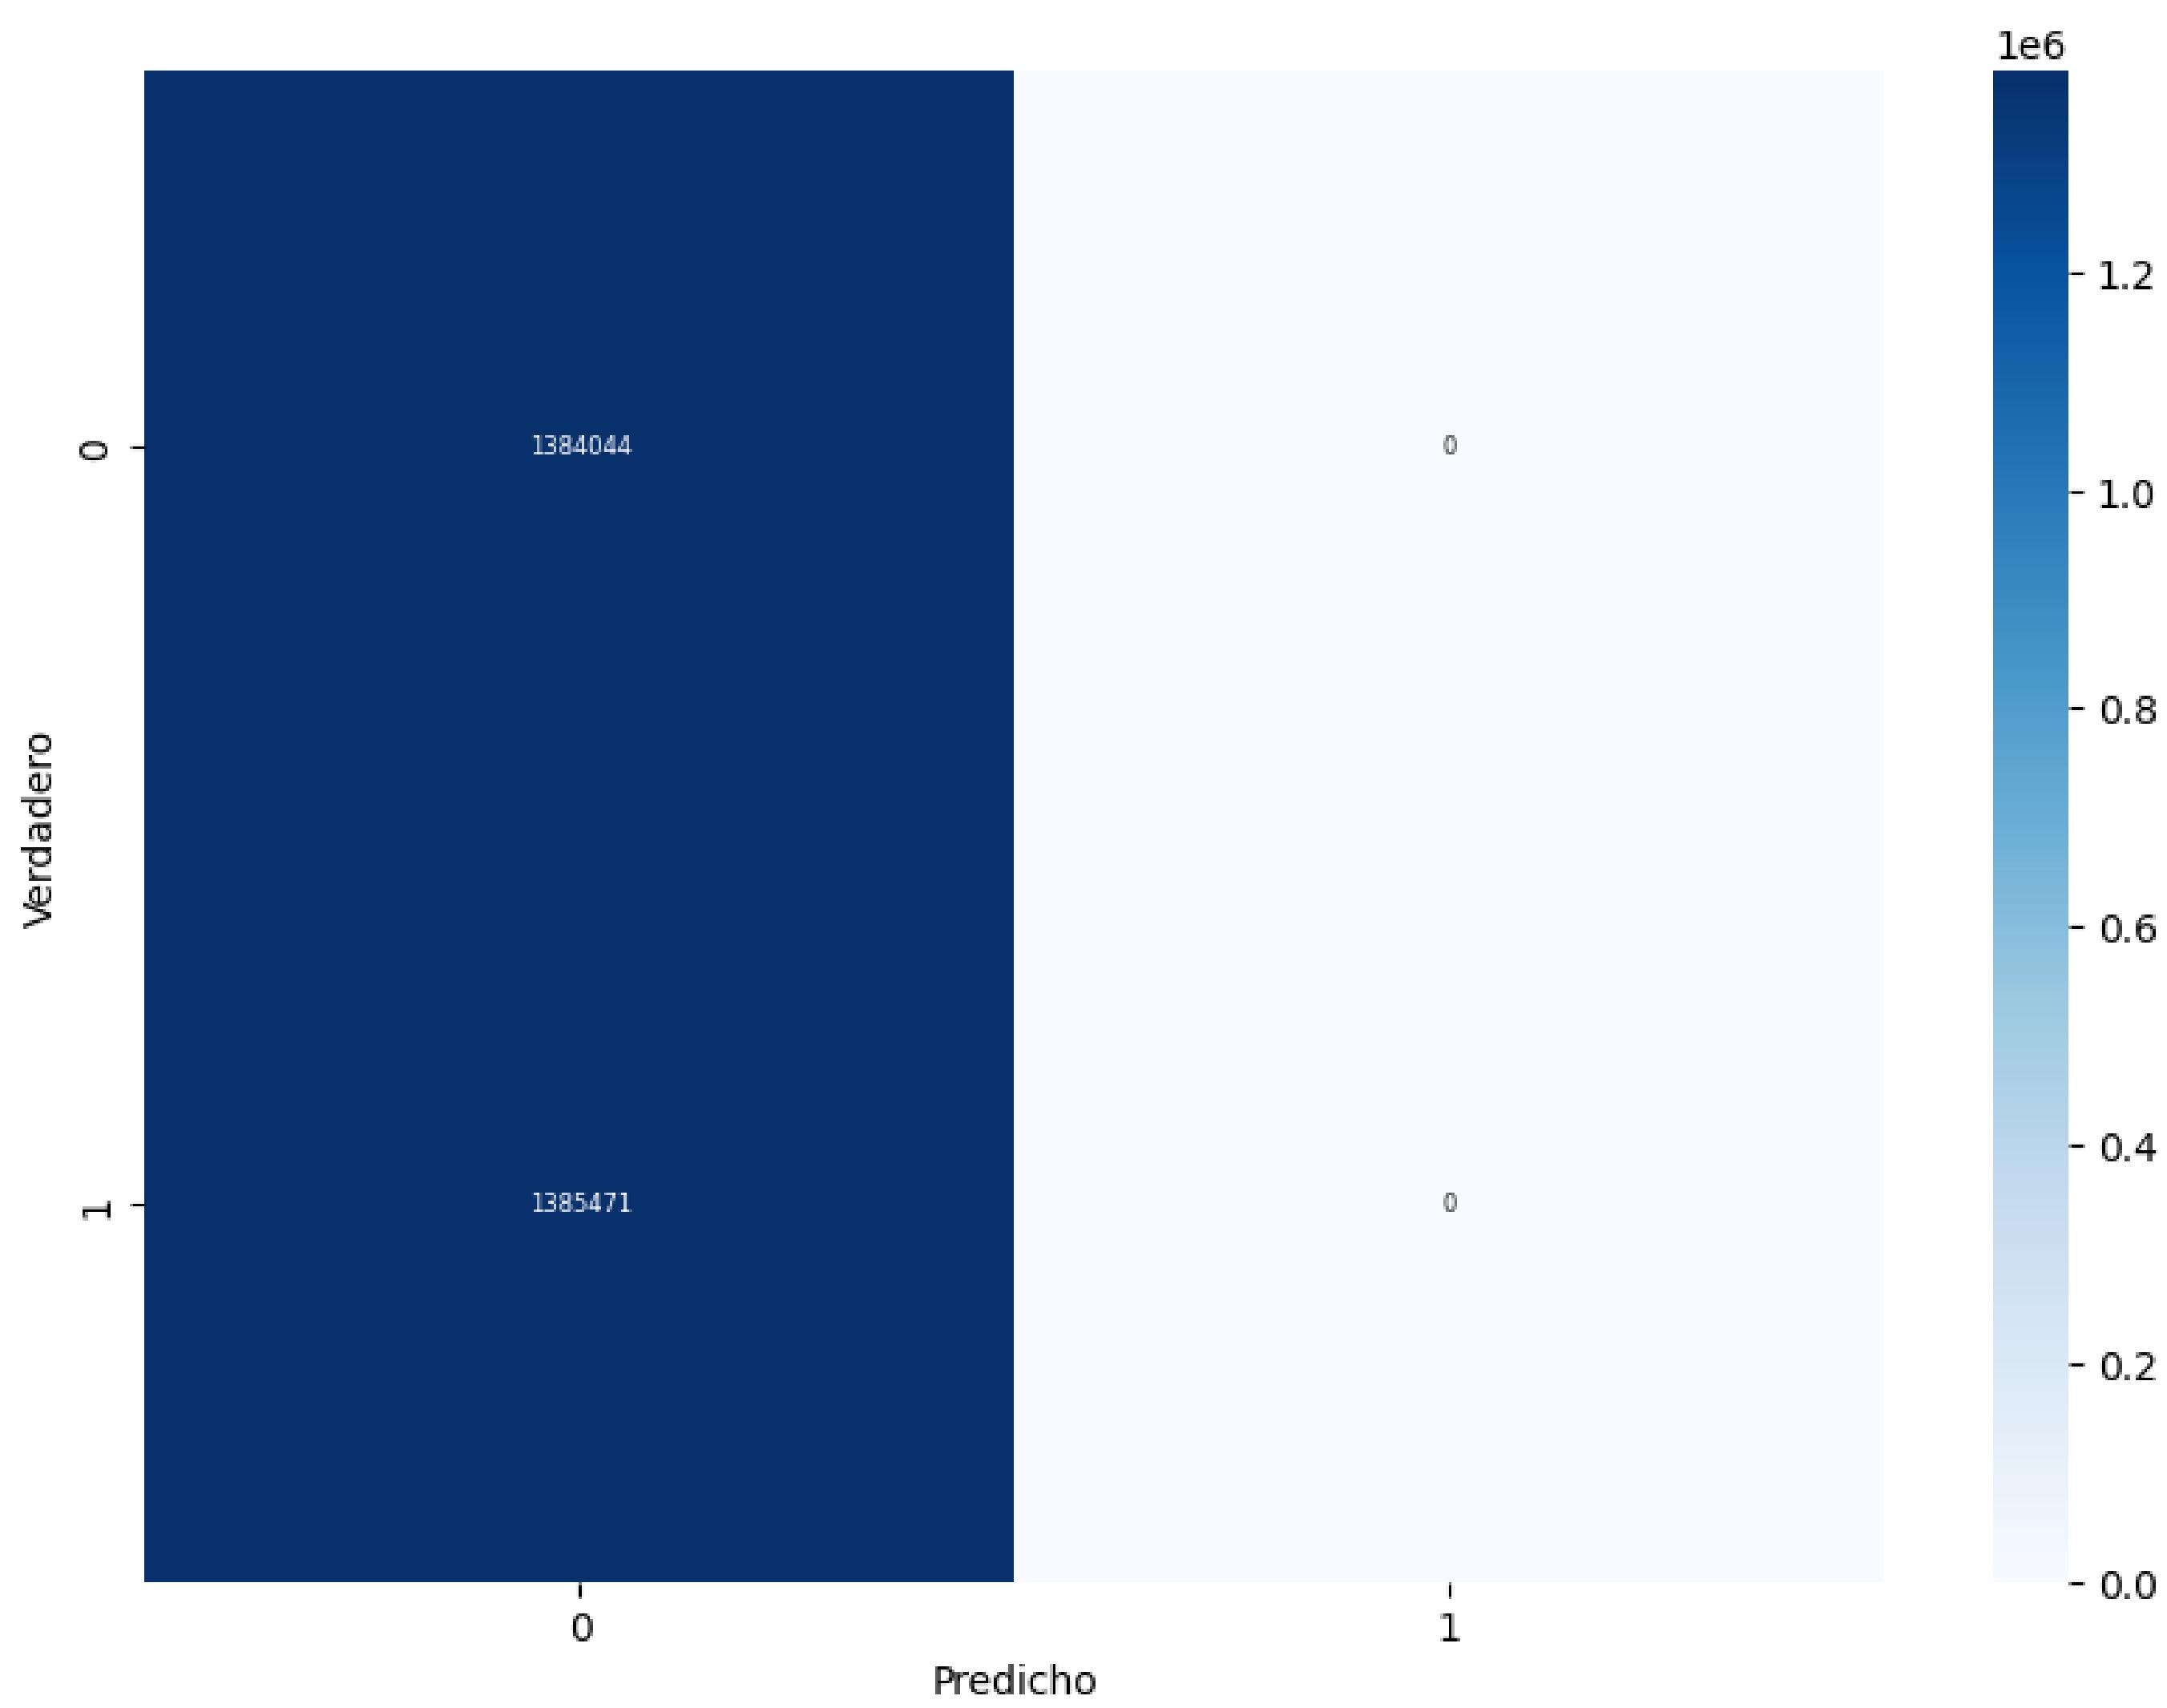
accuracy			0.50	2769515
----------	--	--	------	---------

macro avg	0.25	0.50	0.33	2769515
-----------	------	------	------	---------

weighted avg	0.25	0.50	0.33	2769515
--------------	------	------	------	---------







# CONCLUSIONES FINALES

El modelo que mejor se adapta a nuestro objetivo de negocio es el modelo RandomForest en multiproceso sin resampling con un 58% de aciertos en la clase Malicious, aunque falle en reconocer los archivos benignos en un 25%. Los modelos de DL el mejor ha sido con resampling consiguiendo un 50% de verdaderos positivos sin fallos , lo que ocurre que no podemos valorar la otra clase al no ser reconocida . ,

En resumen, , un casi 60% de positivos verdaderos totalmente fiables, es una buena proporción para nuestro objetivo, ahora le voy a pasar el test de prueba a este modelo para su ultima comprobacion

## TEST DE VALIDACION

```
# Entrenar el modelo
with parallel_backend('threading'):
    model.fit(X_train, y_train)

# Evaluar el modelo
y_pred = model.predict(X_val)
accuracy = accuracy_score(y_val, y_pred)

print(classification_report(y_pred, y_val))
```

	precision	recall	f1-score	support
0	0.58	1.00	0.74	808056
1	1.00	0.74	0.85	2228272
accuracy			0.81	3036328
macro avg	0.79	0.87	0.79	3036328
weighted avg	0.89	0.81	0.82	3036328

## TEST DE PRUEBA

```
with parallel_backend('threading'):
    model.fit(X_train, y_train)

# Evaluar el modelo
y_pred = model.predict(X_test)
accuracy = accuracy_score(y_test, y_pred)

print(classification_report(y_pred, y_test))
```

✓ 4m 28.7s

	precision	recall	f1-score	support
0	0.58	1.00	0.74	808574
1	1.00	0.74	0.85	2227754
accuracy			0.81	3036328
macro avg	0.79	0.87	0.79	3036328
weighted avg	0.89	0.81	0.82	3036328



**finally  
FINISHED**

**THANK YOU  
SO MUCH**

